

# Database Design Tuning



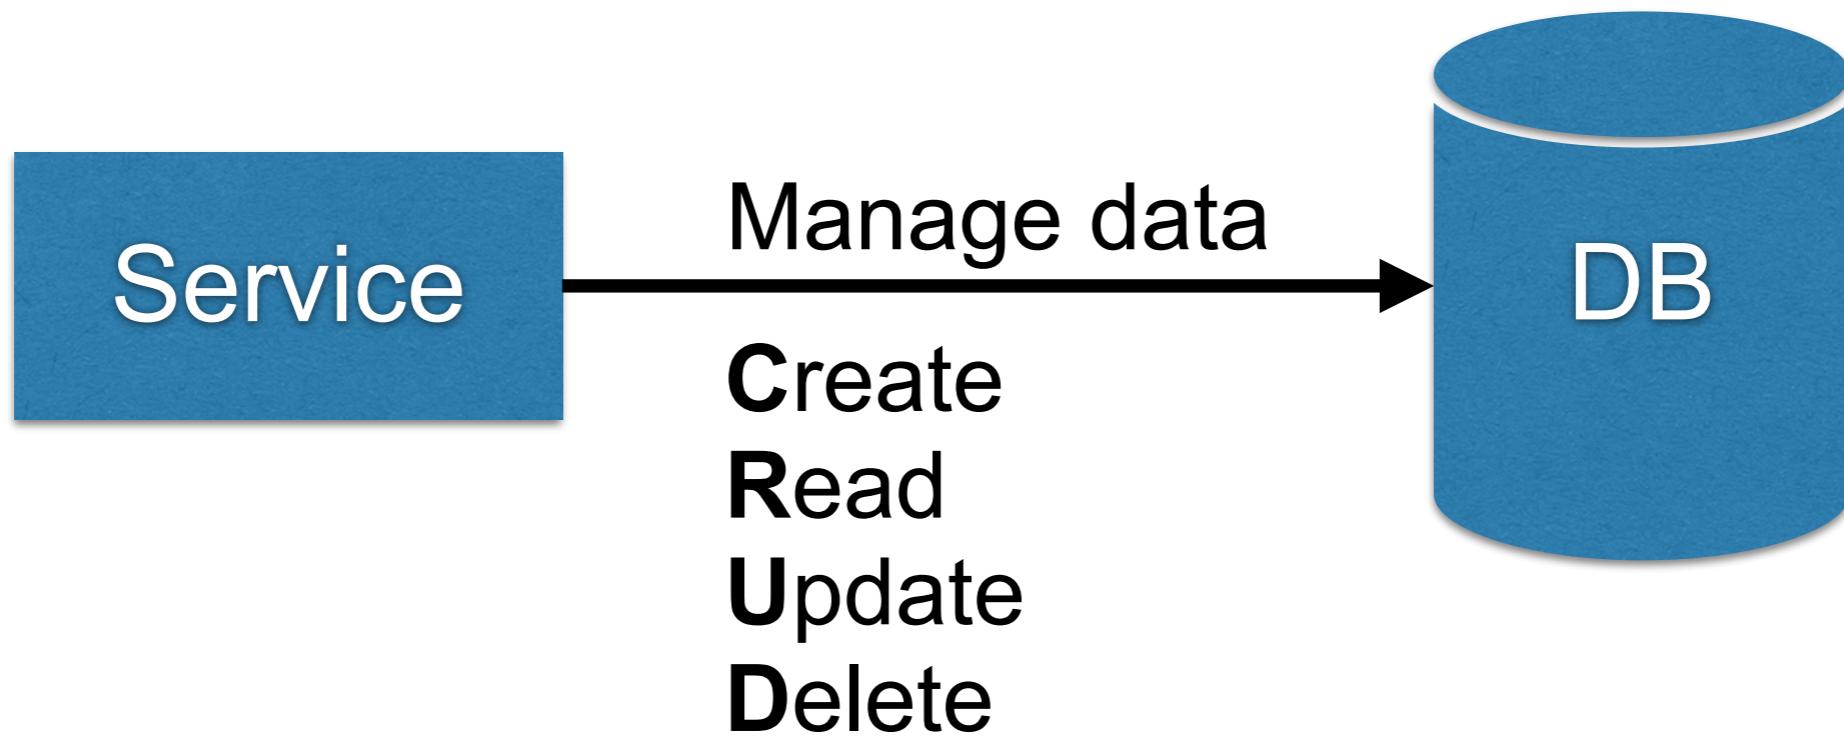
**[https://github.com/up1/  
workshop-database](https://github.com/up1/workshop-database)**



# Database

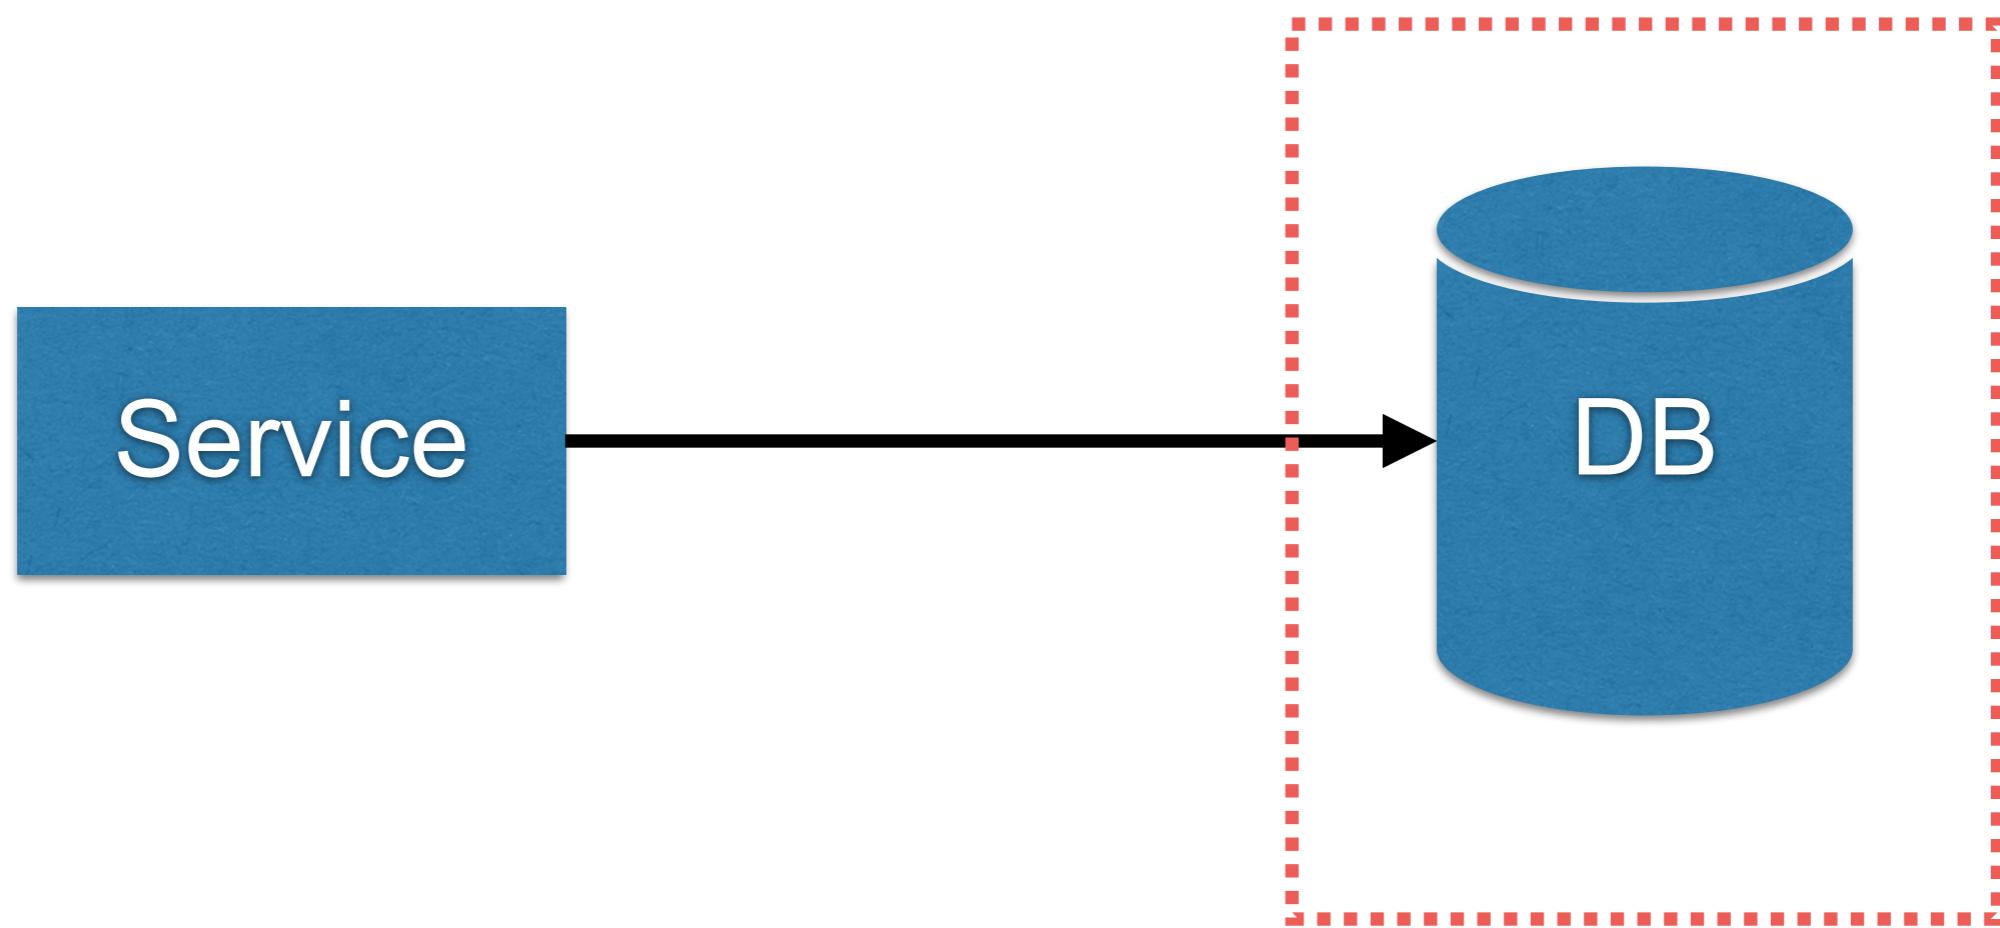


# Database



# Working with Database

Directly to database

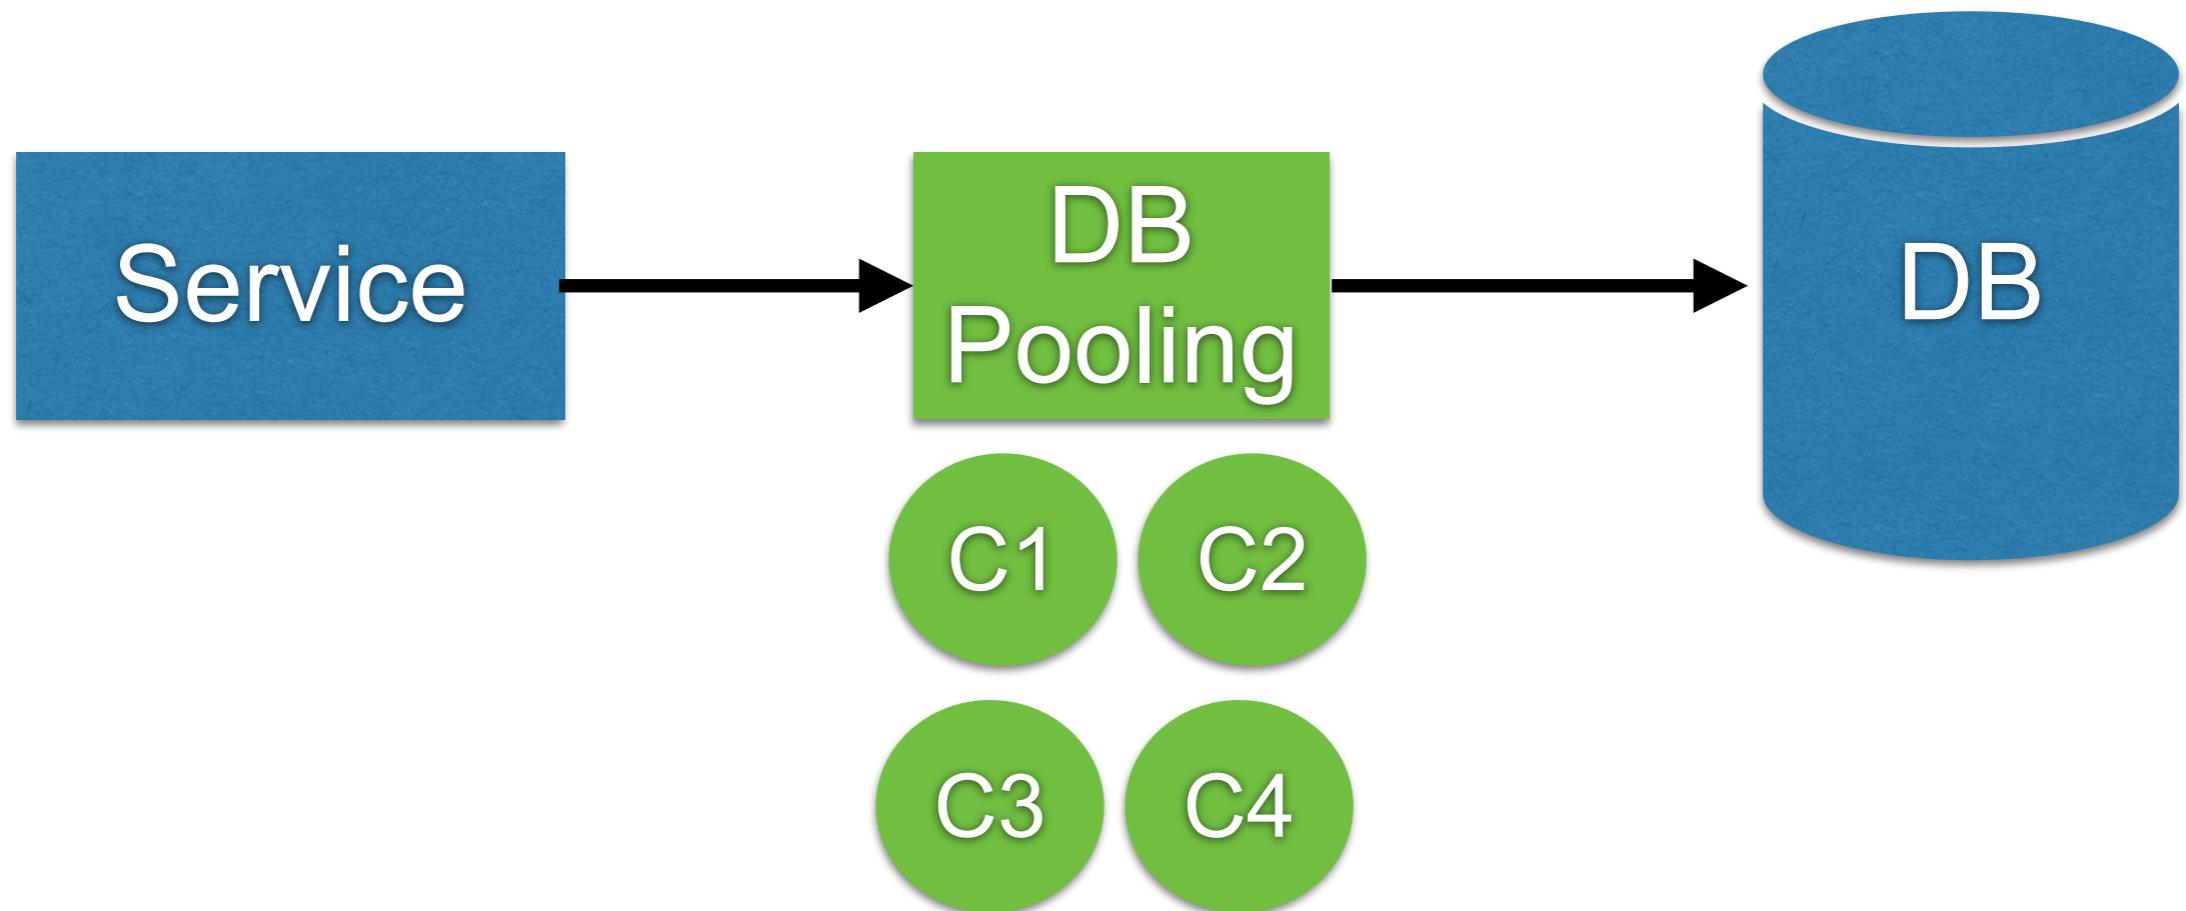


Database modeling  
Database architecture



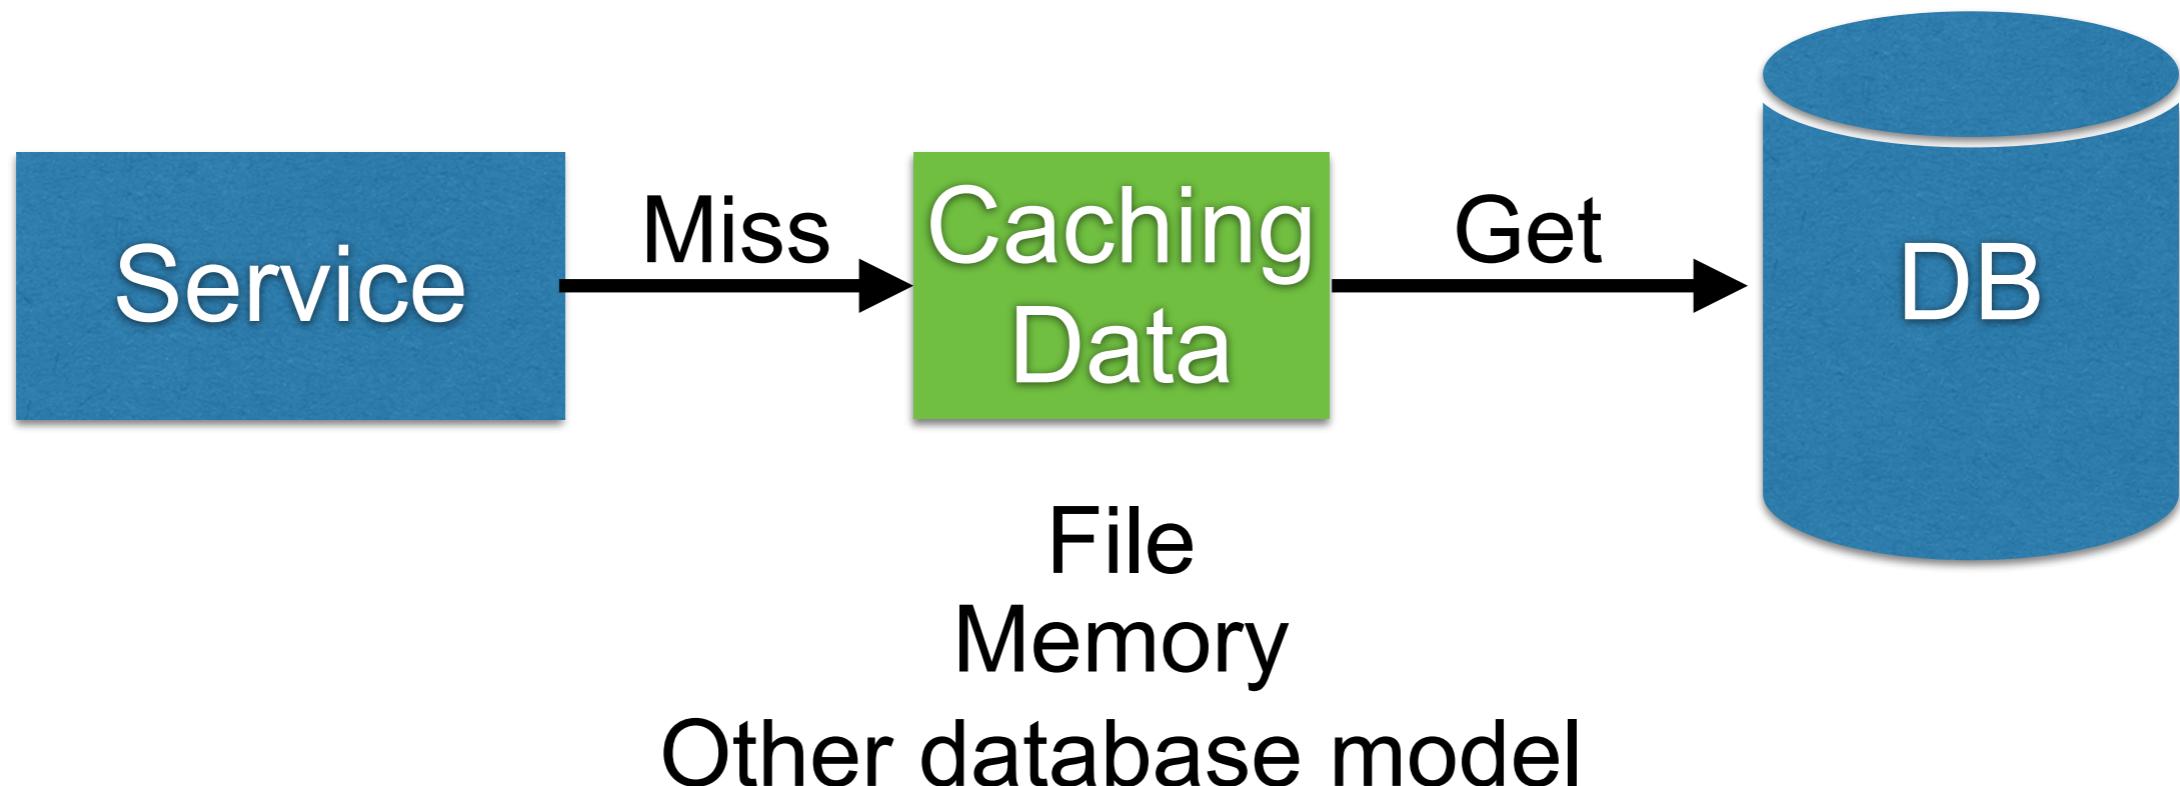
# Working with Database

## Connection pool



# Working with Database

## Caching data layer



# Database ?

Structured data collection  
Records  
Relationships



# Database Management System ?

Software systems for storing and querying databases



# User's perspective

Modeling data  
Querying data  
Store data

Read

Write



# Database Model



# Database Model ?

419 systems in ranking, July 2023

| Rank     |          |          | DBMS  | Database Model   | Score    |          |          |
|----------|----------|----------|---|--|----------|----------|----------|
| Jul 2023 | Jun 2023 | Jul 2022 |   |  | Jul 2023 | Jun 2023 | Jul 2022 |
| 1.       | 1.       | 1.       | Oracle                 | Relational, Multi-model       | 1256.01  | +24.54   | -24.28   |
| 2.       | 2.       | 2.       | MySQL                  | Relational, Multi-model       | 1150.35  | -13.59   | -44.53   |
| 3.       | 3.       | 3.       | Microsoft SQL Server  | Relational, Multi-model       | 921.60   | -8.47    | -20.53   |
| 4.       | 4.       | 4.       | PostgreSQL             | Relational, Multi-model       | 617.83   | +5.01    | +1.96    |
| 5.       | 5.       | 5.       | MongoDB                | Document, Multi-model         | 435.49   | +10.13   | -37.49   |
| 6.       | 6.       | 6.       | Redis                  | Key-value, Multi-model        | 163.76   | -3.59    | -9.86    |
| 7.       | 7.       | 7.       | IBM Db2   | Relational, Multi-model     | 139.81   | -5.07    | -21.40   |
| 8.       | 8.       | 8.       | Elasticsearch   | Search engine, Multi-model  | 139.59   | -4.16    | -14.74   |
| 9.       | 9.       | 9.       | Microsoft Access  | Relational   | 130.72   | -3.73    | -14.37   |
| 10.      | 10.      | 10.      | SQLite               | Relational   | 130.20   | -1.02    | -6.48    |
| 11.      | 11.      | ↑ 13.    | Snowflake            | Relational   | 117.69   | +3.55    | +18.53   |
| 12.      | 12.      | ↓ 11.    | Cassandra            | Wide column  | 106.53   | -2.03    | -7.88    |
| 13.      | 13.      | ↓ 12.    | MariaDB              | Relational, Multi-model     | 96.10    | -1.21    | -16.42   |
| 14.      | 14.      | 14.      | Splunk  | Search engine  | 87.12    | -2.34    | -11.09   |
| 15.      | ↑ 16.    | 15.      | Microsoft Azure SQL Database  | Relational, Multi-model     | 78.96    | -0.01    | -5.94    |

<https://db-engines.com/en/ranking>



# Database Model ?

419 systems in ranking, July 2023

| Rank     |          |          | DBMS  | Database Model   | Score    |          |          |
|----------|----------|----------|---|--|----------|----------|----------|
| Jul 2023 | Jun 2023 | Jul 2022 |   |  | Jul 2023 | Jun 2023 | Jul 2022 |
| 1.       | 1.       | 1.       | Oracle                 | Relational, Multi-model       | 1256.01  | +24.54   | -24.28   |
| 2.       | 2.       | 2.       | MySQL                  | Relational, Multi-model       | 1150.35  | -13.59   | -44.53   |
| 3.       | 3.       | 3.       | Microsoft SQL Server  | Relational, Multi-model       | 921.60   | -8.47    | -20.53   |
| 4.       | 4.       | 4.       | PostgreSQL             | Relational, Multi-model       | 617.83   | +5.01    | +1.96    |
| 5.       | 5.       | 5.       | MongoDB                | Document, Multi-model         | 435.49   | +10.13   | -37.49   |
| 6.       | 6.       | 6.       | Redis                  | Key-value, Multi-model        | 163.76   | -3.59    | -9.86    |
| 7.       | 7.       | 7.       | IBM Db2   | Relational, Multi-model      | 139.81   | -5.07    | -21.40   |
| 8.       | 8.       | 8.       | Elasticsearch   | Search engine, Multi-model  | 139.59   | -4.16    | -14.74   |
| 9.       | 9.       | 9.       | Microsoft Access  | Relational   | 130.72   | -3.73    | -14.37   |
| 10.      | 10.      | 10.      | SQLite               | Relational   | 130.20   | -1.02    | -6.48    |
| 11.      | 11.      | 13.      | Snowflake            | Relational   | 117.69   | +3.55    | +18.53   |
| 12.      | 12.      | 11.      | Cassandra            | Wide column  | 106.53   | -2.03    | -7.88    |
| 13.      | 13.      | 12.      | MariaDB              | Relational, Multi-model     | 96.10    | -1.21    | -16.42   |
| 14.      | 14.      | 14.      | Splunk  | Search engine  | 87.12    | -2.34    | -11.09   |
| 15.      | 16.      | 15.      | Microsoft Azure SQL Database  | Relational, Multi-model     | 78.96    | -0.01    | -5.94    |

<https://db-engines.com/en/ranking>



# Database Model

Relational  
Document  
Key-value  
Search engine  
Wide column  
Graph  
Time series  
Vector



# Database Model

**Relational**

Document  
Key-value

Search engine

Wide column

Graph

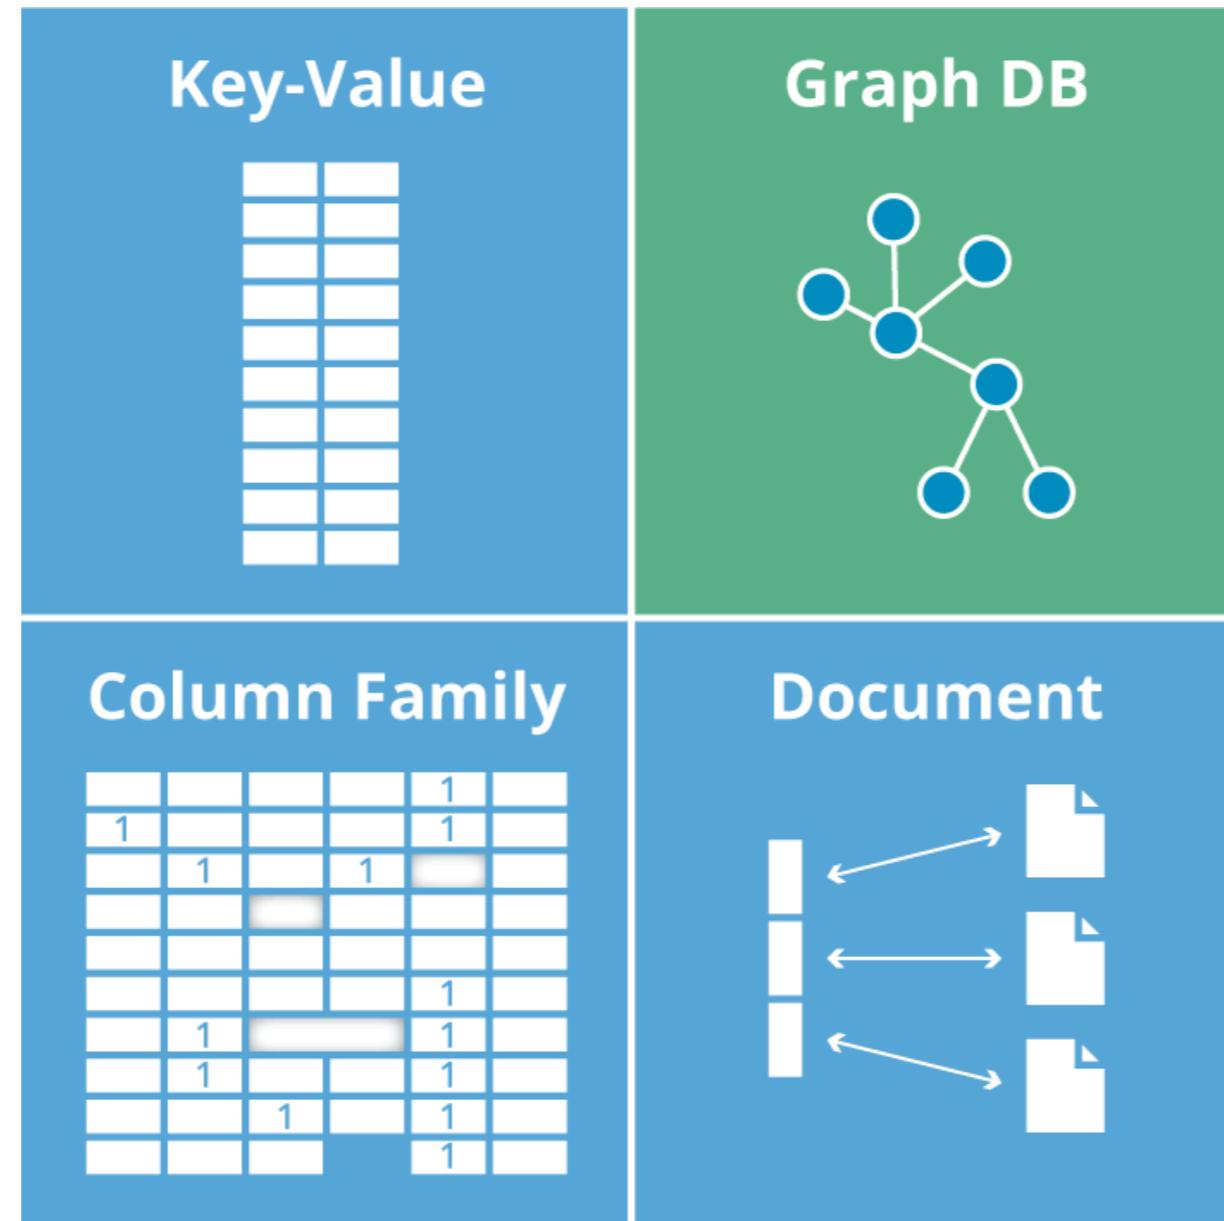
Time series

Content

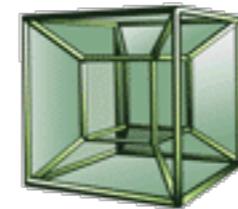
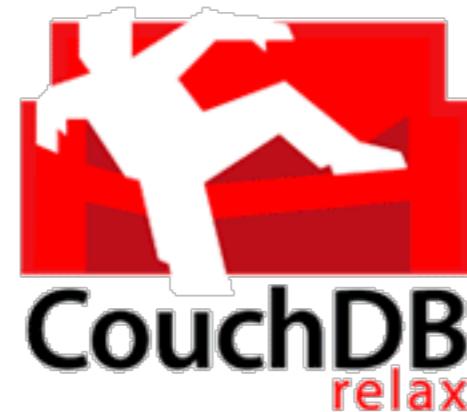
Non-Relational  
NoSQL



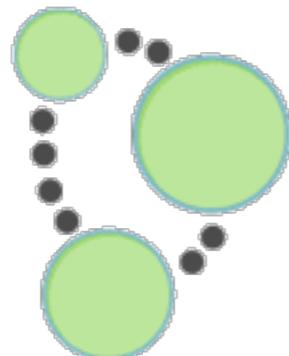
# NoSQL



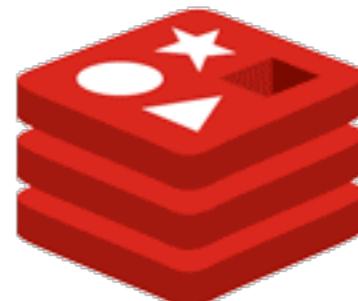
# NoSQL



HYPERTABLE<sup>INC</sup>



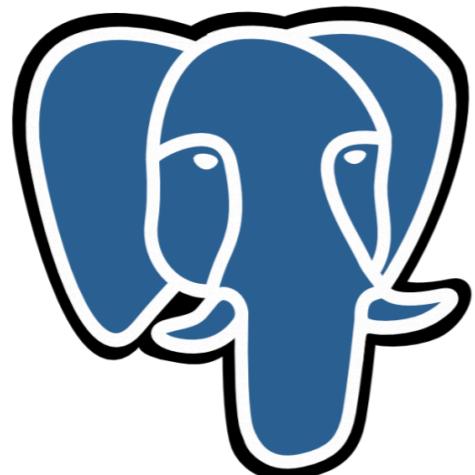
Neo4j



redis

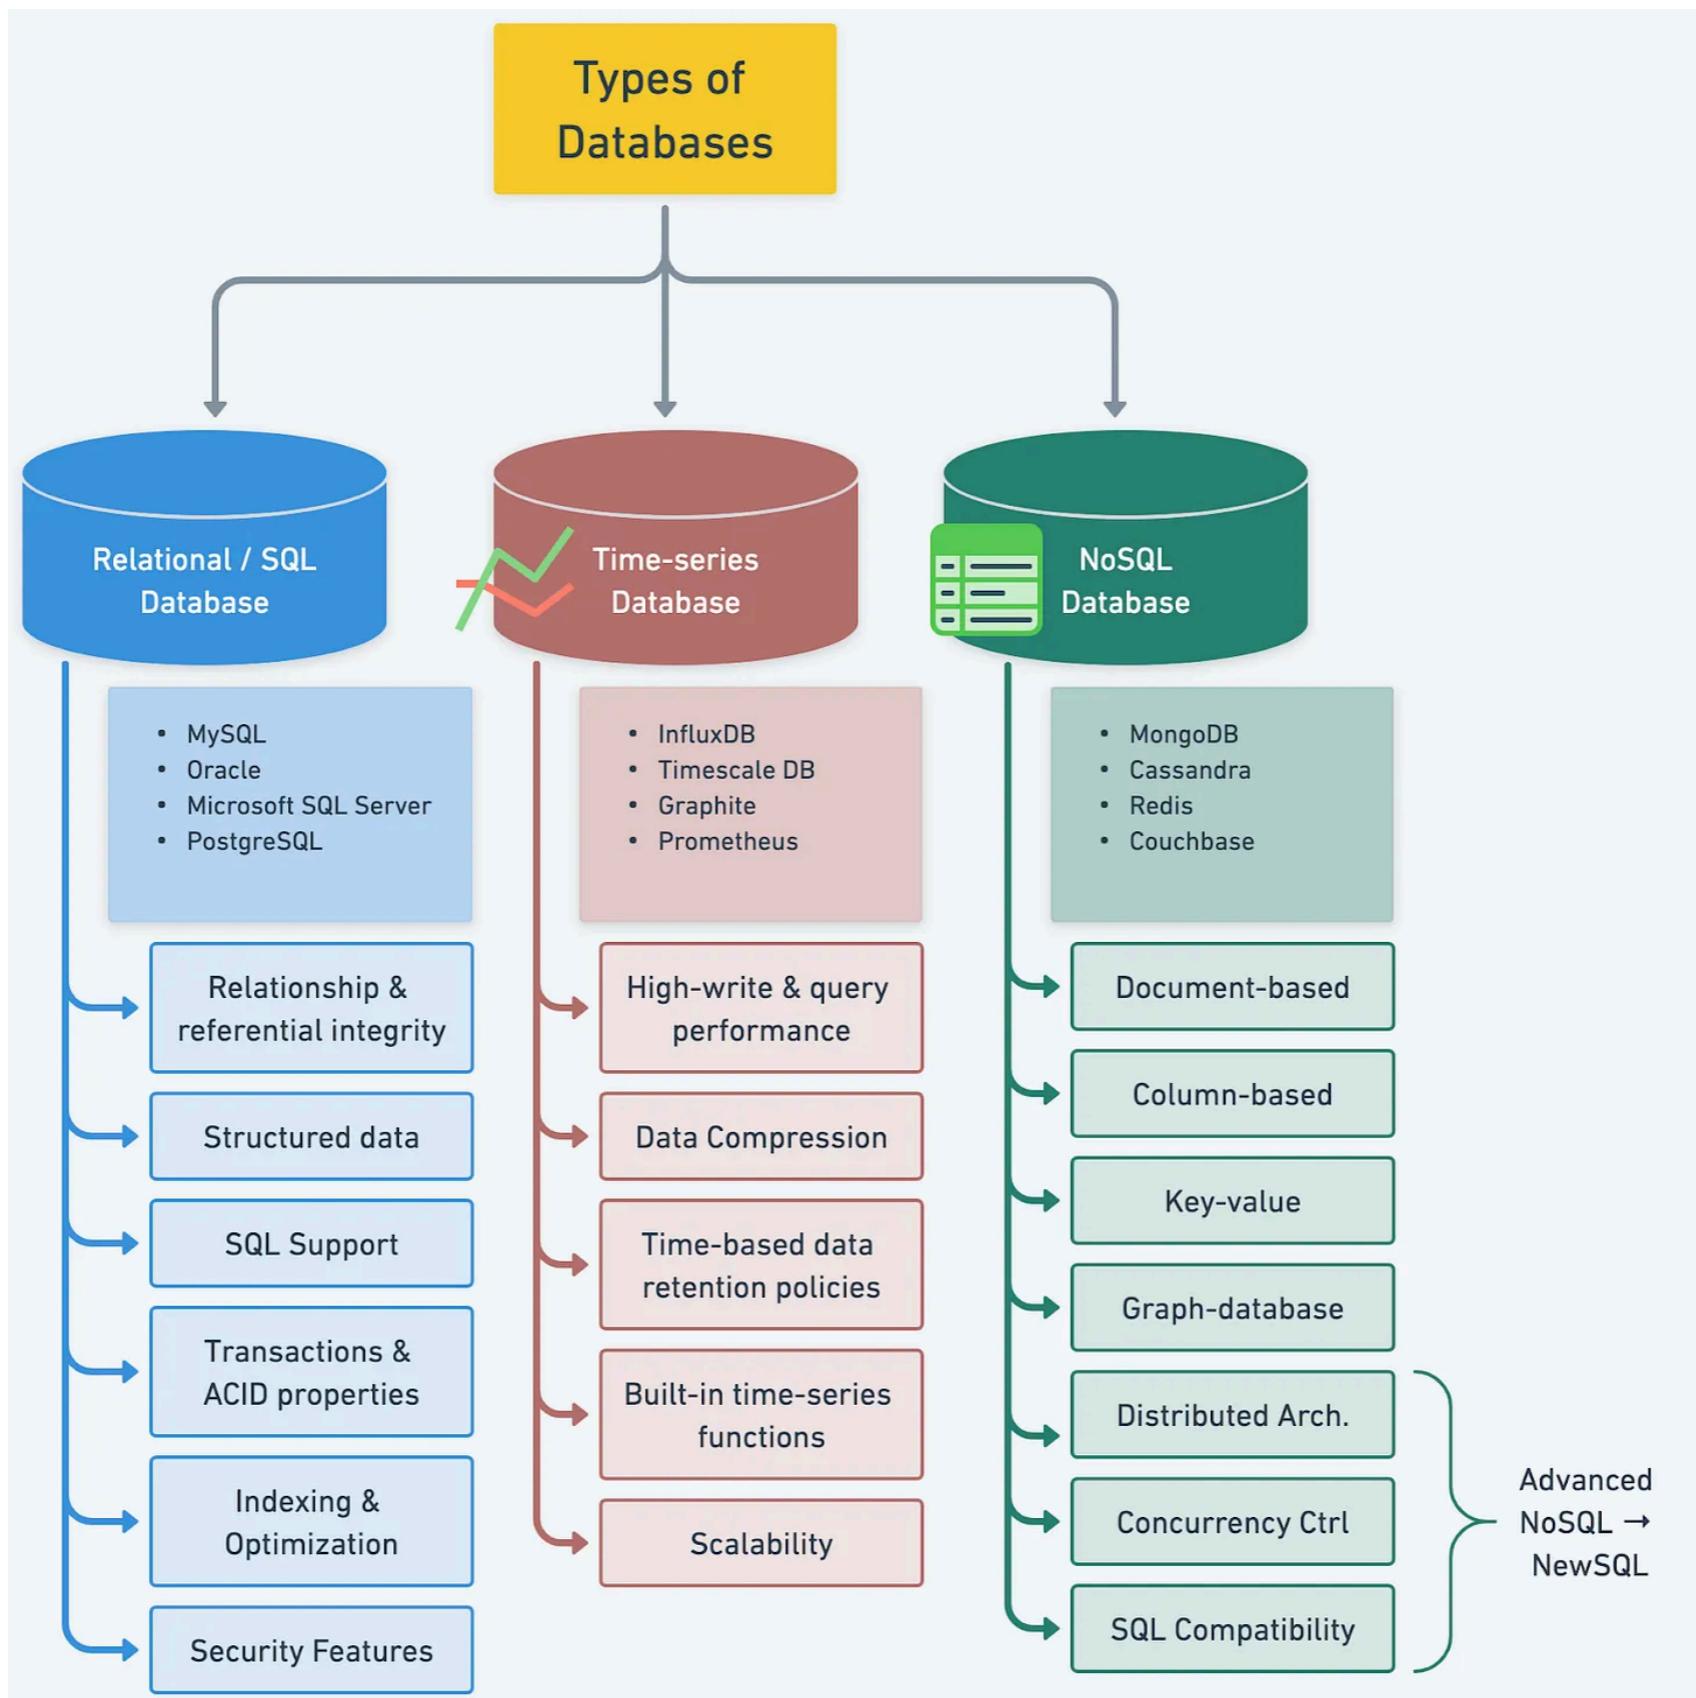


# Relational Database



PostgreSQL





<https://blog.bytebytogo.com/p/understanding-database-types>



# **Relational Database ?**

**Relational DataBase Manament System**

Maintain data in **tables**

**Pre-define structure**

Each table has **Primary key**

**Many columns**

Data is represented in term of **row**



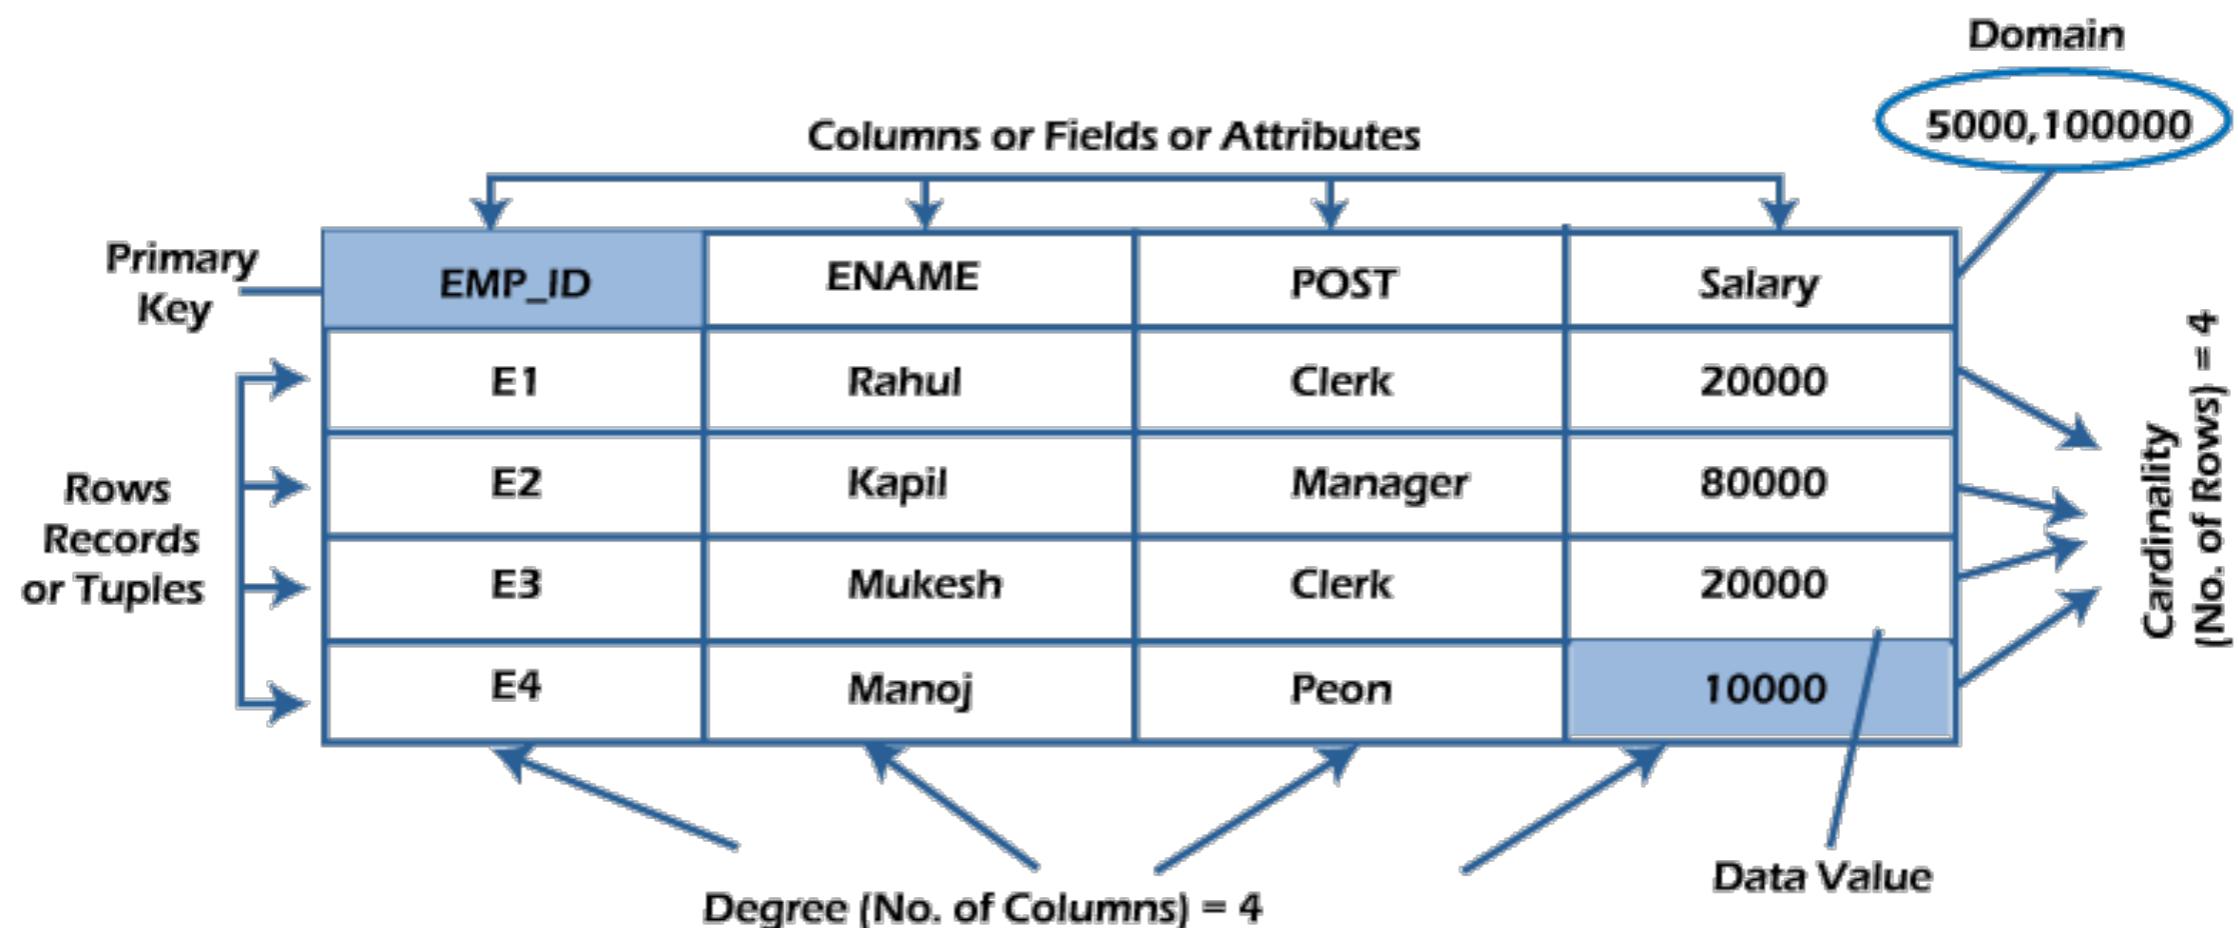
# Relational Database

| Database     | Ease of use | Support complex query | Community support | Licence     |
|--------------|-------------|-----------------------|-------------------|-------------|
| MySQL        | High        | High                  | Strong            | Open source |
| PostgreSQL   | Medium      | Very high             | Strong            | Open source |
| MSSQL Server | Medium      | Very high             | Strong            | Proprietary |
| Oracle       | Low         | Very high             | Moderate          | Proprietary |

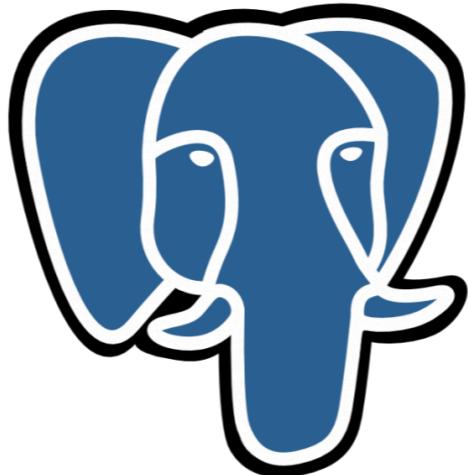


# Relational Database ?

Table = Employee



# Working with PostgreSQL



Postgre<sup>SQL</sup>



# PostgreSQL Database

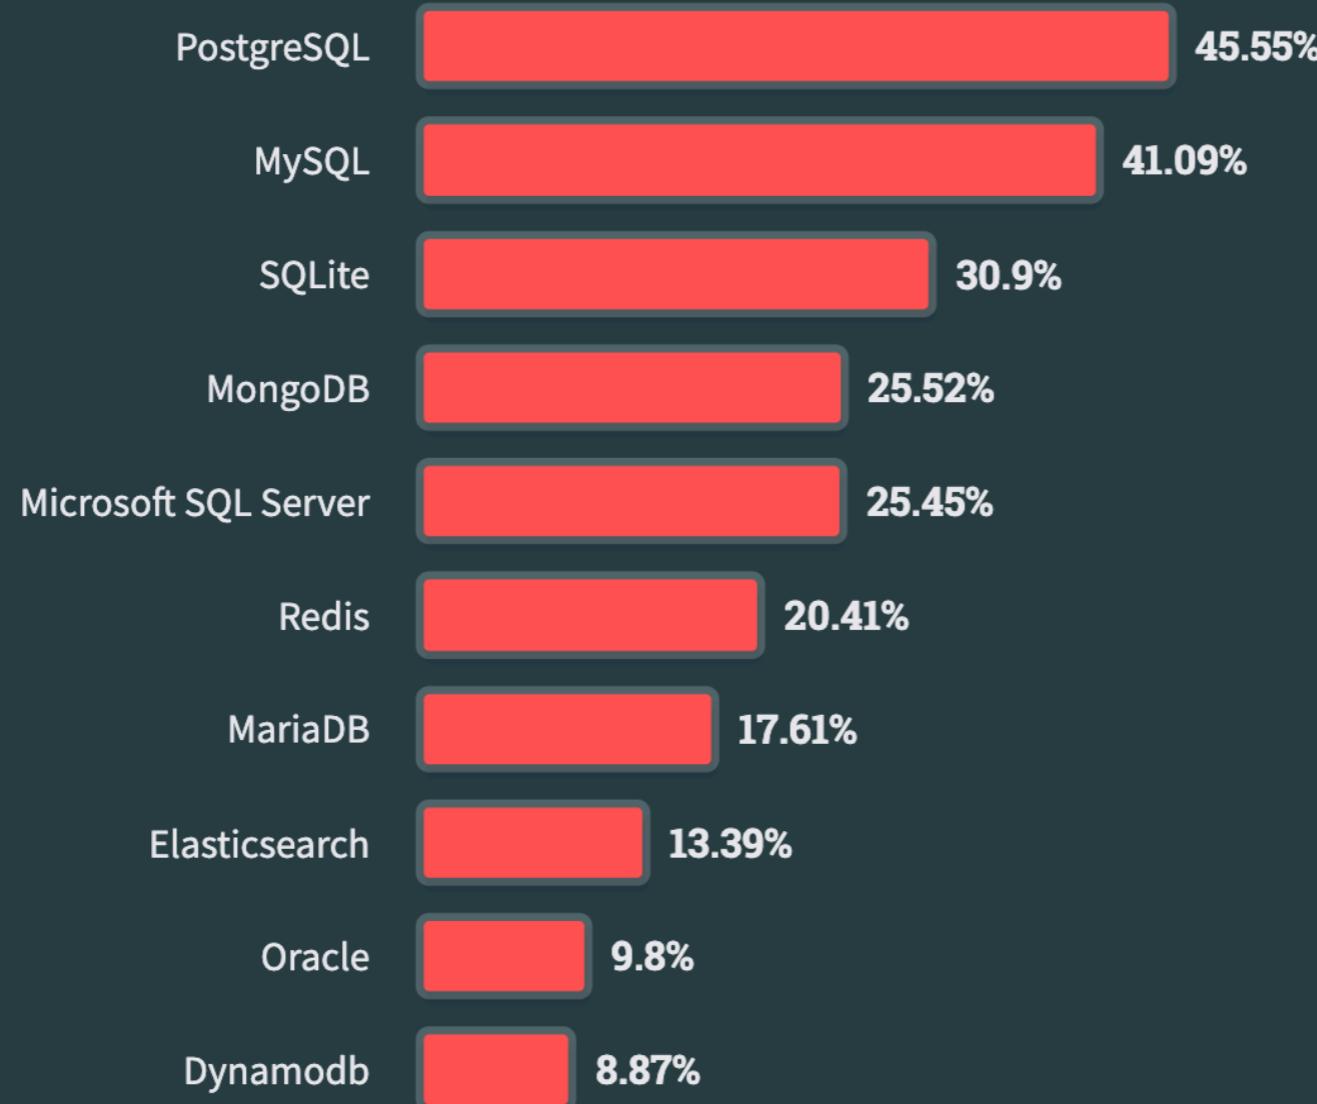
All Respondents

Professional Developers

Learning to Code

Other Coders

76,634 responses



<https://survey.stackoverflow.co/2023/#most-popular-technologies-database>



# Data Modeling

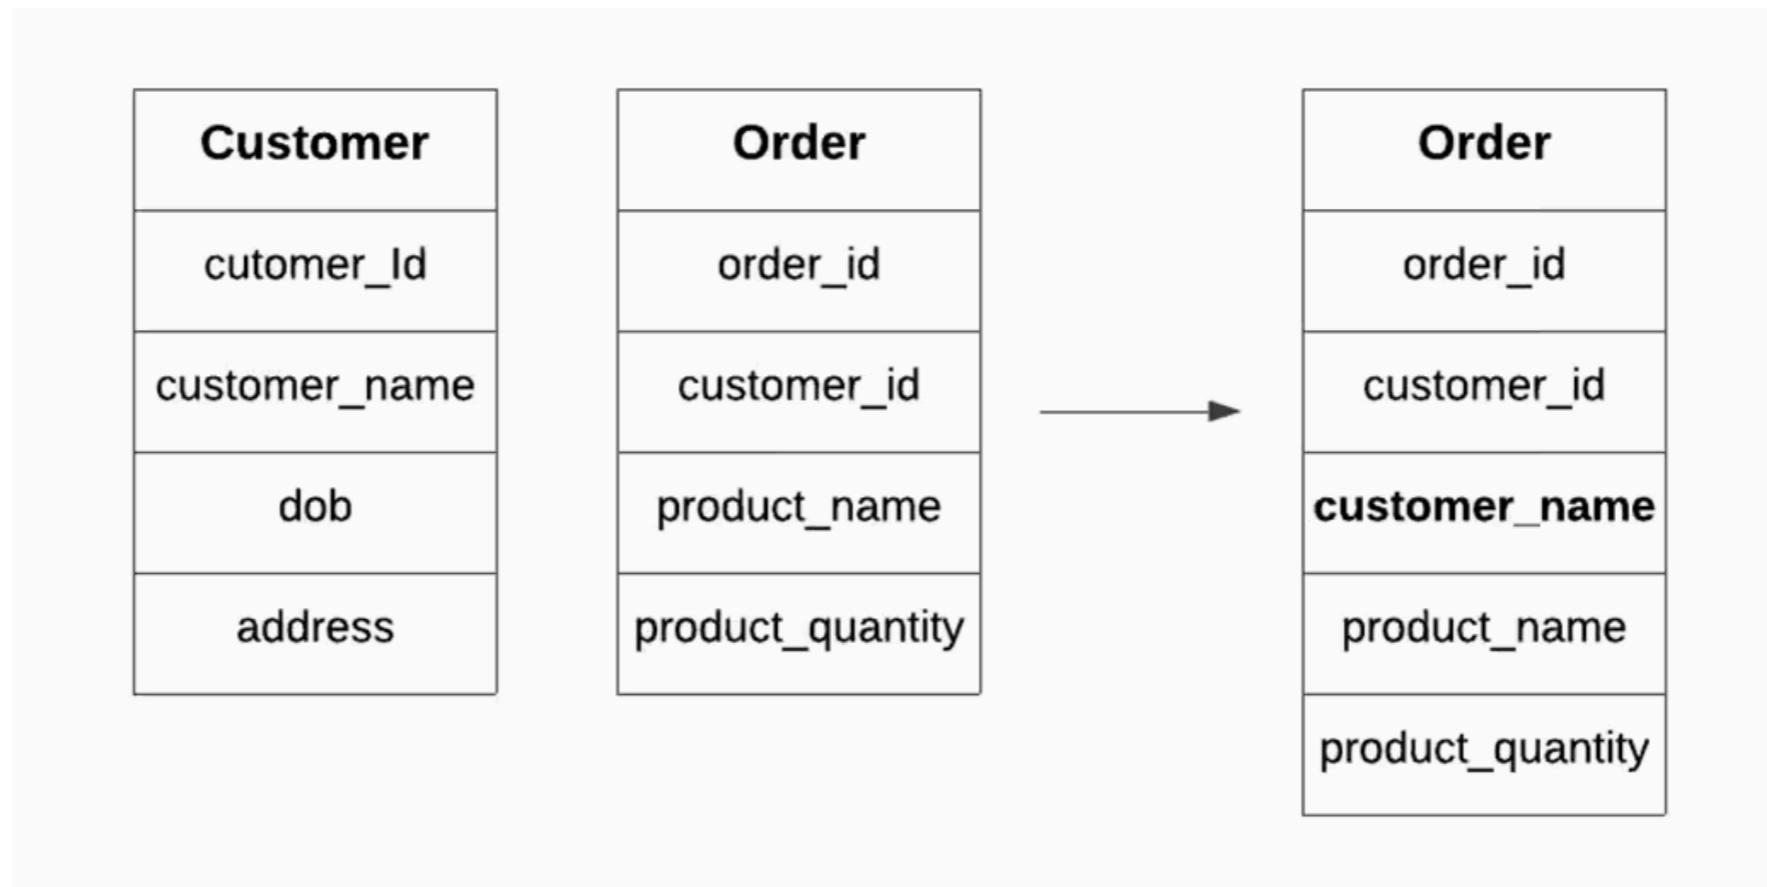
Read

Write



# Data Modeling

## Normalization Denormalization



# Normalization

Cleansing data

Reorganized data

Remove redundant data

Improve efficient for store data

*Organize data in consistent way*



# Normalization

Divide in to several tables  
Link tables with relationships

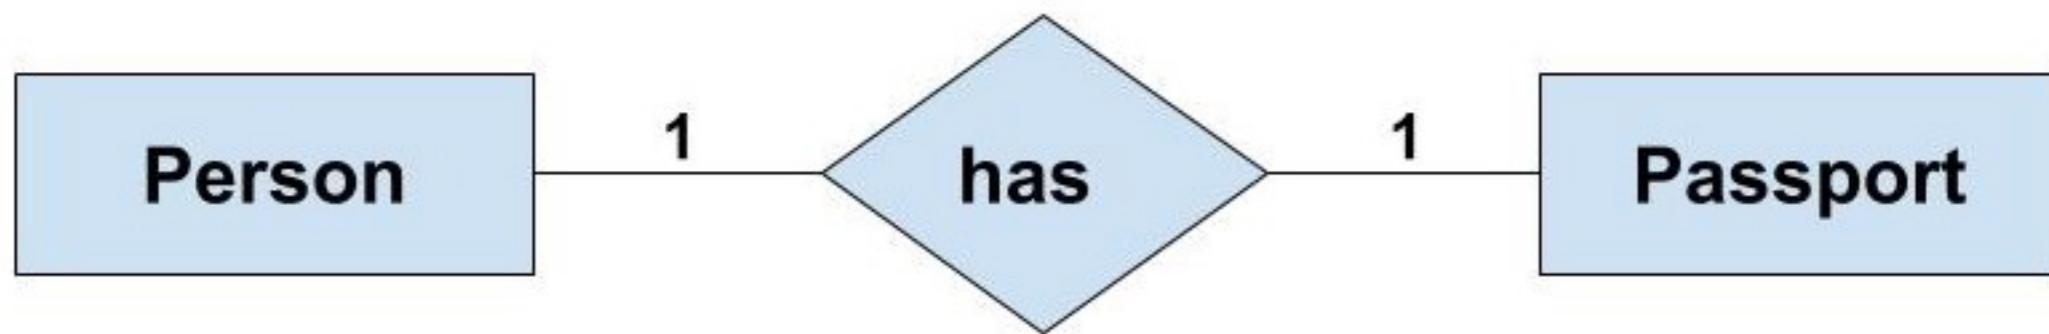


# Relationships



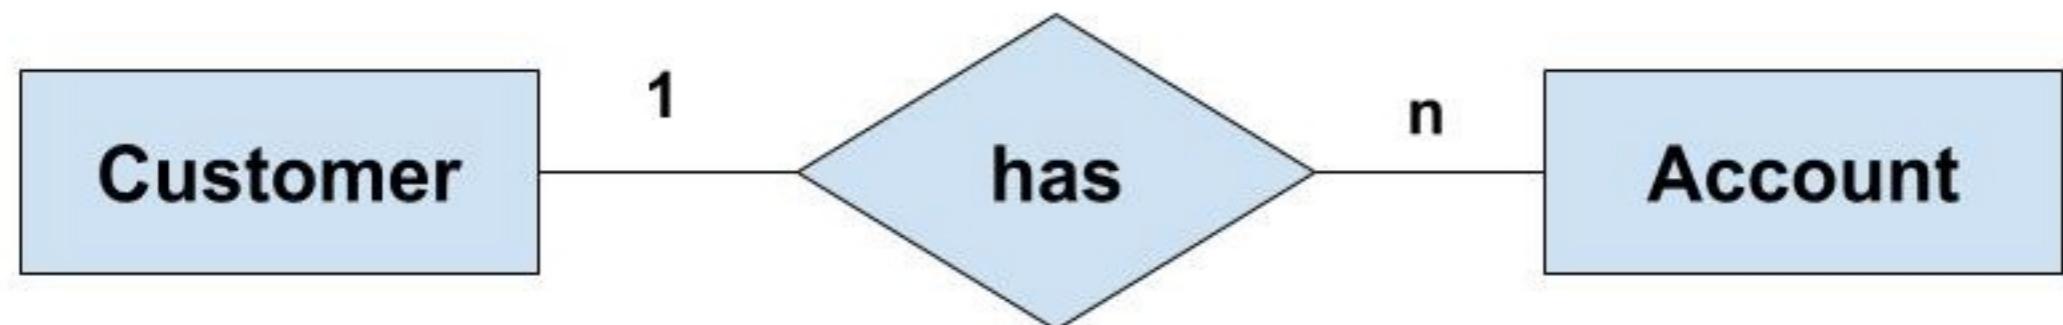
# Relationship between table

One-to-one  
One-to-many  
Many-to-many



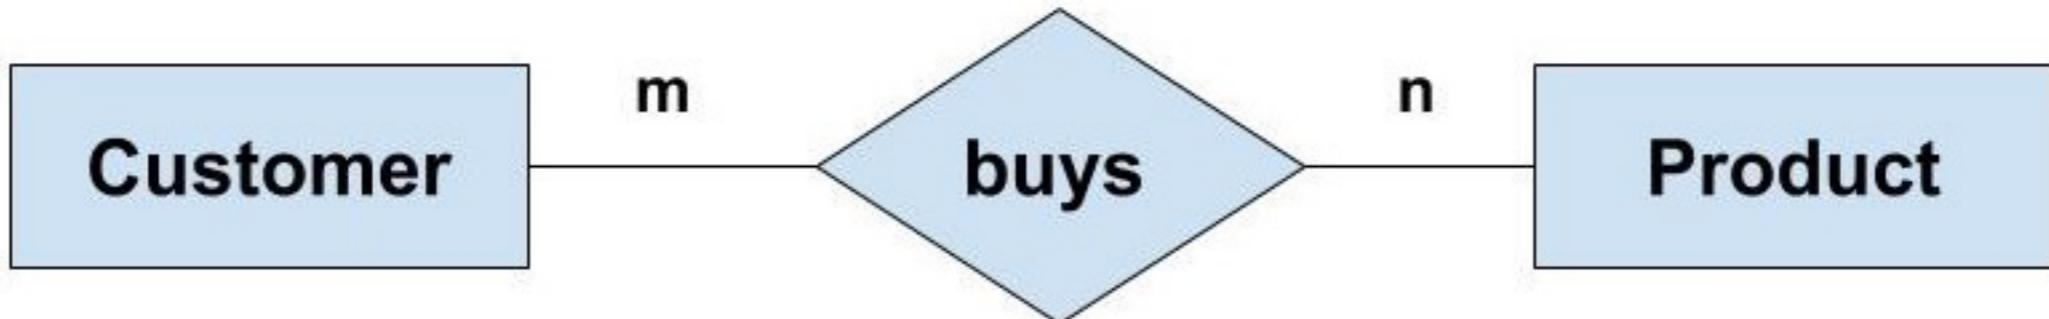
# Relationship between table

One-to-one  
One-to-many  
Many-to-many



# Relationship between table

One-to-one  
One-to-many  
**Many-to-many**

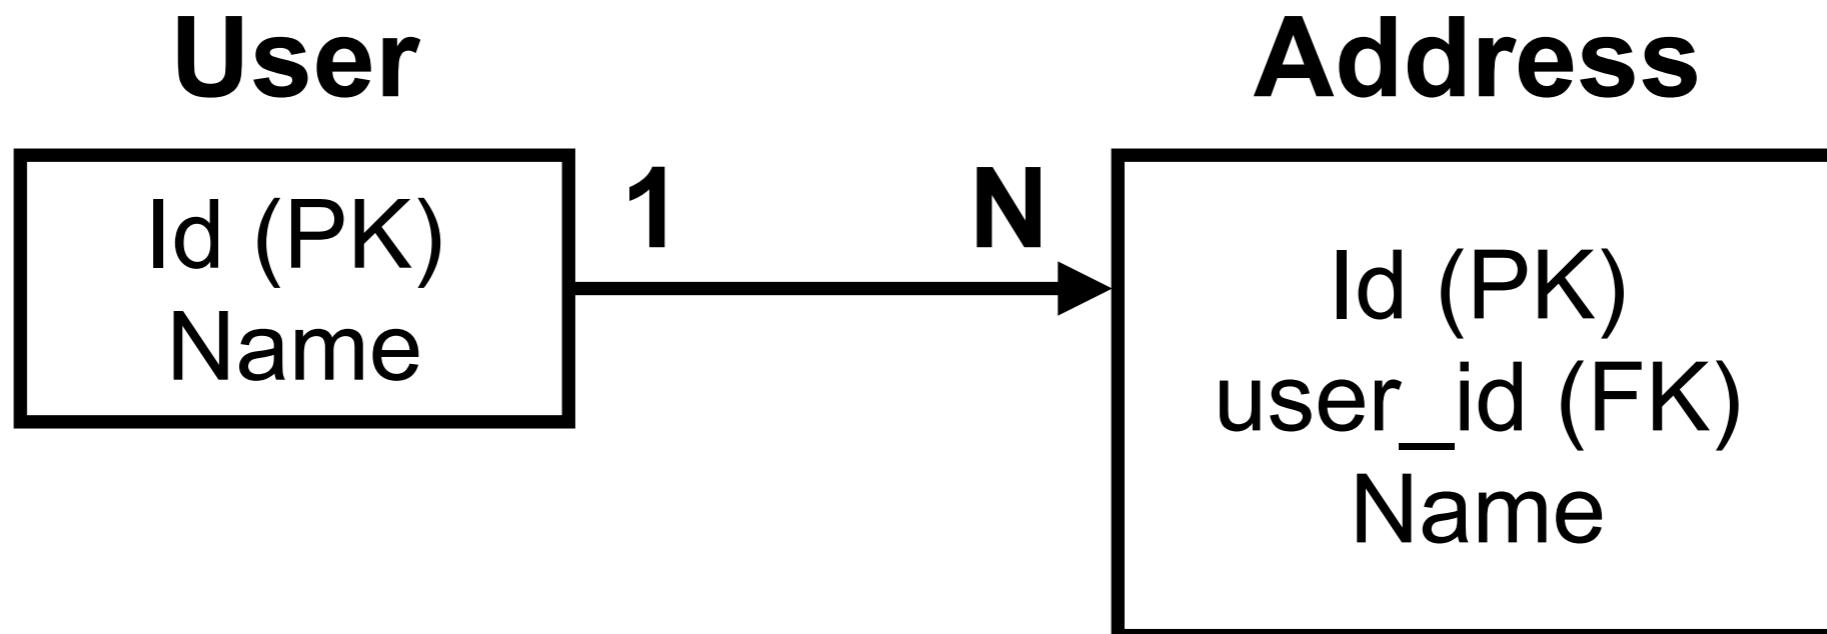


# Relationships

Primary key (PK)

Composite key

Foreign key (FK)



# Normalization forms

**First Normal form (1NF)**

**2NF**

**3NF**

**4NF**

**5NF**

**6NF**



# First Normal Form (1NF)

No repeat data in table

Every column should only have single value

Every row should be unique



# Sample data in 1NF

| Name   | Address   | Gender | T-shirt Size |
|--------|-----------|--------|--------------|
| User 1 | Address 1 | Male   | Large        |
| User 2 | Address 2 | Female | Small        |
| User 2 | Address 2 | Female | Medium       |
| User 3 | Address 3 | Male   | Medium       |



# Second Normal Form (2NF)

Apply 1NF  
Have one primary key



# Problem in 1NF

| Name   | Address   | Gender | T-shirt Size |
|--------|-----------|--------|--------------|
| User 1 | Address 1 | Male   | Large        |
| User 2 | Address 2 | Female | Small        |
| User 2 | Address 2 | Female | Medium       |
| User 3 | Address 3 | Male   | Medium       |



# Sample data in 2NF

| User ID | Name   | Address   | Gender |
|---------|--------|-----------|--------|
| 1       | User 1 | Address 1 | Male   |
| 2       | User 2 | Address 2 | Female |
| 3       | User 3 | Address 3 | Male   |



# Sample data in 2NF

| User ID | T-shirt Size |
|---------|--------------|
| 1       | Large        |
| 2       | Small        |
| 2       | Medium       |
| 3       | Medium       |



# Third Normal Form (3NF)

Apply 2NF

It should only be dependent on the primary key

if any changes to the primary key are made,  
all data that is impacted must be put into a  
new table



# Problem in 2NF

| User ID | Name   | Address   | Gender |
|---------|--------|-----------|--------|
| 1       | User 1 | Address 1 | Male   |
| 2       | User 2 | Address 2 | Female |
| 3       | User 3 | Address 3 | Male   |



# Sample data in 3NF

| User ID | Name   | Address   | Gender |
|---------|--------|-----------|--------|
| 1       | User 1 | Address 1 | 1      |
| 2       | User 2 | Address 2 | 2      |
| 3       | User 3 | Address 3 | 1      |



# Sample data in 3NF

| Gender ID | Gender name |
|-----------|-------------|
| 1         | Male        |
| 2         | Female      |



# Sample data in 3NF

| User ID | T-shirt Size |
|---------|--------------|
| 1       | Large        |
| 2       | Small        |
| 2       | Medium       |
| 3       | Medium       |



# **3NF is normal form ...**



# Benefits of data normalization

Freeing up space

Improve query response time

Reduce data inconsistency

Easier to use and analyze



# Challenges of data normalization

Slower query response rates  
Accurate knowledge is required  
Add complexity for teams  
Denormalization as an alternative !!



# Primary Key



# Foreign Key



# Constraints



# Workshop



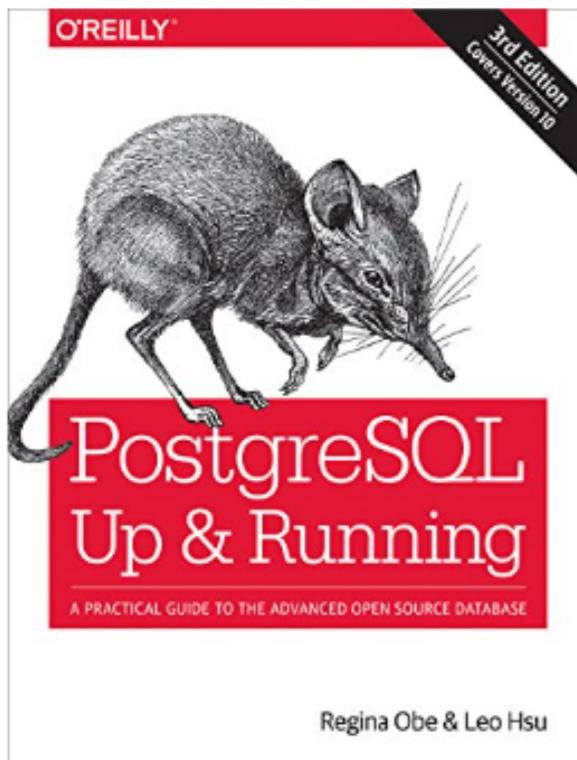
# Book Store

QuickStart  
Guides.

Master SQL and Manage, Analyze and Manipulate Data

\$25<sup>19</sup> ✓p

Kindle Store › Kindle eBooks › Computers & Technology



## PostgreSQL: Up and Running: A Practical Guide to the Advanced Open Source Database 3rd Edition, Kindle Edition

by [Regina O. Obe](#) (Author), [Leo S. Hsu](#) (Author) | Format: Kindle Edition

4.4 99 ratings

[See all formats and editions](#)

Kindle  
**\$10.09 - \$28.99**

Paperback  
\$44.99

[Read with Our Free App](#)

6 Used from \$36.00

20 New from \$34.18

Thinking of migrating to PostgreSQL? This clear, fast-paced introduction helps you understand and use this open source database system. Not only will you learn about the enterprise class features in versions 9.5 to 10, you'll also discover that PostgreSQL is more than a database system—it's an impressive application platform as well.

With examples throughout, this book shows you how to achieve tasks that are difficult or impossible in other databases. This third edition covers new features, such as ANSI-SQL constructs found only in proprietary databases until now: foreign data wrapper

ANSI-SQL constructs found only in PostgreSQL: foreign data wrapper, foreign function support, and more. The book also covers the latest PostgreSQL 10 features, including the new JSONB data type, improved query performance, and enhanced security features.

[▼ Read more](#)

[Read sample](#)

### Follow the Author



Regina Obe

[Follow](#)

ISBN-13



978-1491963418

Edition

#

3rd

Sticky notes



[On Kindle Scribe](#)

Publisher



O'Reilly Media

Publication date



October 10, 2017

Language



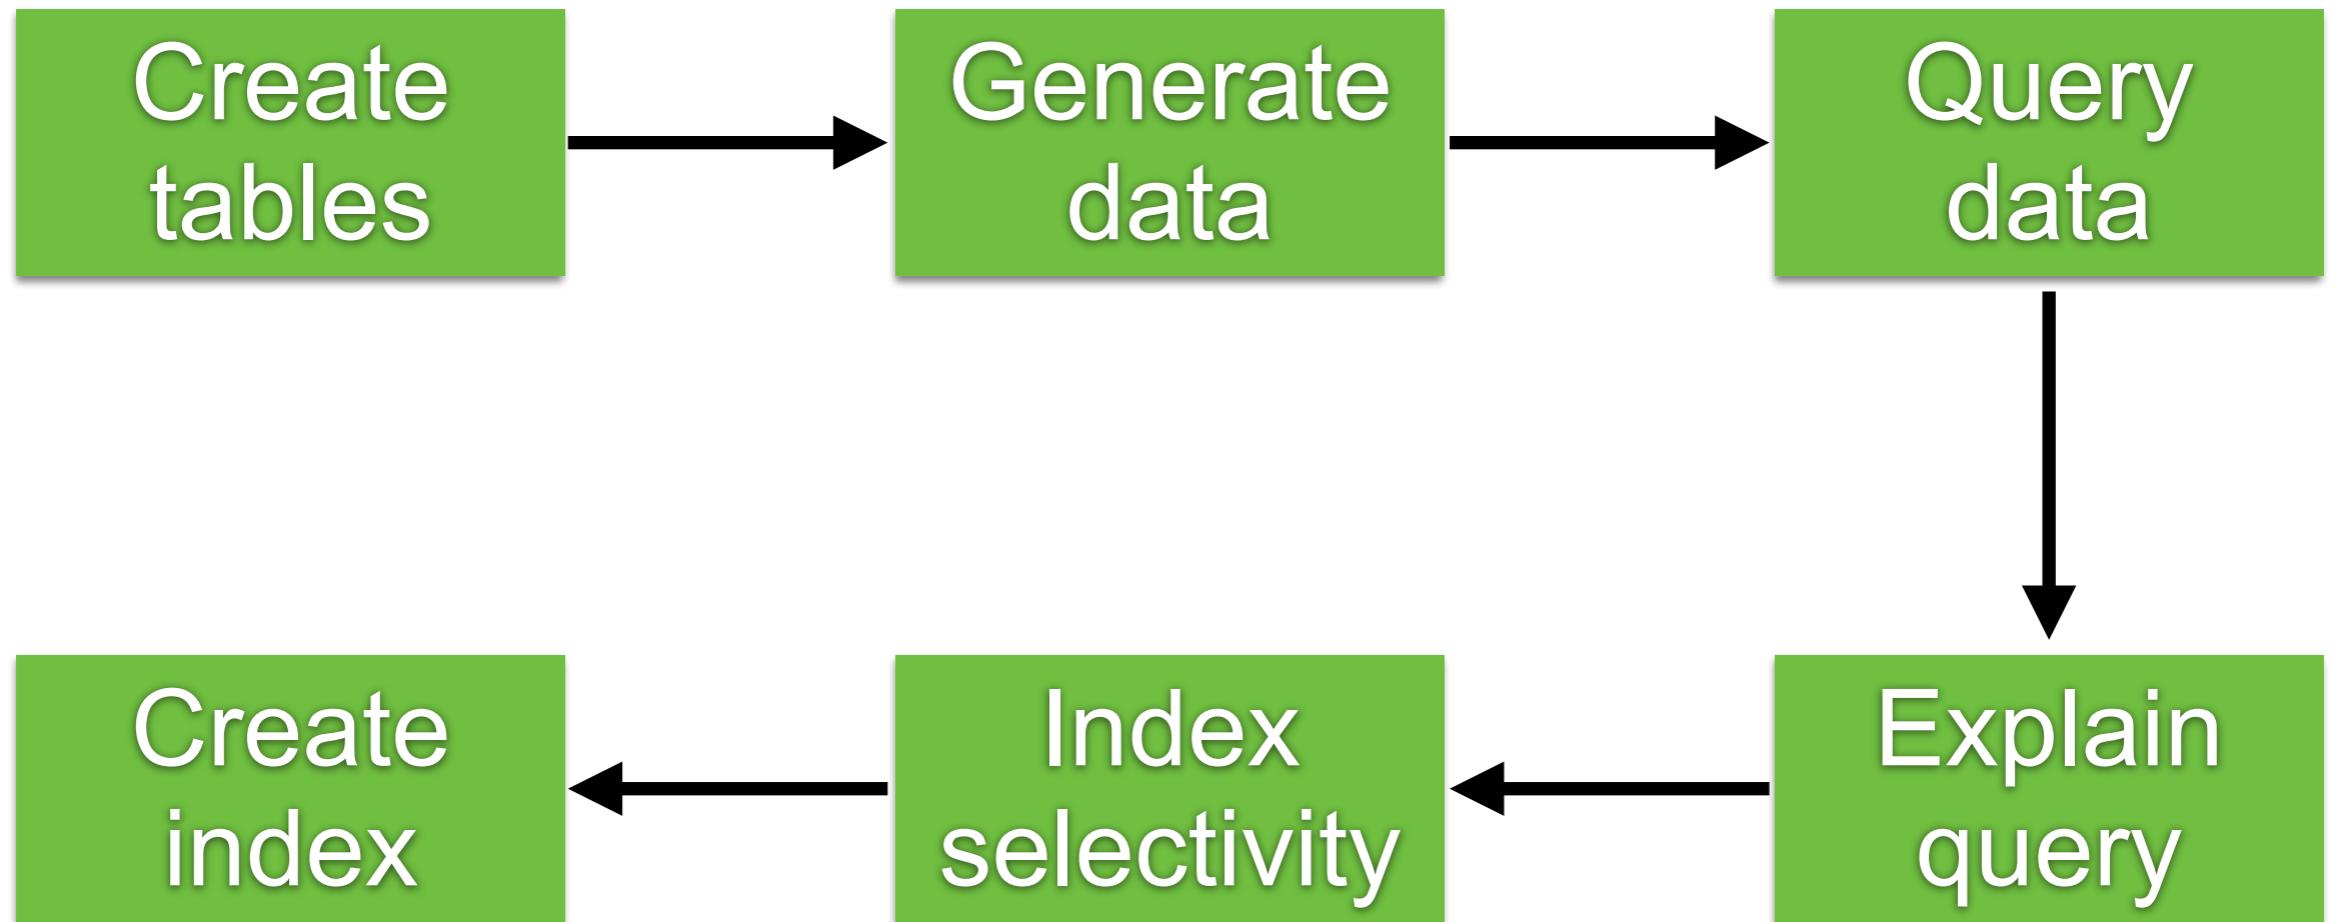
English



# Book Store



# Step in workshop

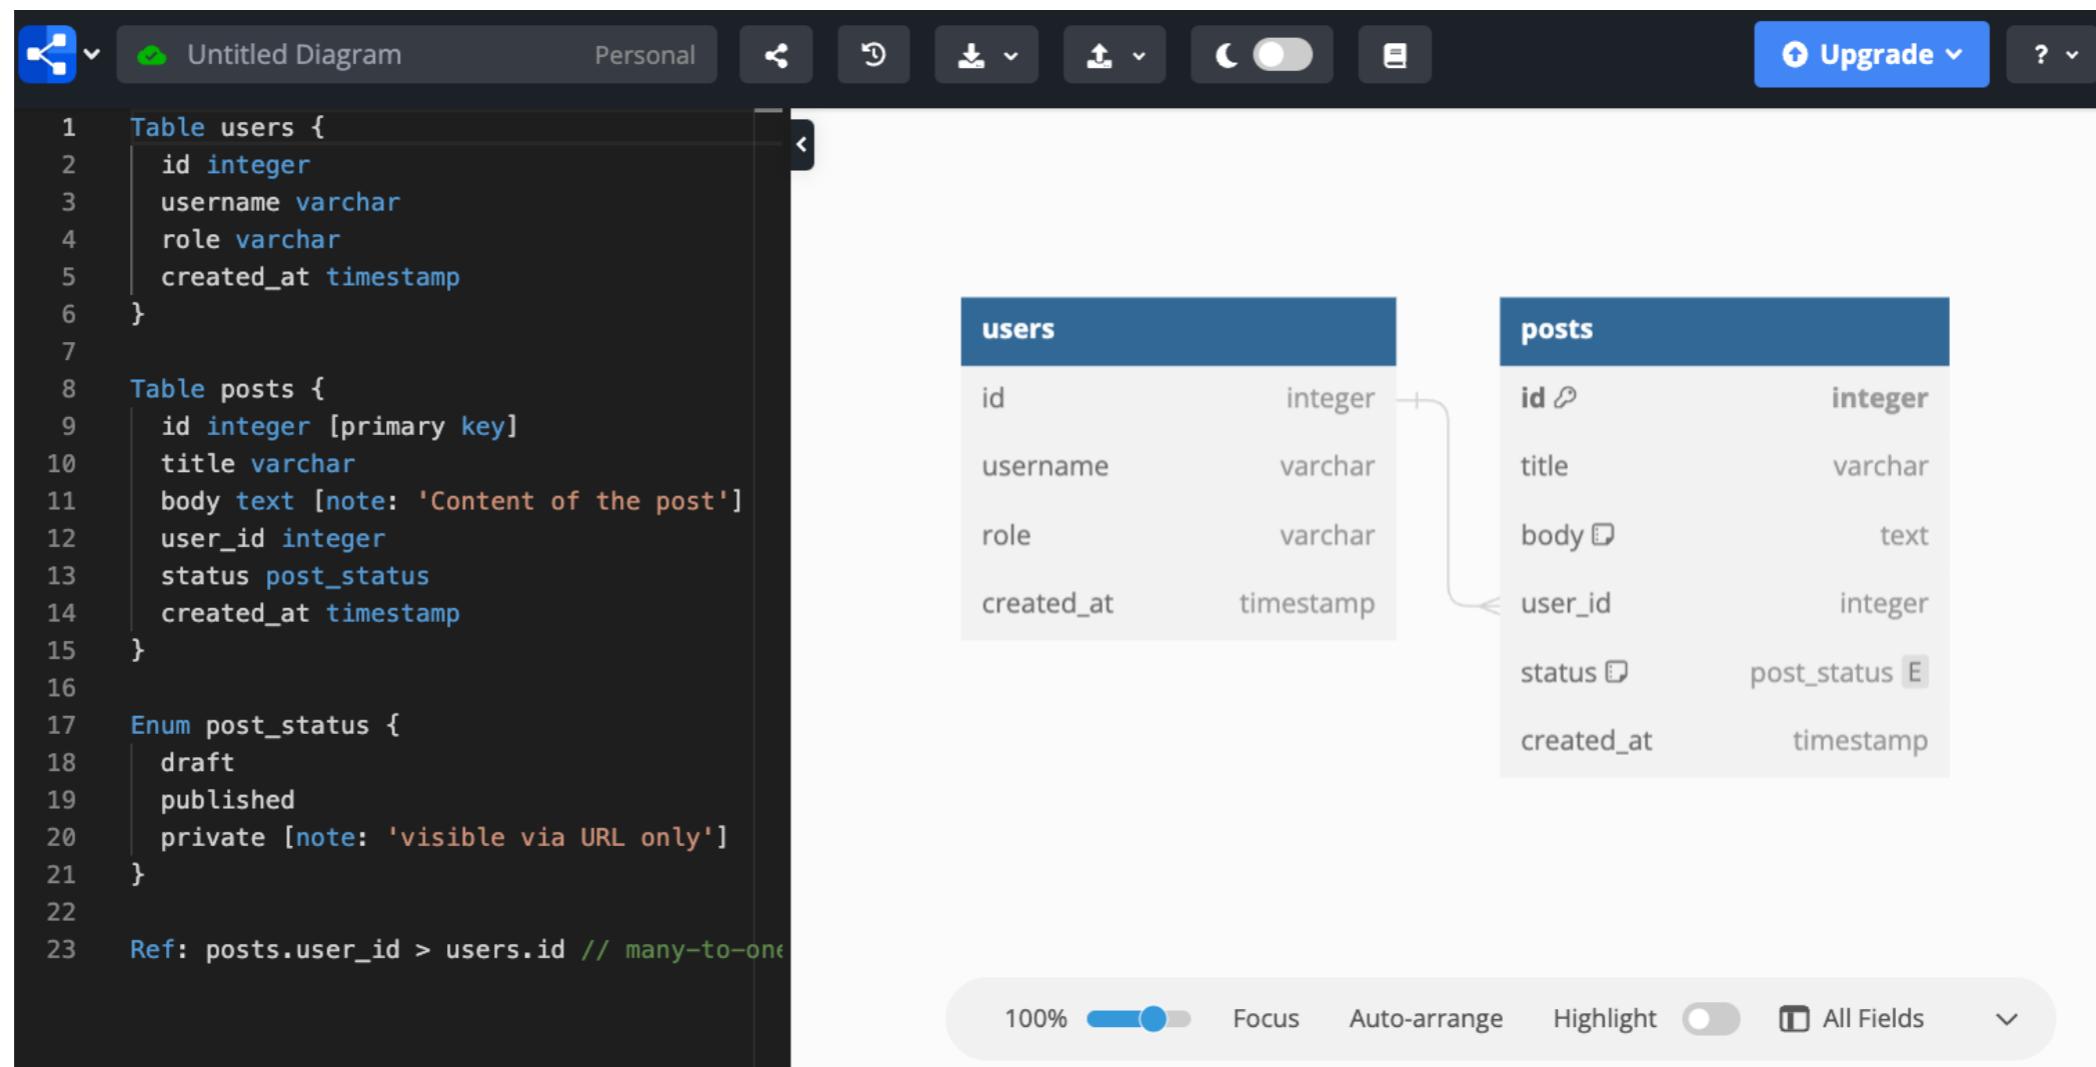


# Design Data Model



# Design table with DBML

## Database Markup Language



<https://dbml.dbdiagram.io/home/>



# Convert SQL to DBML

\$npm install -g @dbml/cli

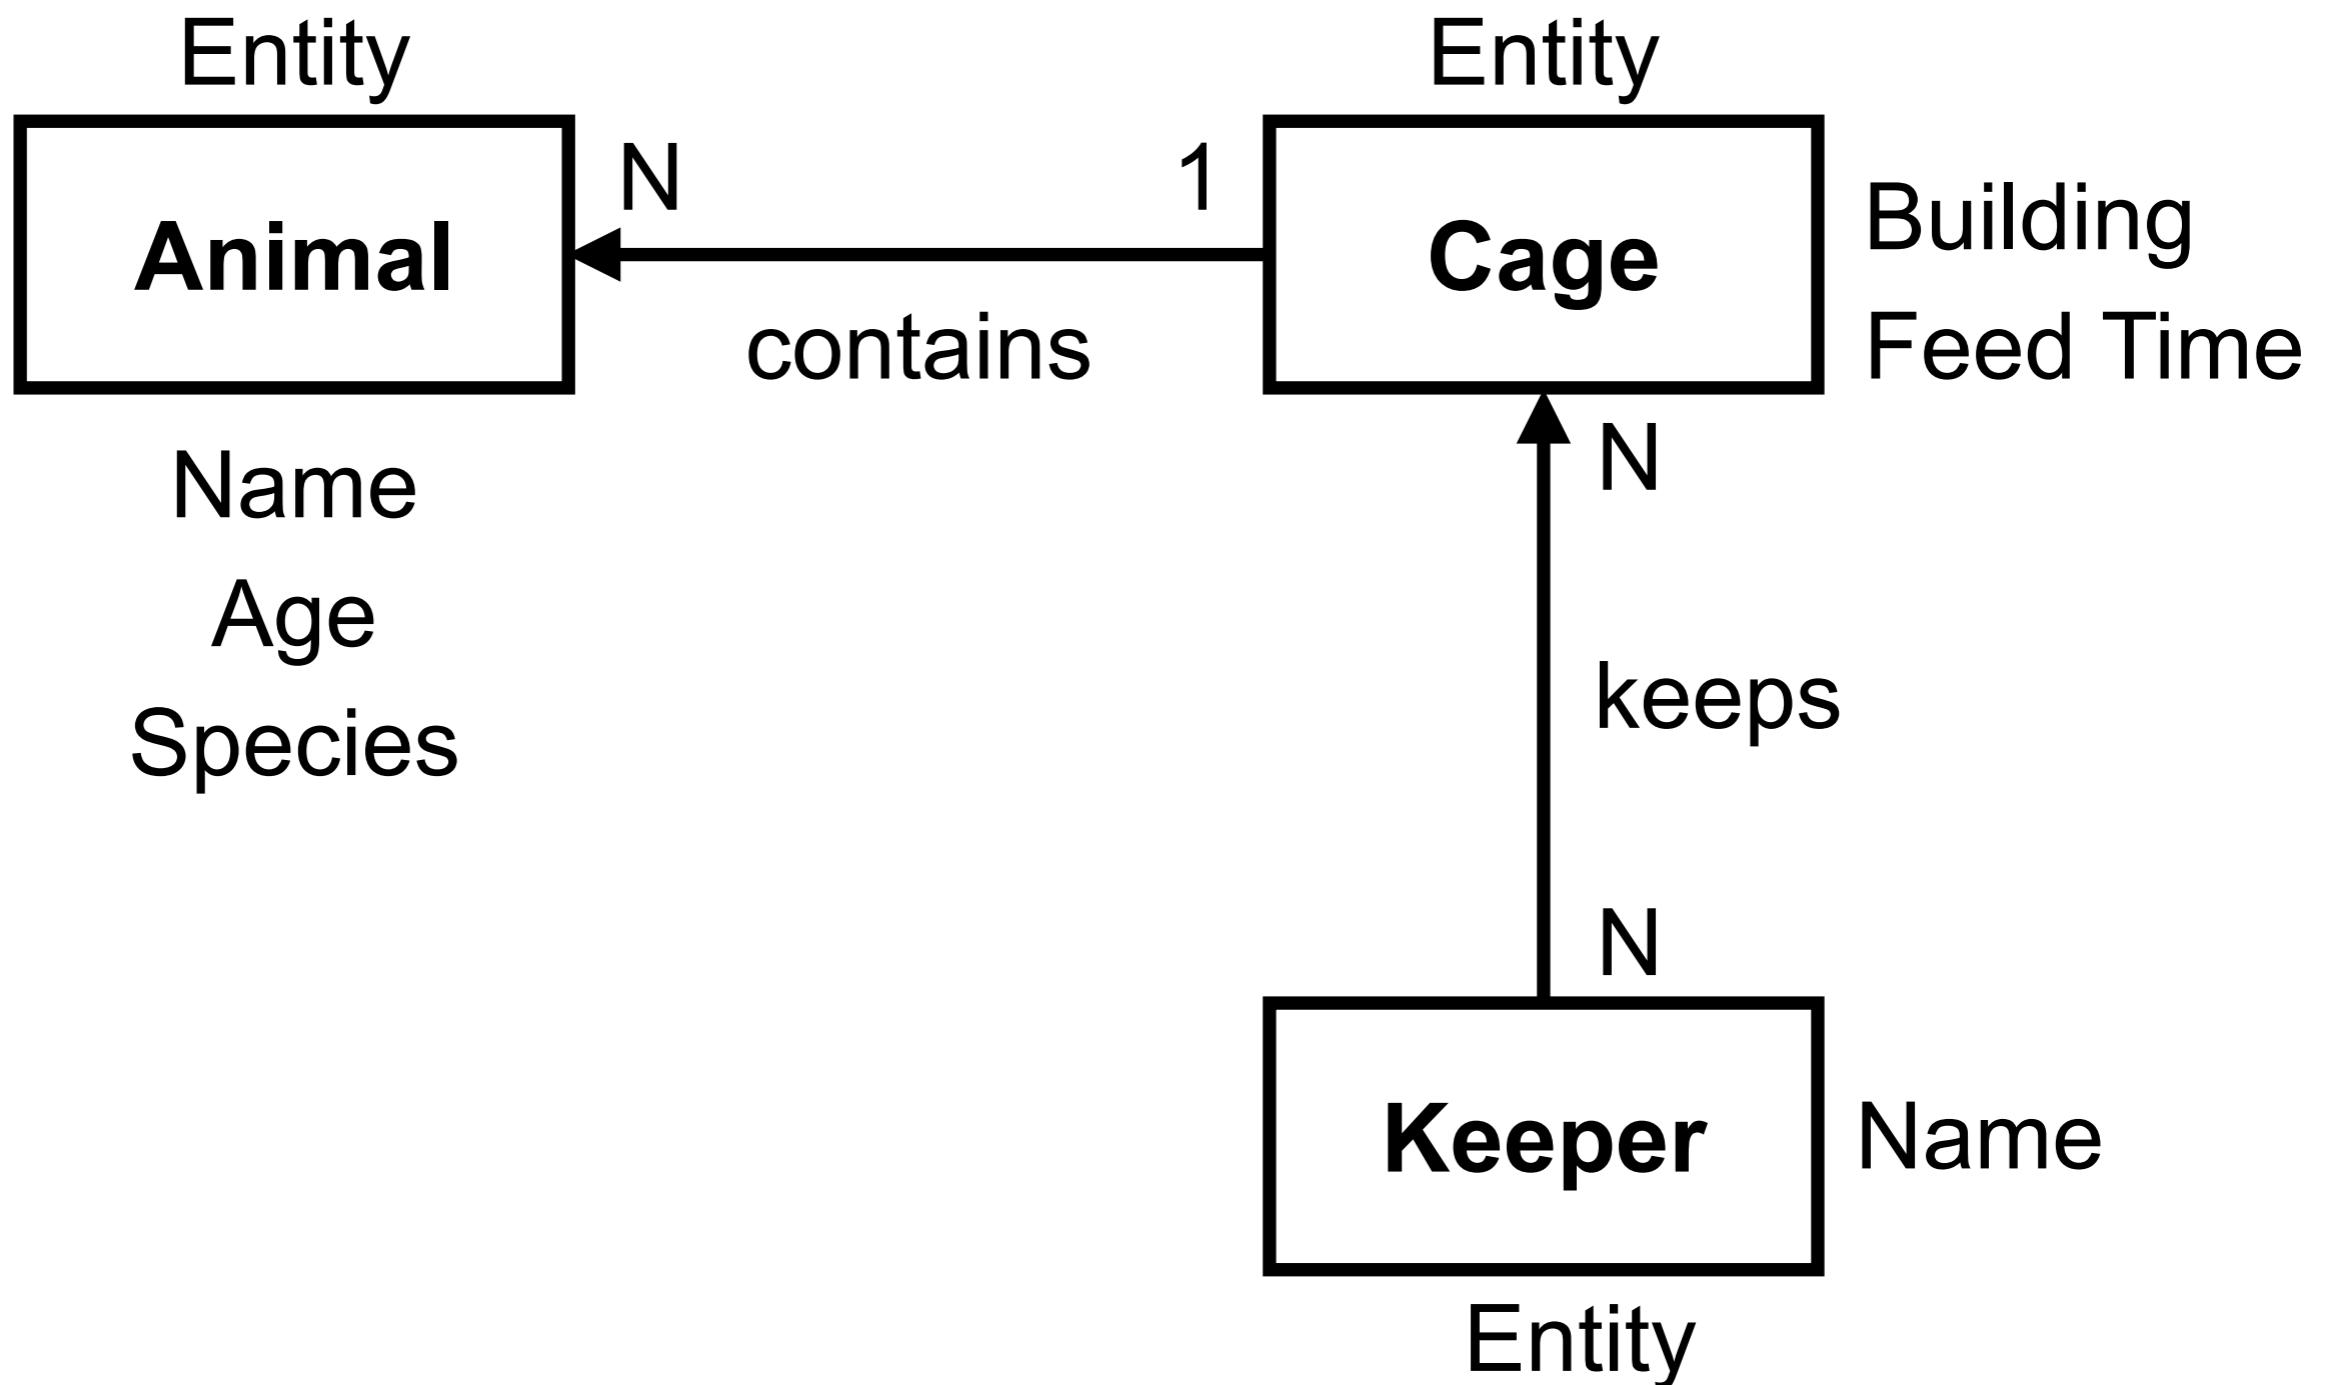


<https://dbml.dbdiagram.io/cli/#convert-a-sql-file-to-dbml>



# Design with entity relation diagram





# How to manage RDBMS ?



# **Database Language**

**Data Definition Language (DDL)**

**Data Manipulation Language (DML)**

**Data Control Language (DCL)**

**Transaction Control Language (TCL)**



# Data Definition Language (DDL)

Create  
Alter  
Drop  
**Truncate**



# Data Manipulation Language (DML)

Select  
Insert  
Update  
Delete



# Delete vs Truncate



# Delete vs Truncate

## Delete

Row-level operation

Mark every row matching the WHERE-clause as deleted

Slow operation

## Truncate

Table operation

Empty the table

Start a new data file



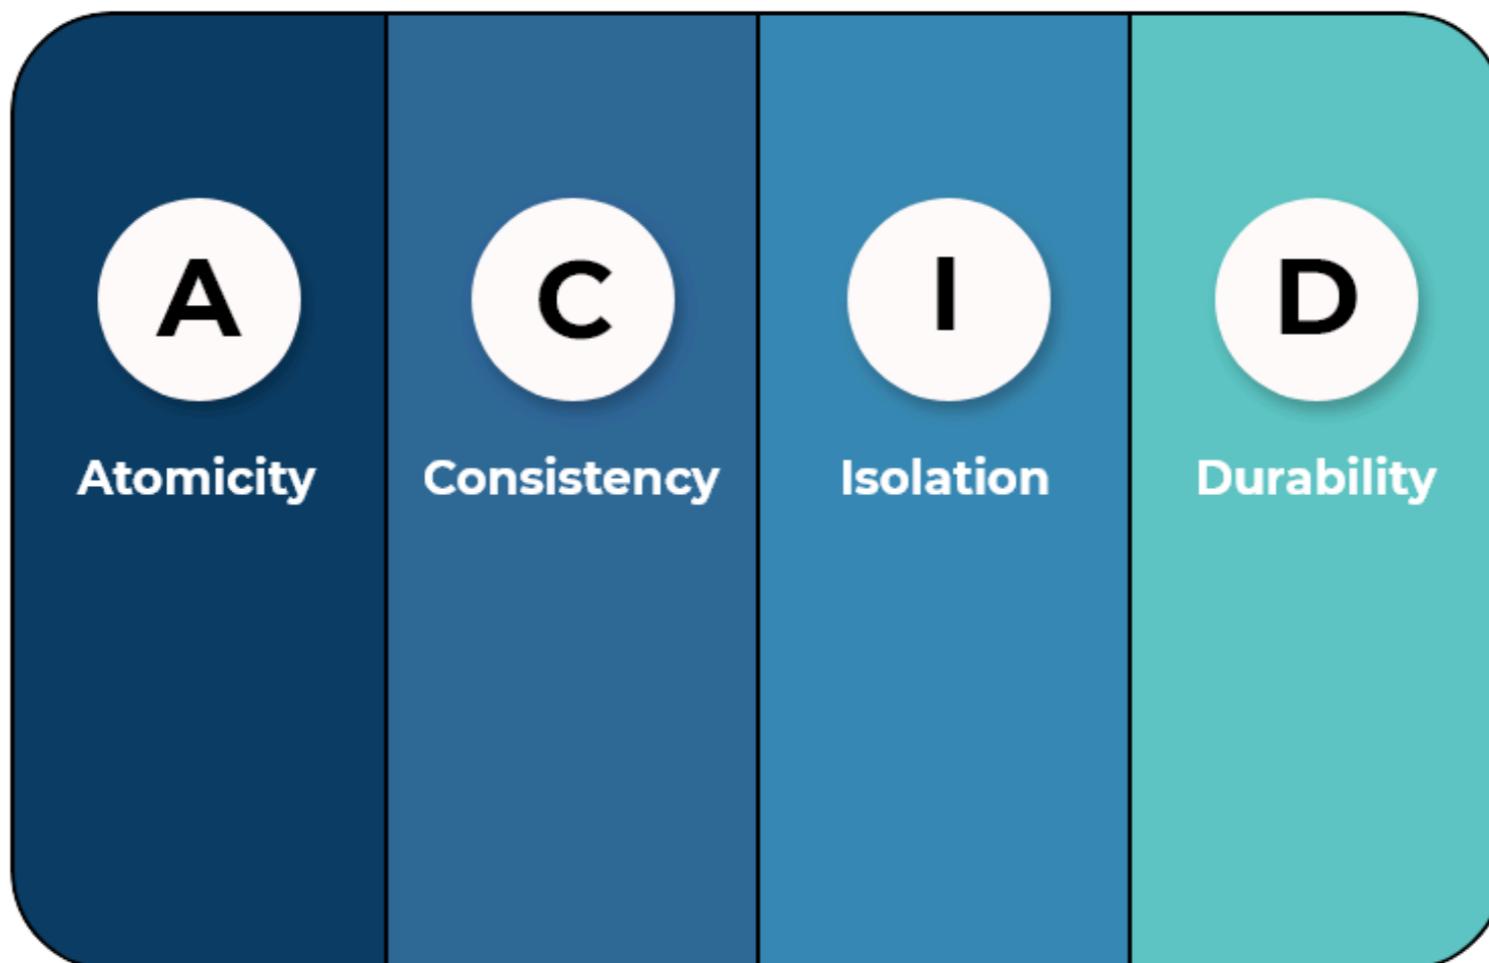
# ACID in RDBMS

<https://en.wikipedia.org/wiki/ACID>

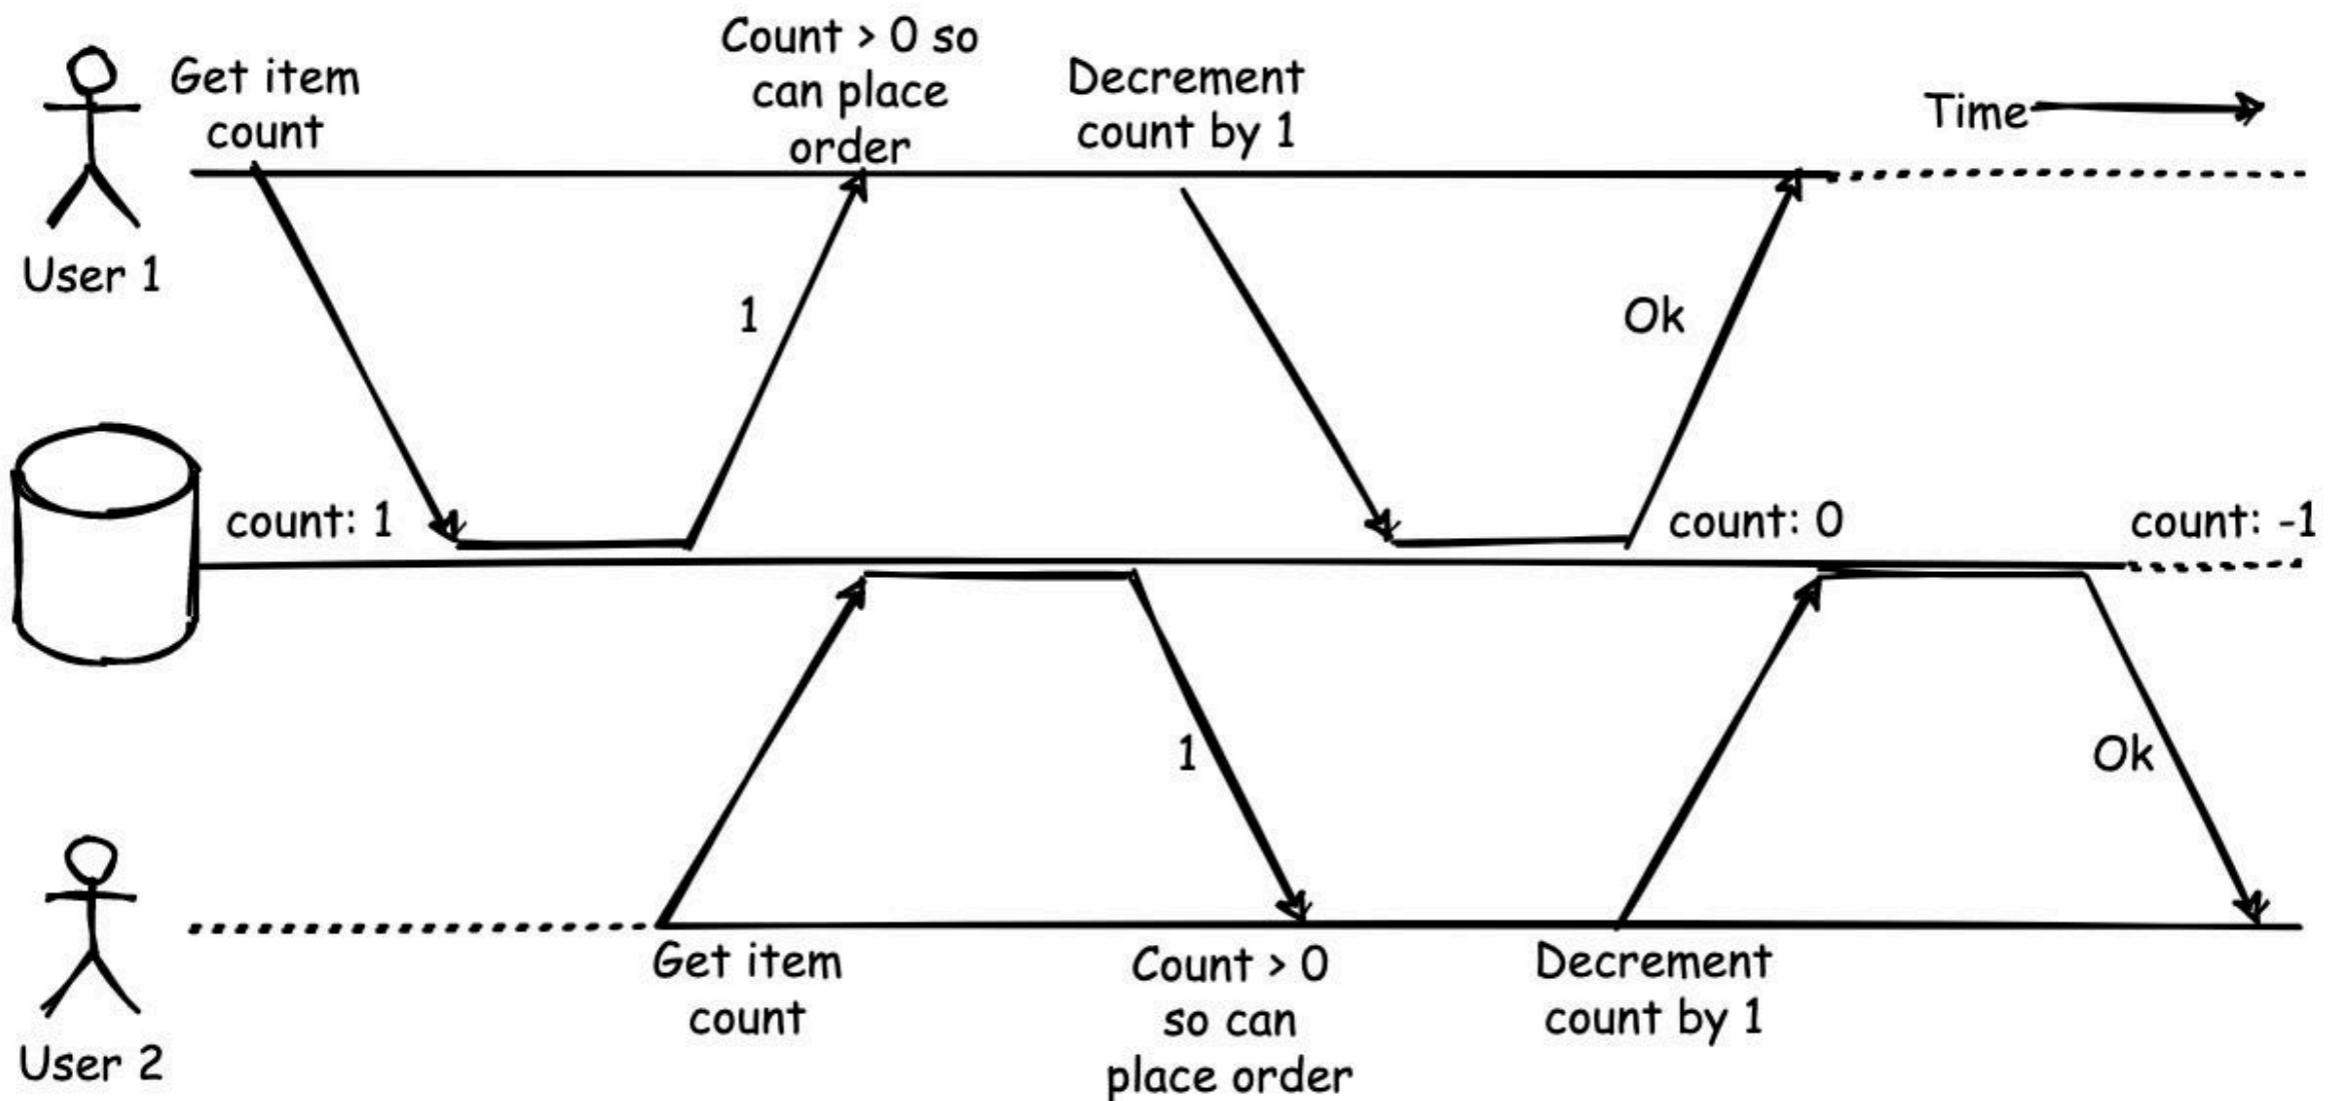


# ACID in RDBMS

Ensure database in valid state even in the event of unexpected errors



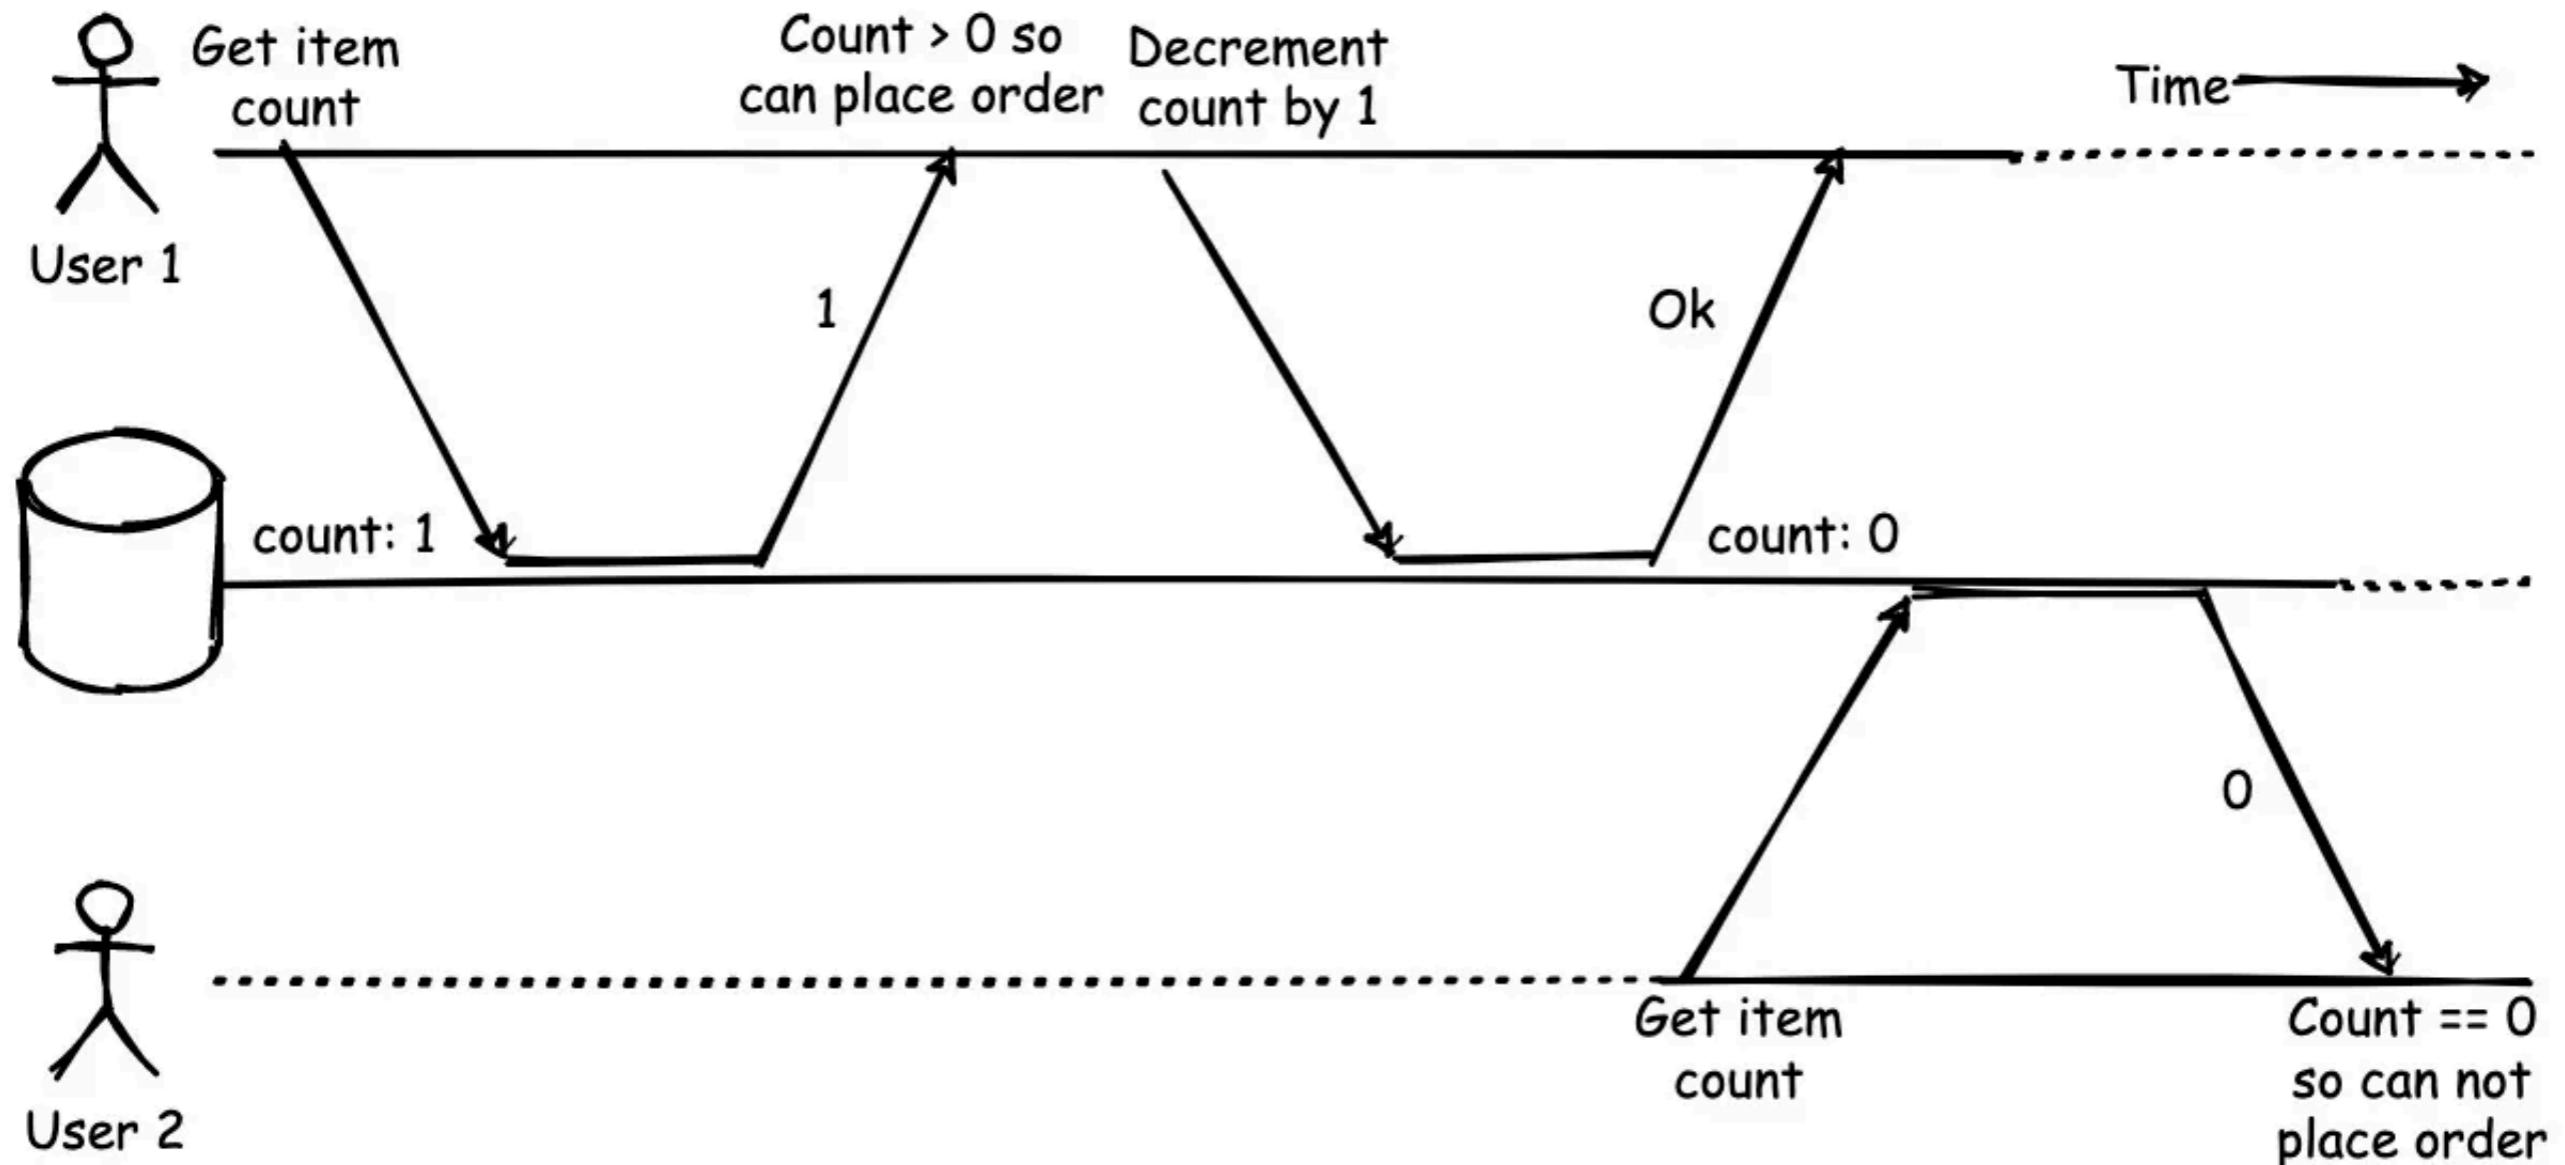
# Race condition



[https://en.wikipedia.org/wiki/Race\\_condition#Example](https://en.wikipedia.org/wiki/Race_condition#Example)



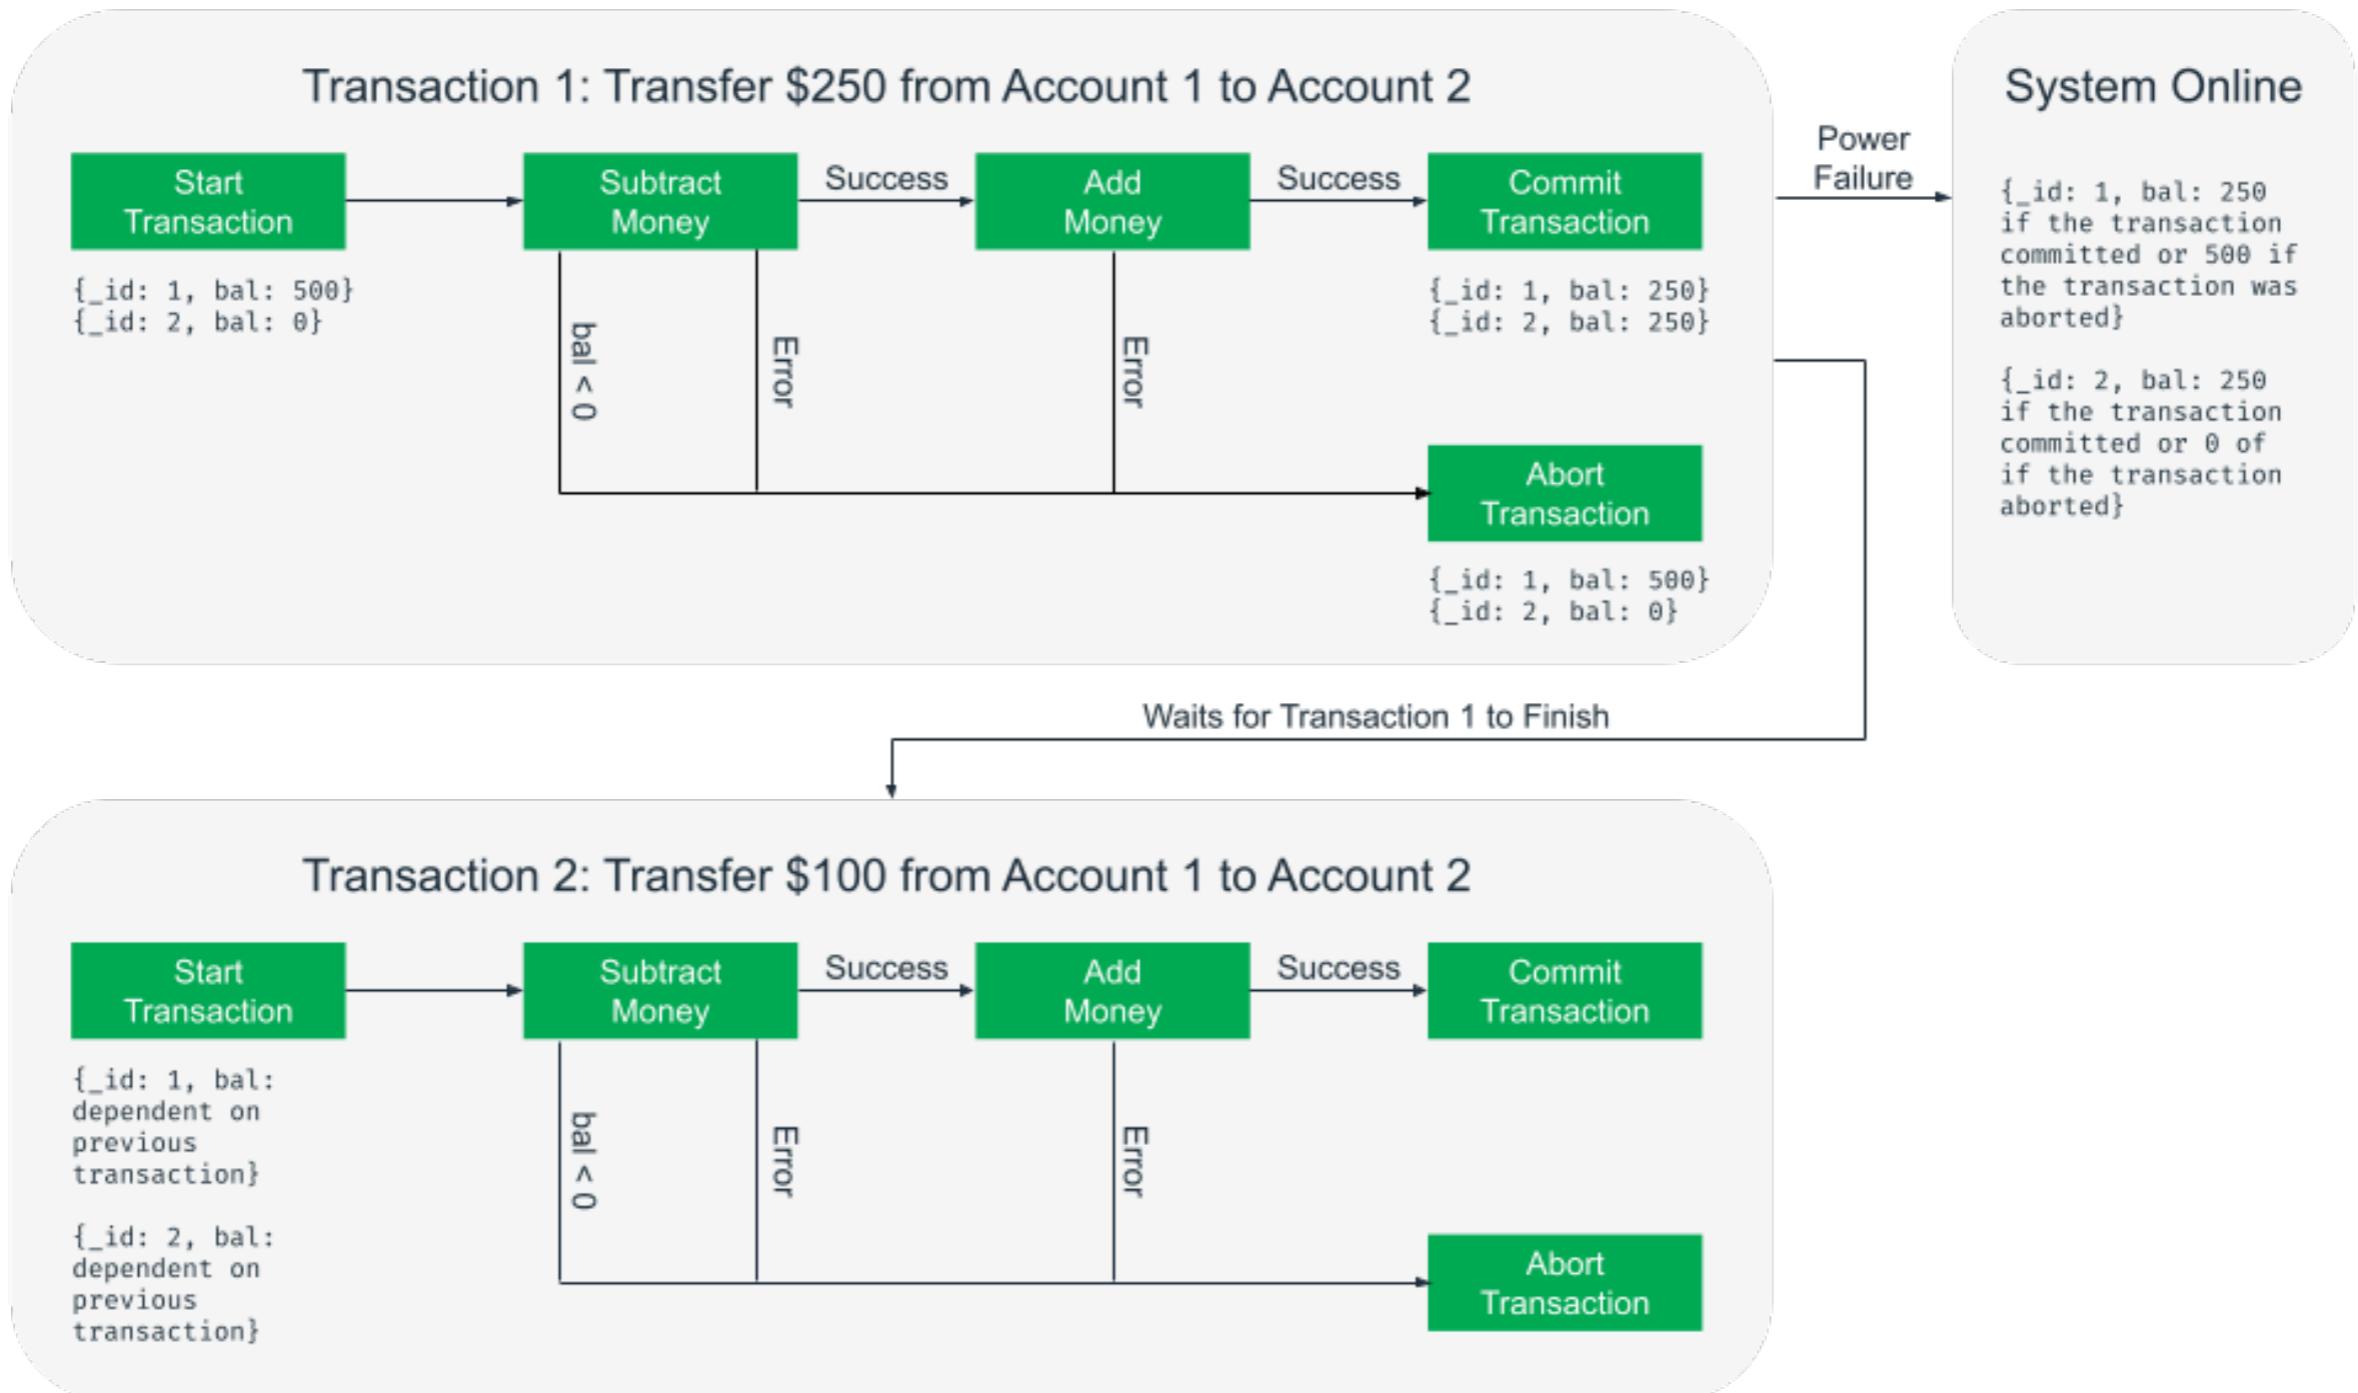
# Sequencial



<https://www.mongodb.com/basics/acid-transactions>



# ACID in RDBMS



<https://www.mongodb.com/basics/acid-transactions>



# StructureQueryLanguage



# **SQL**

Maintain data in **tables**

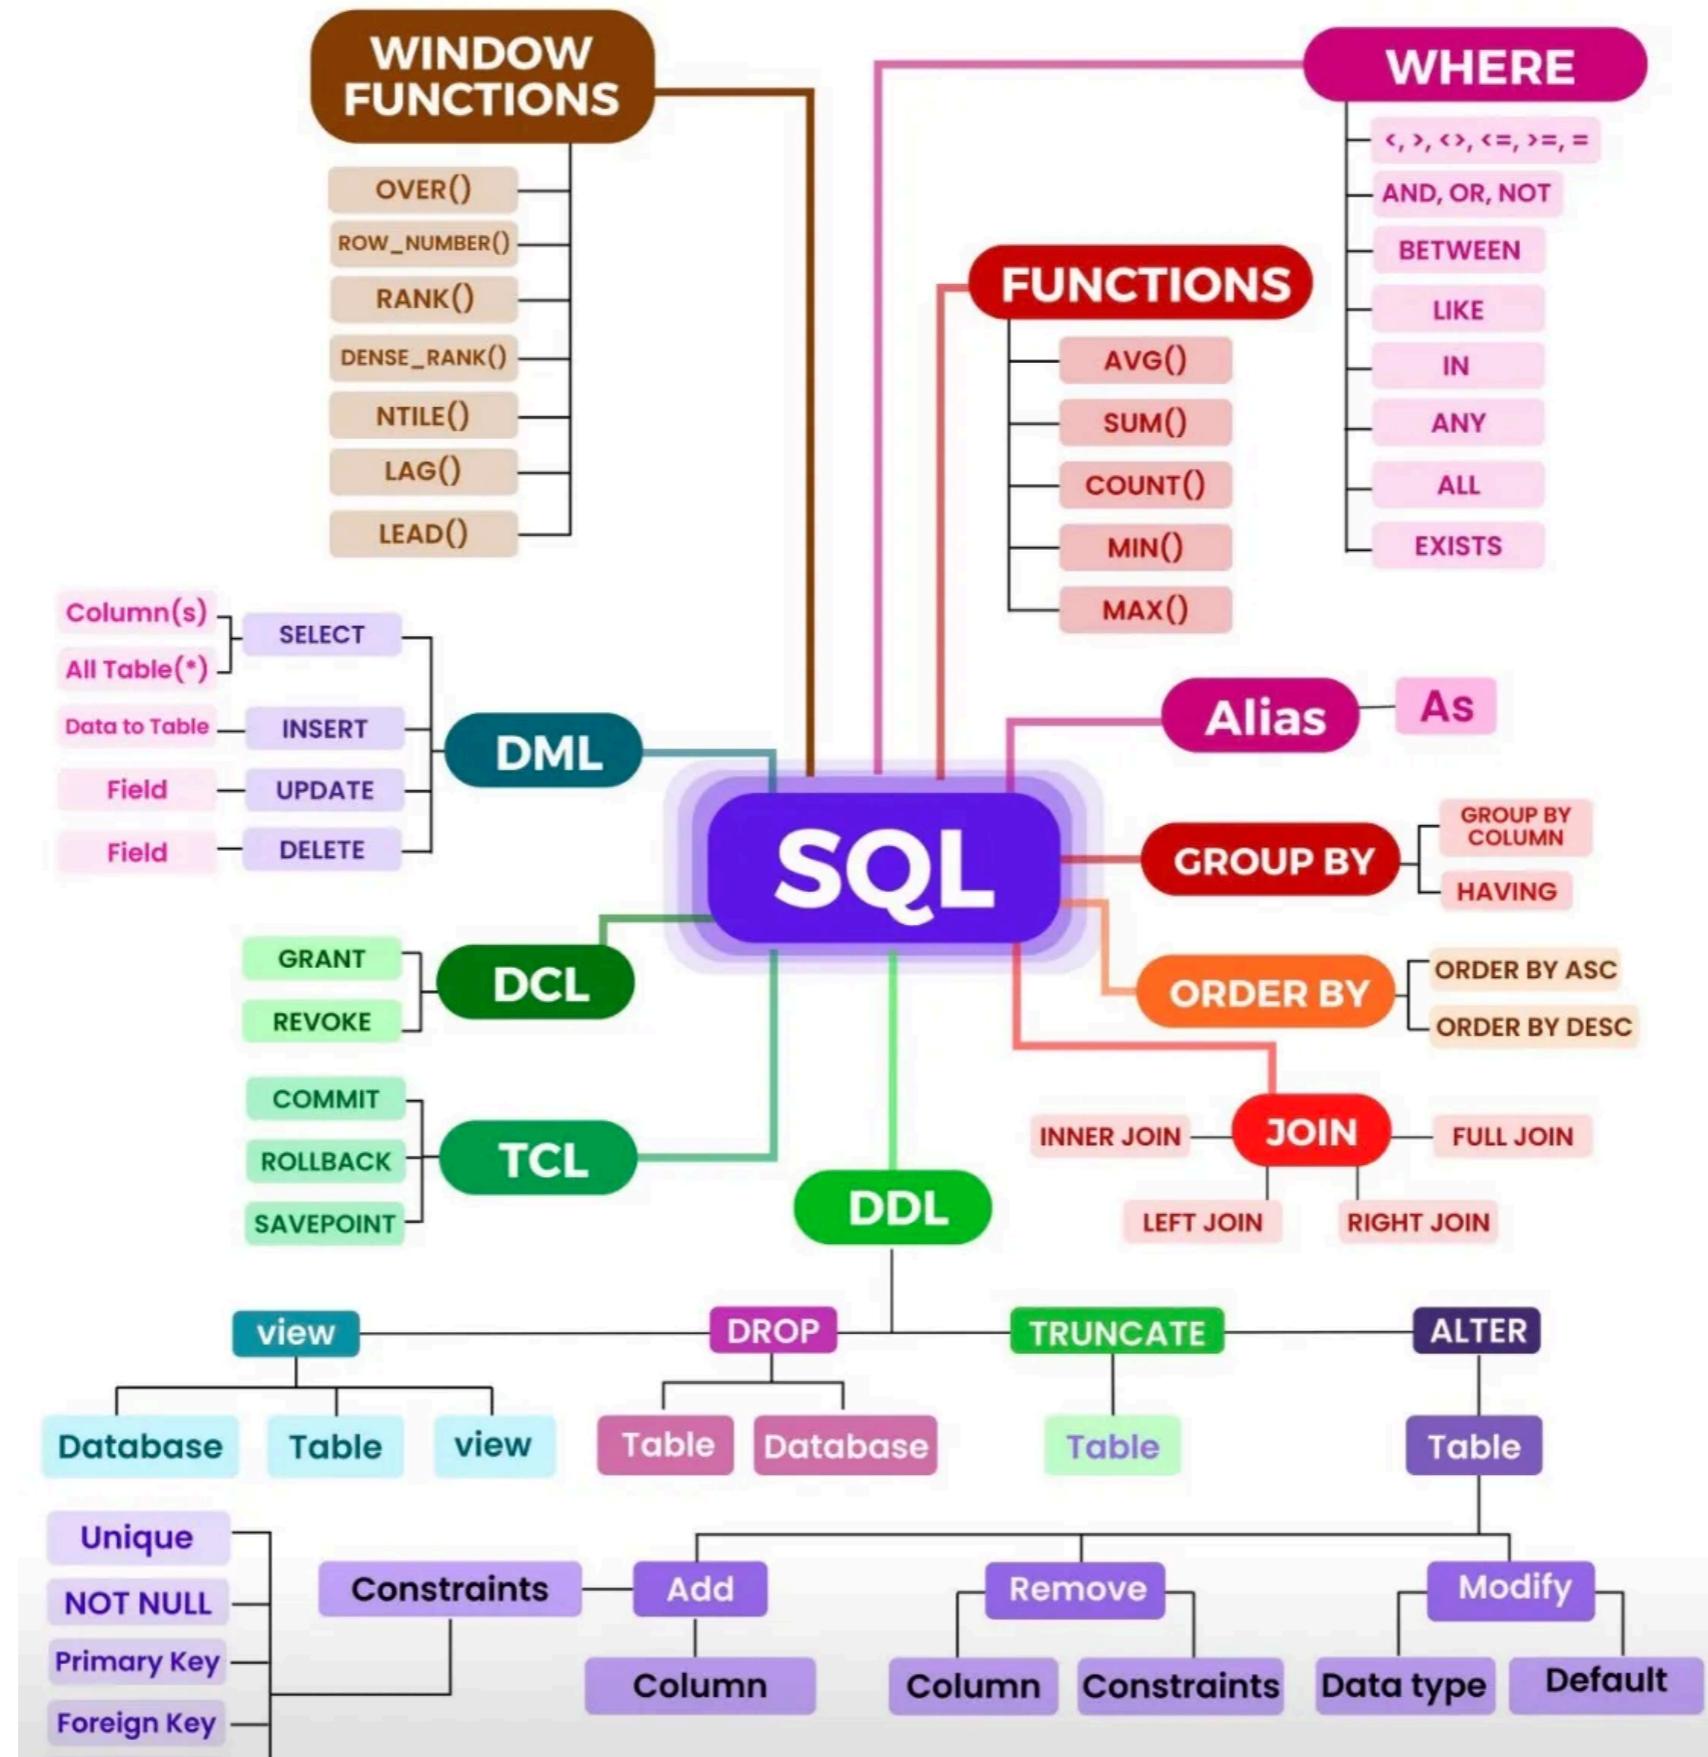
**Pre-define structure**

Each table has **Primary key**

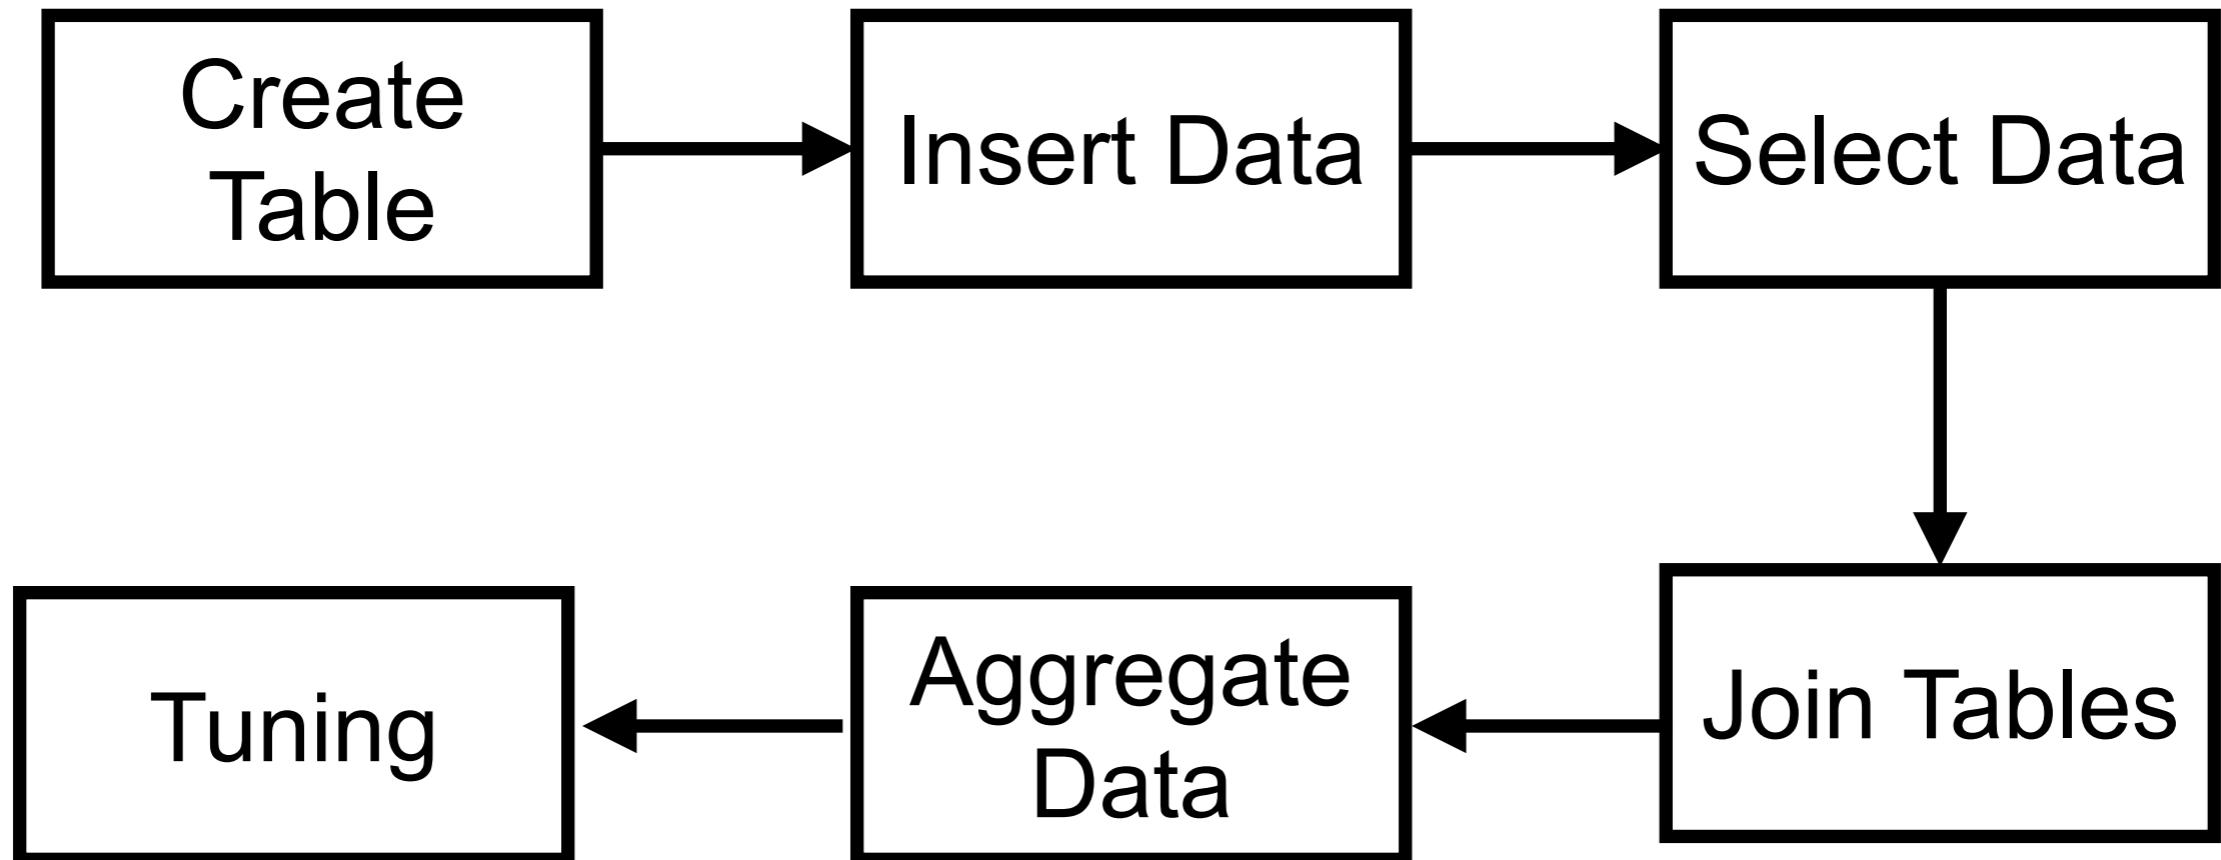
**Many columns**

Data is represented in term of **row**





# Learning paths



# Create table



# Select Data



**SELECT** columns  
**FROM** table  
**WHERE** condition A  
**AND** condition B  
**GROUP BY** columns  
**ORDER BY** columns **ASC|DESC**



# Select Data process

Table A

Table B

Table C



# Select Data process

Table A

Table B

Table C

FROM + JOIN  
(merge)

Table A, B, C



# Select Data process

Table A

Table B

Table C

FROM + JOIN  
(merge)

Table A, B, C

WHERE  
(filter)

Table A, B, C



# Select Data process

Table A

Table B

Table C

FROM + JOIN  
(merge)

Table A, B, C

WHERE  
(filter)

Table A, B, C

Group By  
(Group)



# Select Data process

Table A

Table B

Table C

FROM + JOIN  
(merge)

Table A, B, C

WHERE  
(filter)

Table A, B, C

Group By  
(Group)



Having  
(filter)

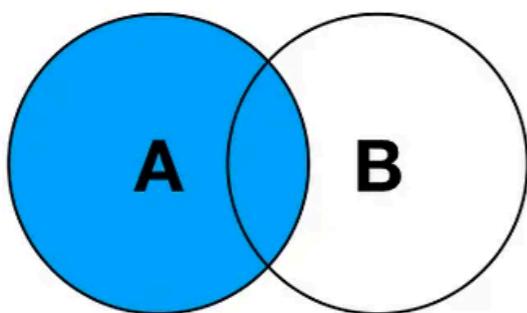


# Joining

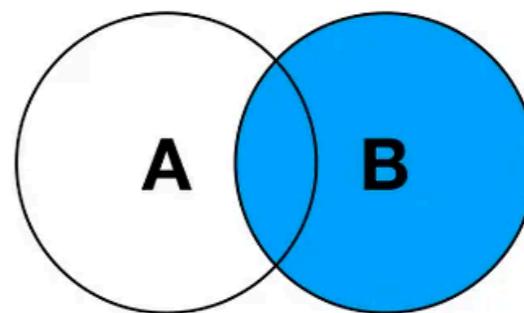
<https://www.postgresql.org/docs/current/tutorial-join.html>



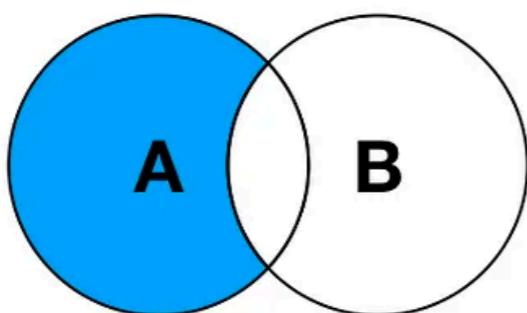
# SQL JOINS



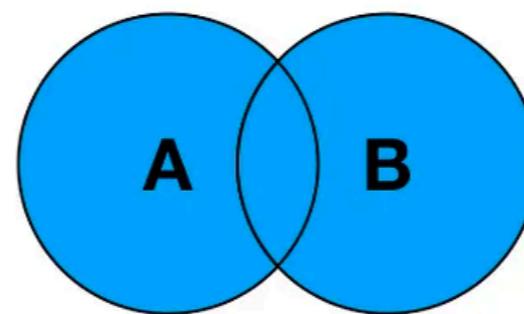
**LEFT JOIN**



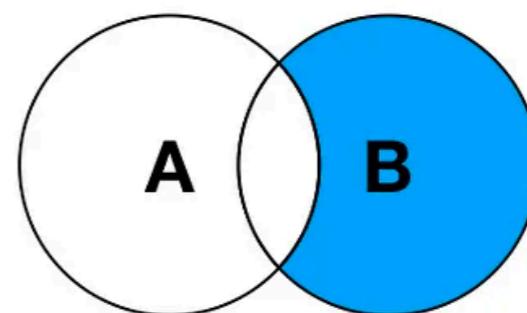
**RIGHT JOIN**



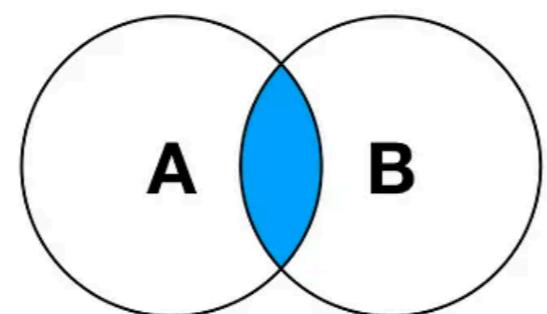
**LEFT JOIN EXCLUDING  
INNER JOIN**



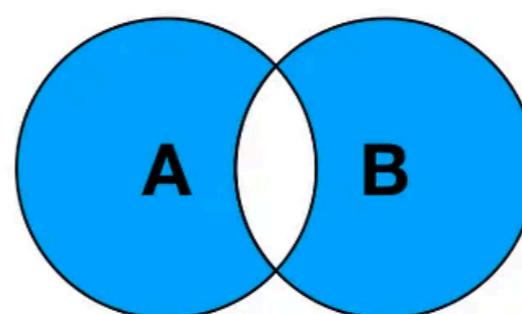
**FULL OUTER JOIN**



**RIGHT JOIN EXCLUDING  
INNER JOIN**



**INNER JOIN**



**FULL OUTER JOIN EXCLUDING  
INNER JOIN**



Self-join

# Join in PostgreSQL

Left join

Inner join

Right join

Full outer  
join



# Workshop

Basket A

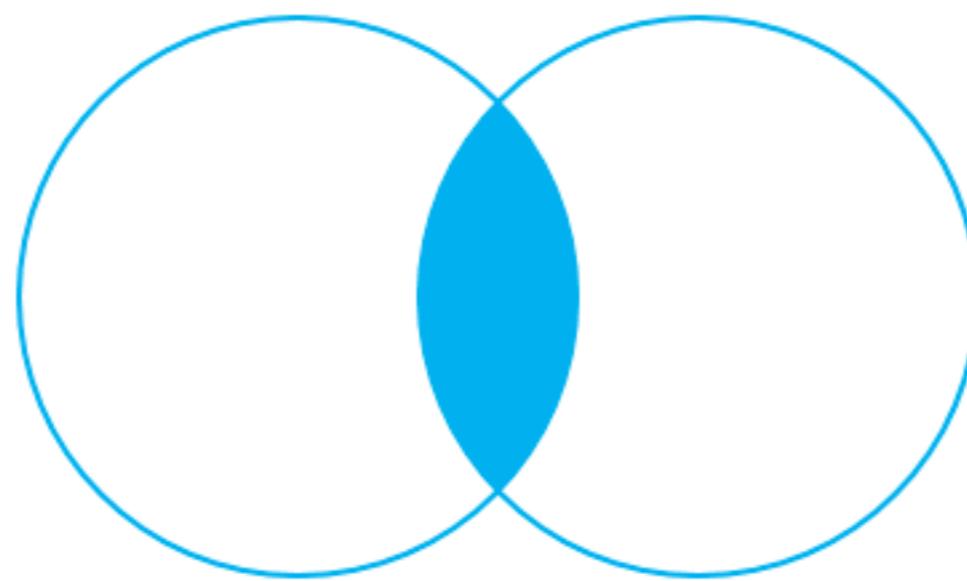
Apple  
Orange  
Banana  
Cucumber

Basket B

Orange  
Apple  
Watermelon  
Pear



# Inner join



INNER JOIN



# Inner join

Basket A

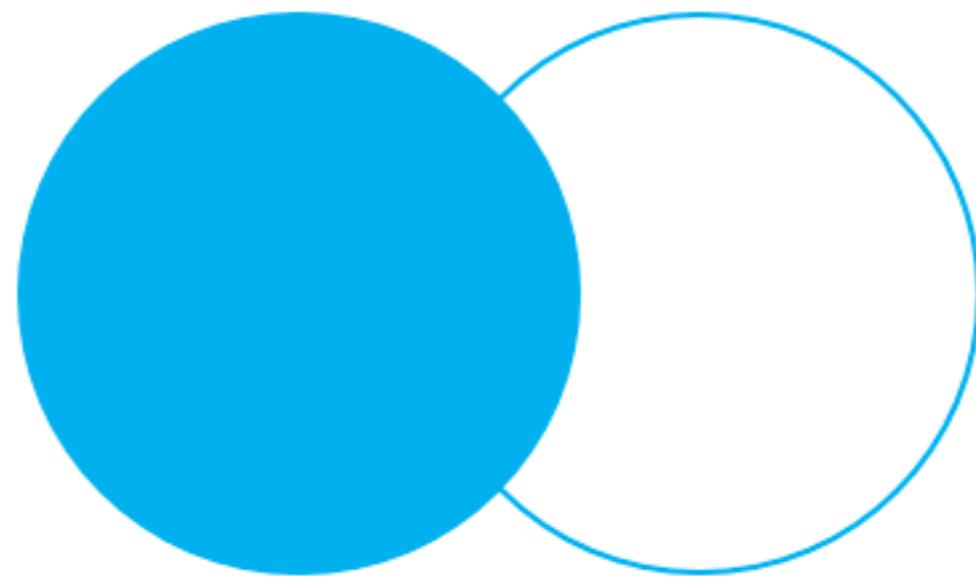
Apple  
Orange  
Banana  
Cucumber

Basket B

Orange  
Apple  
Watermelon  
Pear



# Left join (1)



LEFT OUTER JOIN



# Left join (1)

Basket A

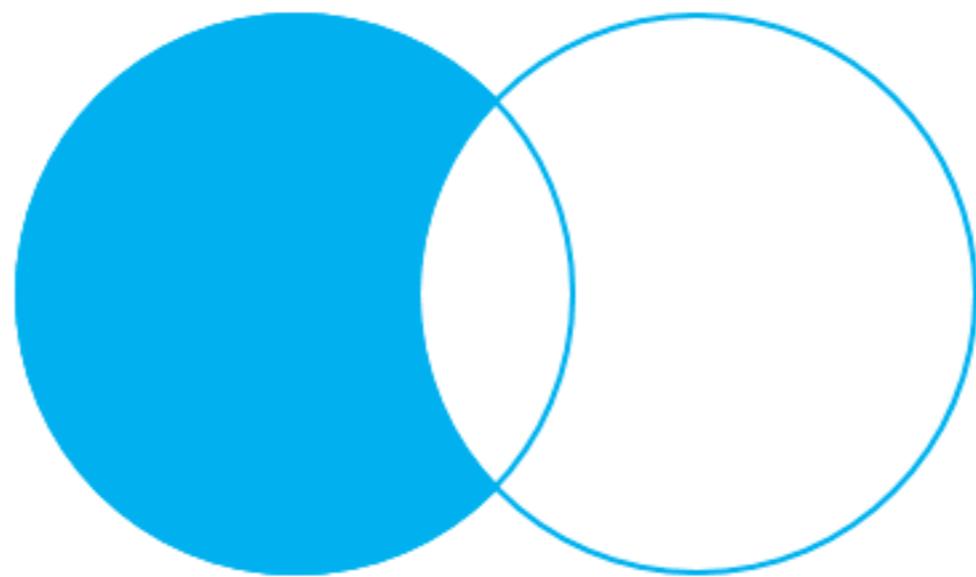
Apple  
Orange  
Banana  
Cucumber

Basket B

Orange  
Apple  
Watermelon  
Pear



# Left join (2)



LEFT OUTER JOIN – only  
rows from the left table



# Left join (2)

Basket A

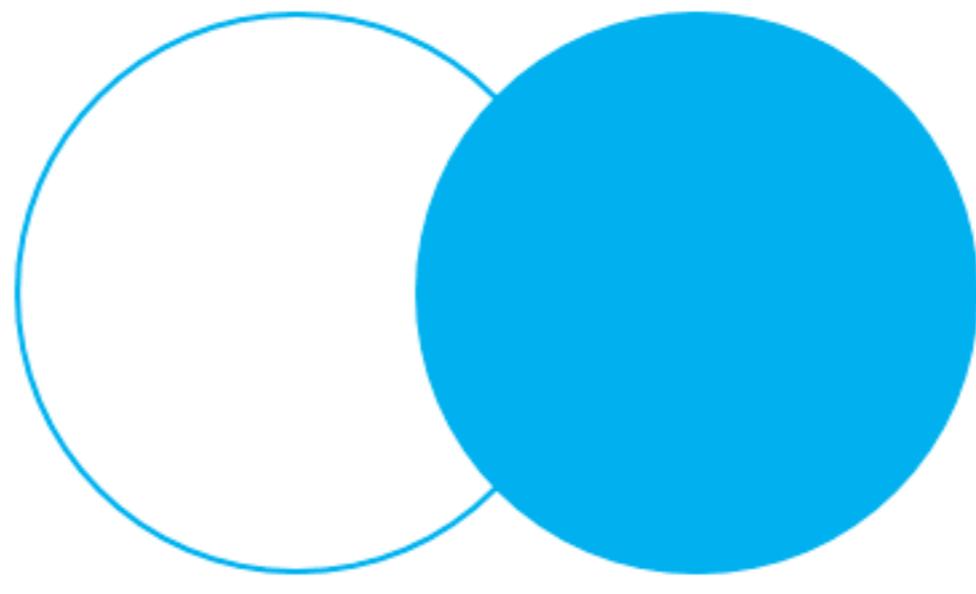
Apple  
Orange  
**Banana**  
**Cucumber**

Basket B

Orange  
Apple  
Watermelon  
Pear



# Right join (1)



RIGHT OUTER JOIN



# Right join (1)

Basket A

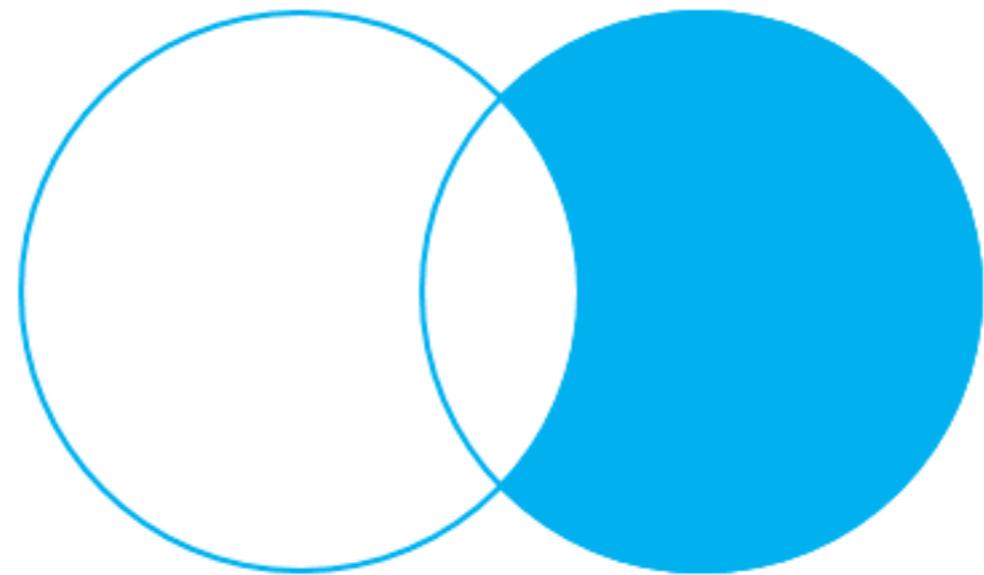
Apple  
Orange  
**Banana**  
**Cucumber**

Basket B

Orange  
Apple  
**Watermelon**  
**Pear**



# Right join (2)



RIGHT OUTER JOIN – only  
rows from the right table



# Right join (2)

Basket A

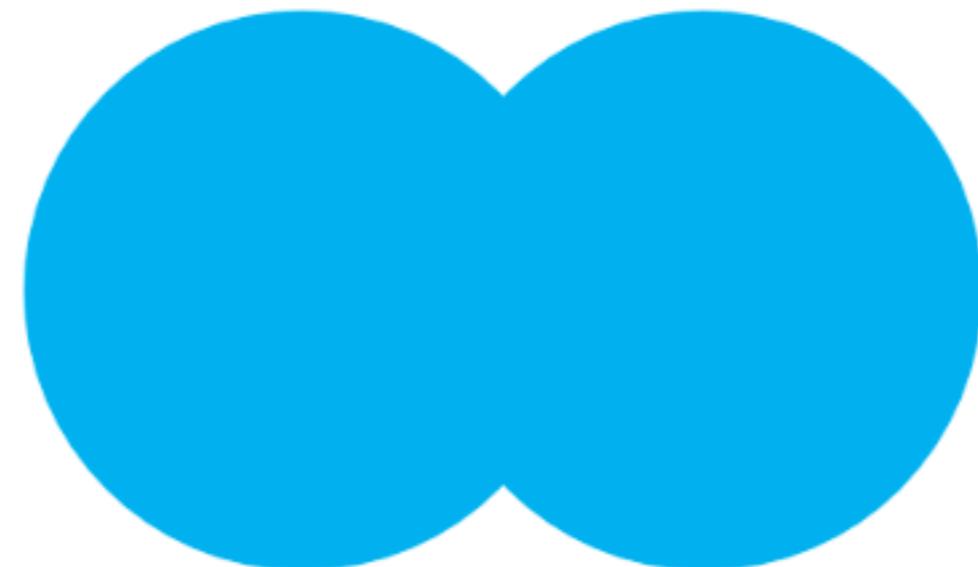
Apple  
Orange  
Banana  
Cucumber

Basket B

Orange  
Apple  
**Watermelon**  
**Pear**



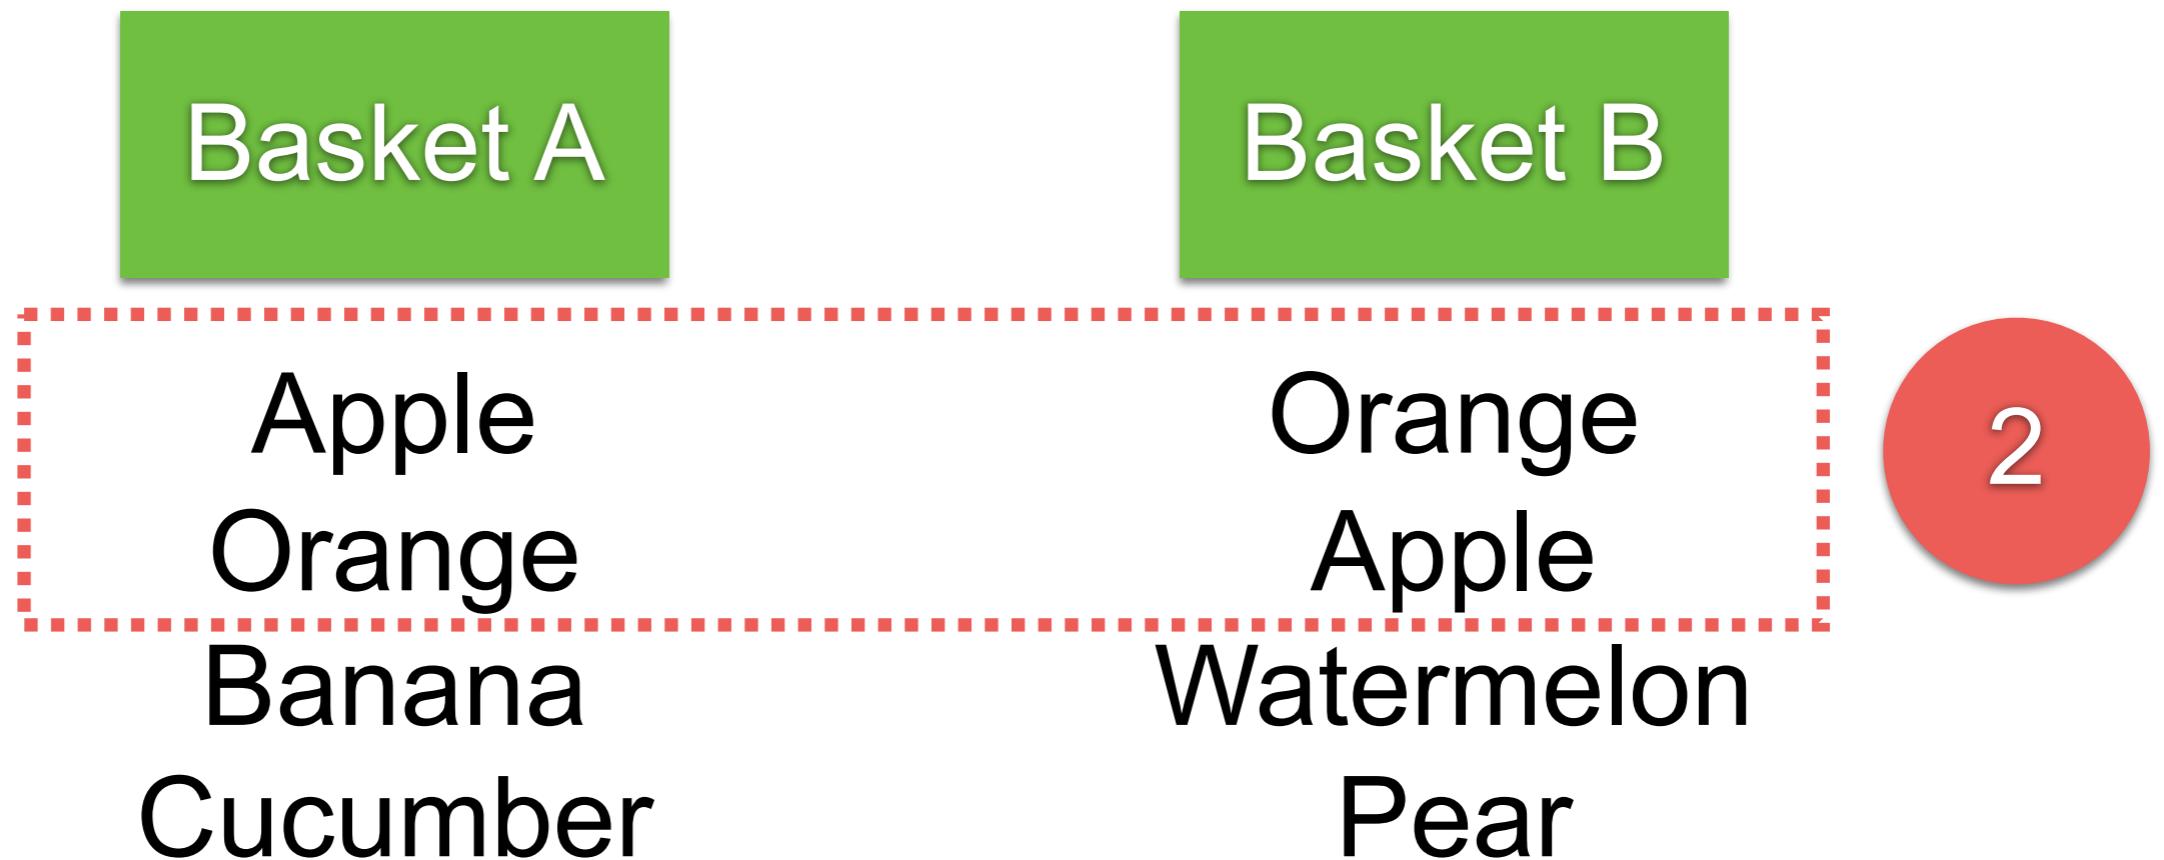
# Full outer join (1)



FULL OUTER JOIN



# Full outer join (1)



# Full outer join (1)

Basket A

2

Apple  
Orange  
Banana  
Cucumber

Basket B

Orange  
Apple  
Watermelon  
Pear



# Full outer join (1)

Basket A

Apple  
Orange  
Banana  
Cucumber

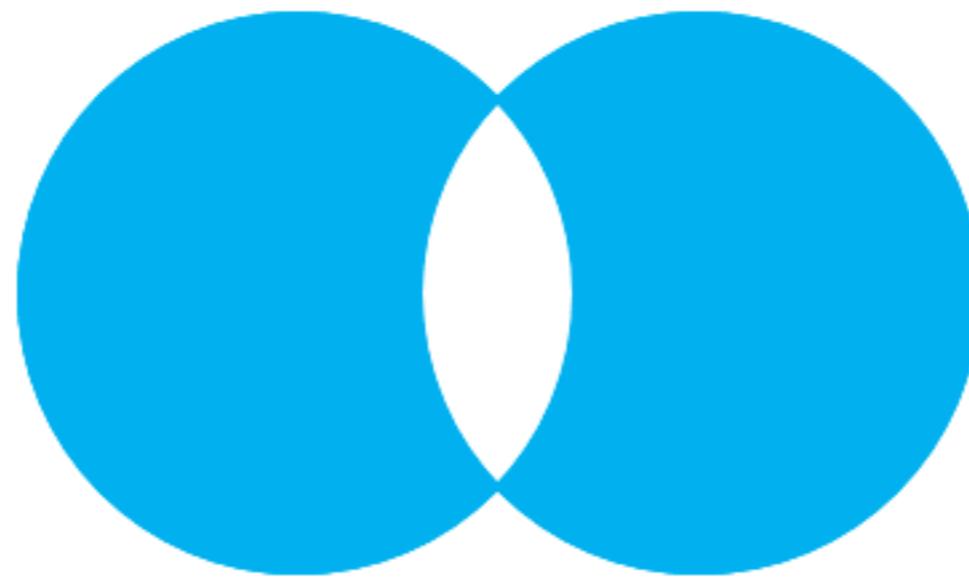
Basket B

Orange  
Apple  
Watermelon  
Pear

2



# Full outer join (2)



FULL OUTER JOIN – only  
rows unique to both tables



# Full outer join (2)

Basket A

2

Apple  
Orange  
Banana  
Cucumber

Basket B

Orange  
Apple  
Watermelon  
Pear



# Full outer join (2)

Basket A

Apple  
Orange  
Banana  
Cucumber

Basket B

Orange  
Apple  
Watermelon  
Pear

2



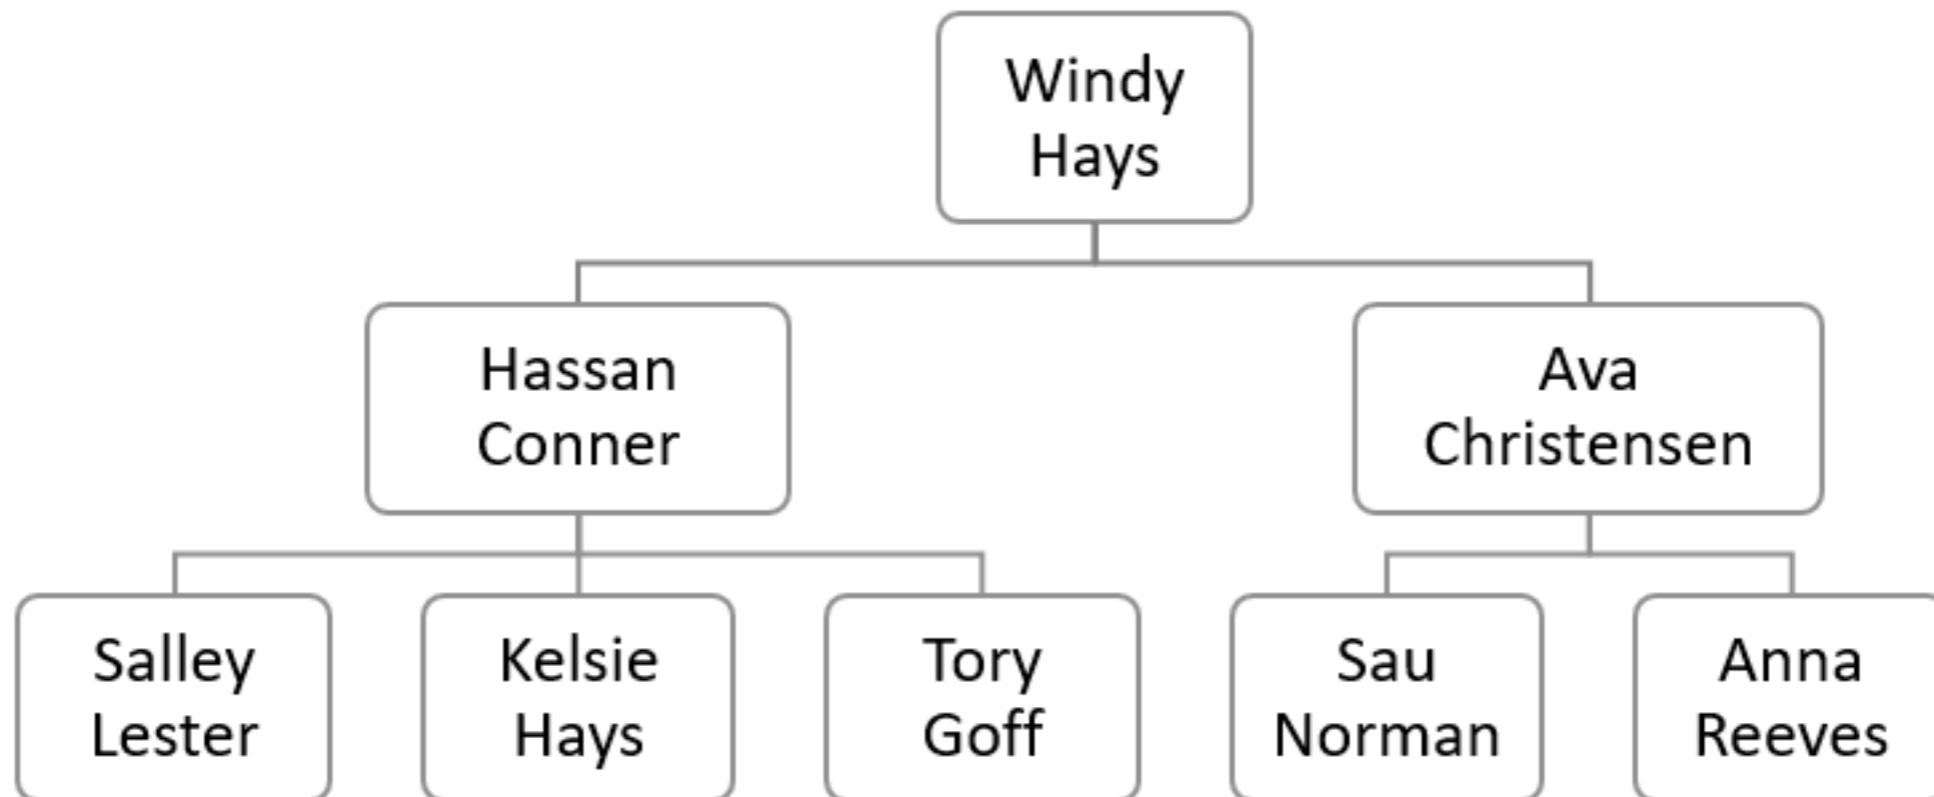
# Workshop :: Joins

[https://github.com/up1/workshop-database/blob/main/  
workshop/joins.md](https://github.com/up1/workshop-database/blob/main/workshop/joins.md)



# Self-join

Regular join that joins a table to itself  
Query hierarchical data



# Workshop :: Self-Join

[https://github.com/up1/workshop-database/blob/main/  
workshop/self-join.md](https://github.com/up1/workshop-database/blob/main/workshop/self-join.md)



# Problems ?



# Problems

Slow query

Accurate knowledge is required

Add complexity

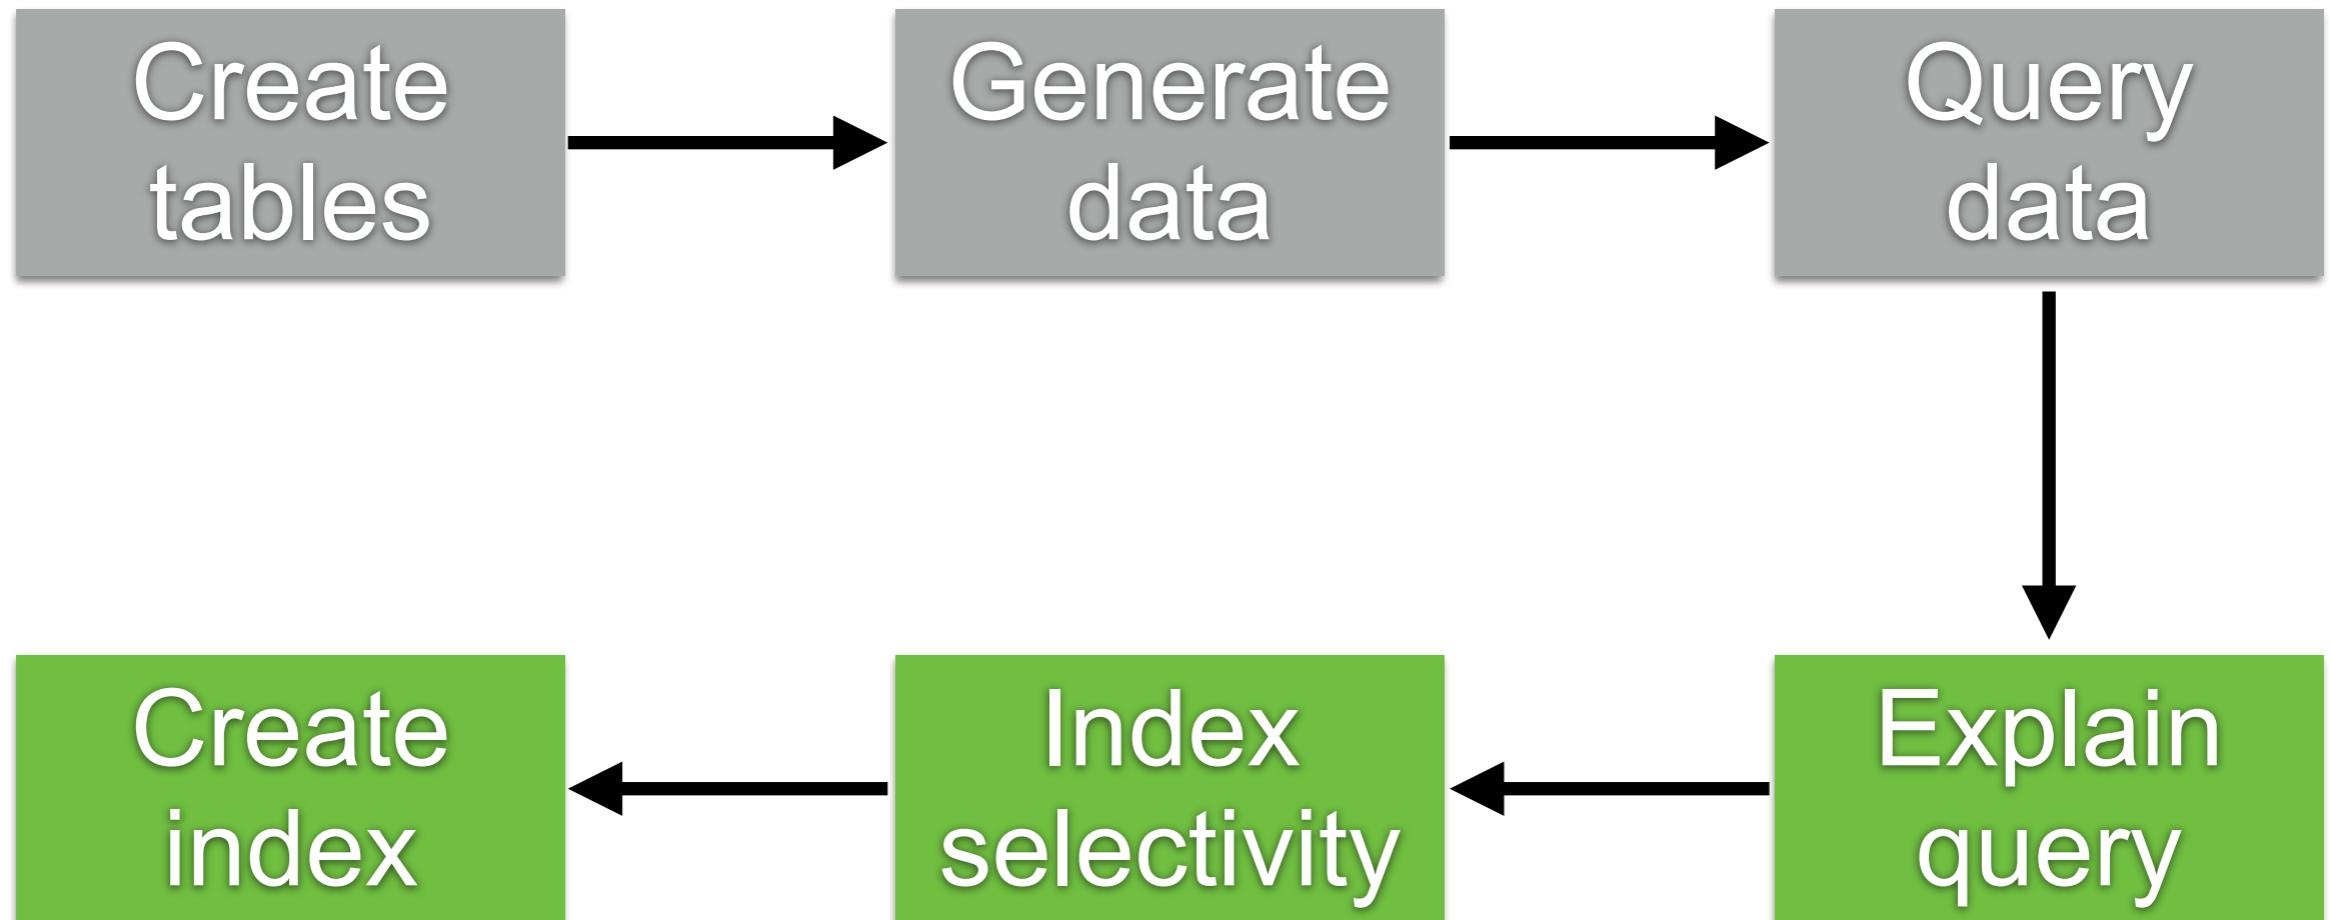
Need denormalization ?



# Improvement to query data

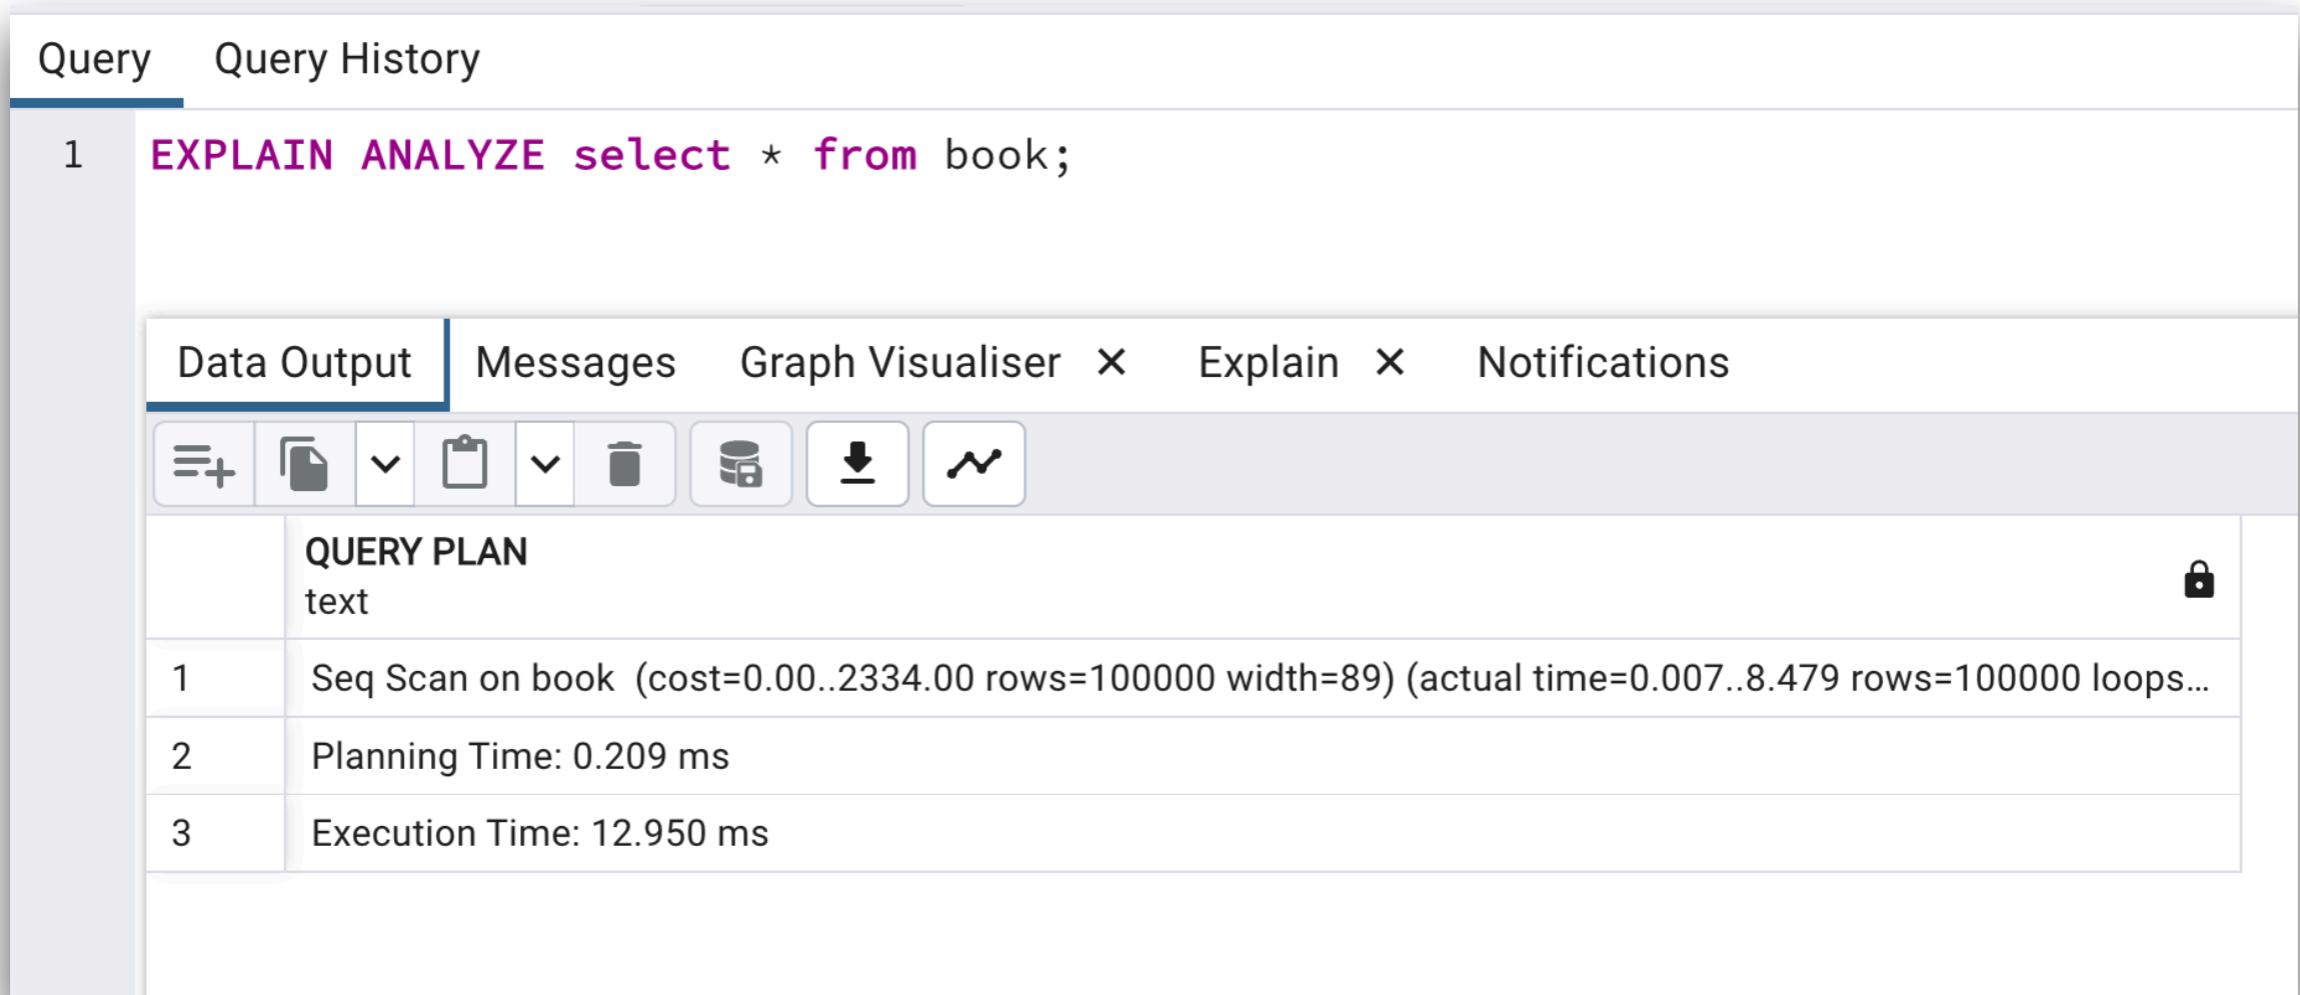


# Step in workshop



# Query plan in PostgreSQL

**EXPLAIN ANALYZE <SQL>**



The screenshot shows a PostgreSQL query tool interface. At the top, there are tabs for "Query" and "Query History". Below the tabs, a query is entered: "EXPLAIN ANALYZE select \* from book;". The "Data Output" tab is selected in the navigation bar, which also includes "Messages", "Graph Visualiser", "Explain", and "Notifications". Below the navigation bar is a toolbar with various icons. The main area displays the query plan and execution statistics:

|      | QUERY PLAN  |           |
|------|---|-----------|
| text |   | lock icon |
| 1    | Seq Scan on book (cost=0.00..2334.00 rows=100000 width=89) (actual time=0.007..8.479 rows=100000 loops=1) |           |
| 2    | Planning Time: 0.209 ms   |           |
| 3    | Execution Time: 12.950 ms   |           |

<https://github.com/up1/workshop-database/tree/main/workshop#query-plan>



# Indexing

<https://www.postgresql.org/docs/current/indexes-types.html>



# Index

## Symbols

2PC, 219, 346, 347, 349–351

## A

A\* pathfinding algorithms, 64, 84  
ACID, 197, 208, 217, 219, 321  
ActiveMQ, 92  
adjacency lists, 77  
Advanced Message Queuing Protocol, 125  
Aeron, 402  
aggregation window, 164, 176  
Airbnb, 193, 199, 337  
Amazon, 137, 199, 317  
Amazon API Gateway, 303  
Amazon Web Services, 251, 302  
AML/CFT, 318  
AMM, 409  
AMQP, 125  
Apache James, 231  
append-only, 362, 365  
Apple, 393  
Apple Pay, 315  
application loop, 398  
ask price, 380  
asynchronous, 328  
At-least once, 122  
at-least once, 93, 122  
at-least-once, 331  
at-most once, 93, 122  
at-most-once, 331

atomic commit, 167  
atomic operation, 218  
audit, 360  
Automatic Market Making, 409  
Availability Zone, 253  
Availability Zones, 268  
availability zones, 27  
AVRO, 165  
AWS, 251, 302, 303  
AWS Lambda, 303  
AZ, 268

## B

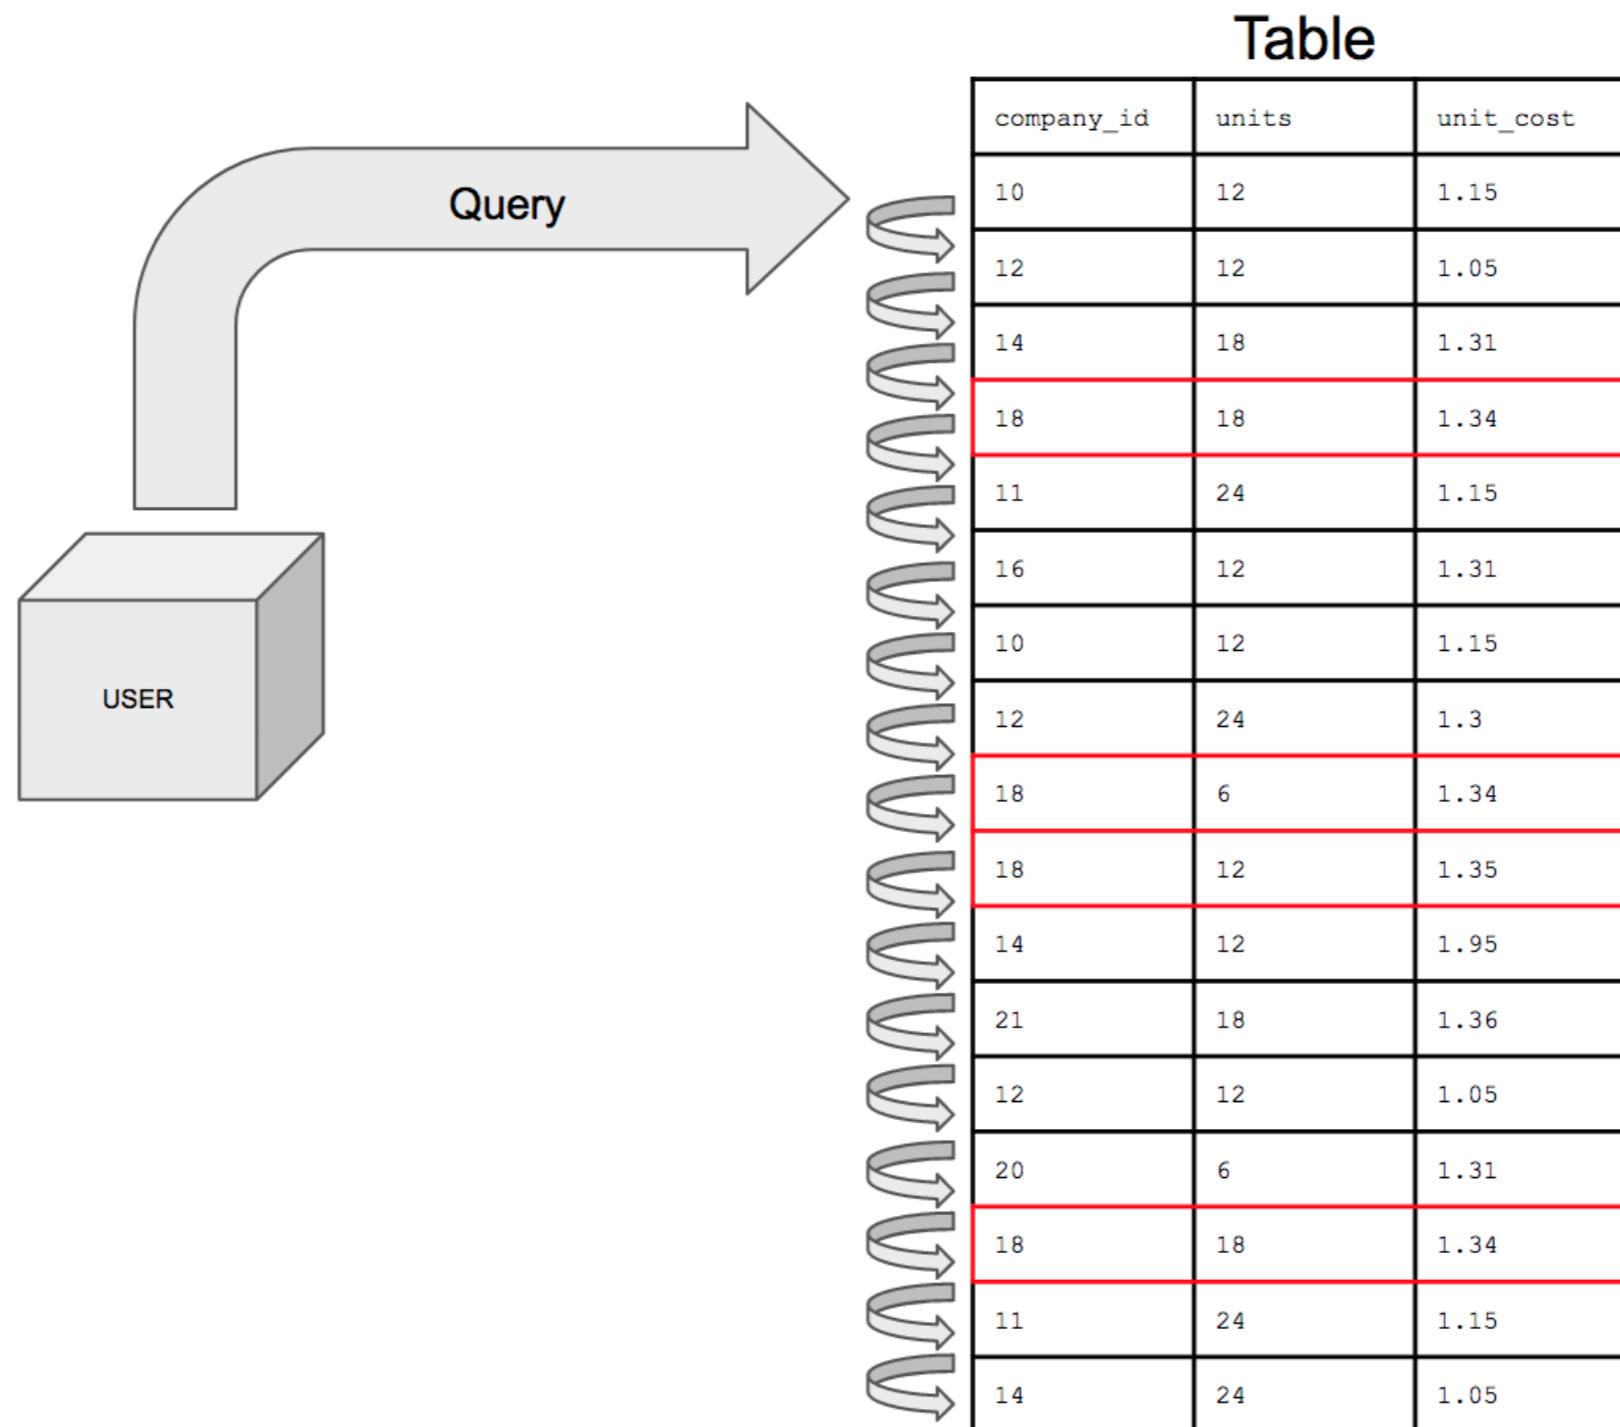
B+ tree, 267  
Backblaze, 272  
base32, 11  
BEAM, 55  
bid price, 380  
Bigtable, 137, 235, 243  
Blue/green deployment, 18  
brokers, 95, 96, 98, 102, 105–107, 113, 118, 120, 122  
buy order, 393

## C

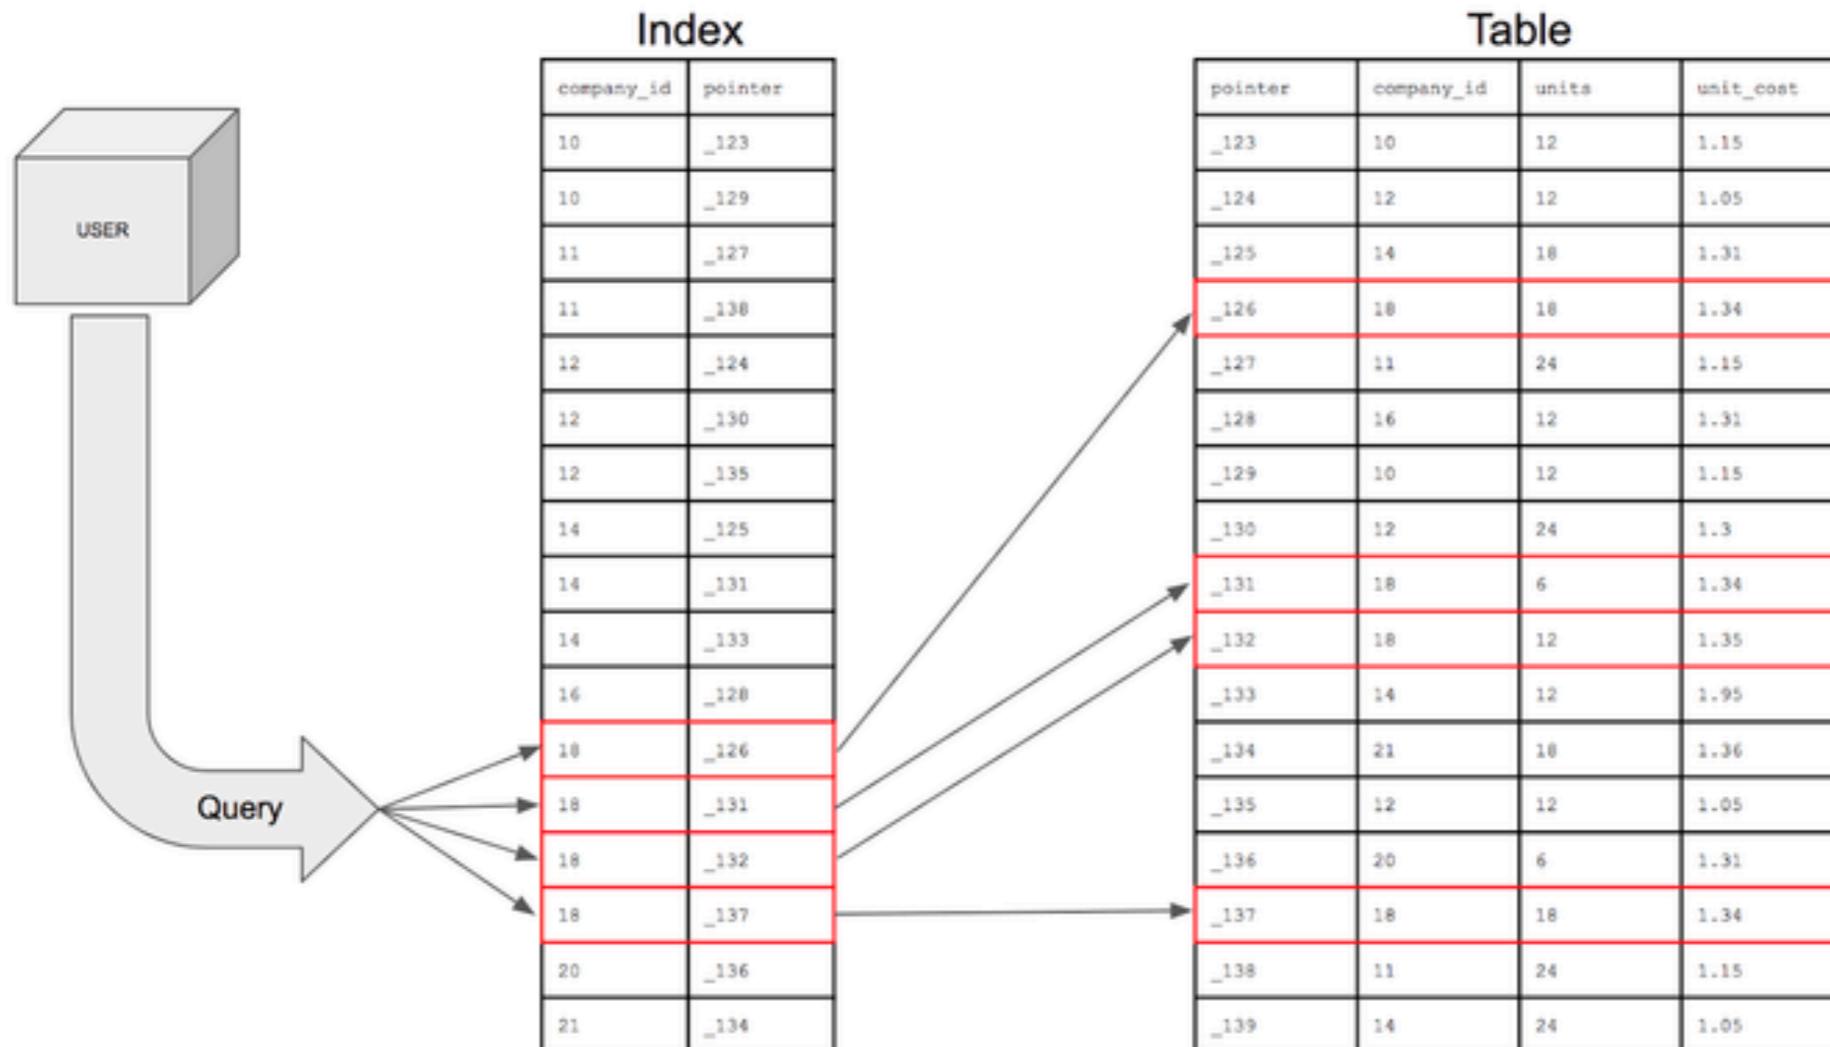
California Consumer Privacy Act, 2  
candlestick chart, 381  
candlestick charts, 384, 387, 396, 407  
CAP theorem, 79



# Without Index (full table scan)



# With Index



# With Index

Table

|         | Column A     | Column B             |
|---------|--------------|----------------------|
| RowId 1 | Some value 1 | Some another value 1 |
| RowId 2 | Some value 2 | Some another value 2 |
| RowId 3 | Some value 3 | Some another value 3 |
| RowId 4 | Some value 1 | Some another value 4 |
| RowId 5 | Some value 5 | Some another value 5 |
| RowId 6 | Some value 6 | Some another value 6 |
| RowId 7 | Some value 1 | Some another value 7 |
| RowId 8 | Some value 8 | Some another value 8 |

Index

|              |                               |
|--------------|-------------------------------|
| Some value 1 | "RowId 1" "RowId 4" "RowId 7" |
| Some value 2 | RowId 2                       |
| Some value 3 | RowId 3                       |
| Some value 5 | RowId 5                       |
| Some value 6 | RowId 6                       |
| Some value 8 | RowId 8                       |



Index by column "Column A"



# Index selectivity

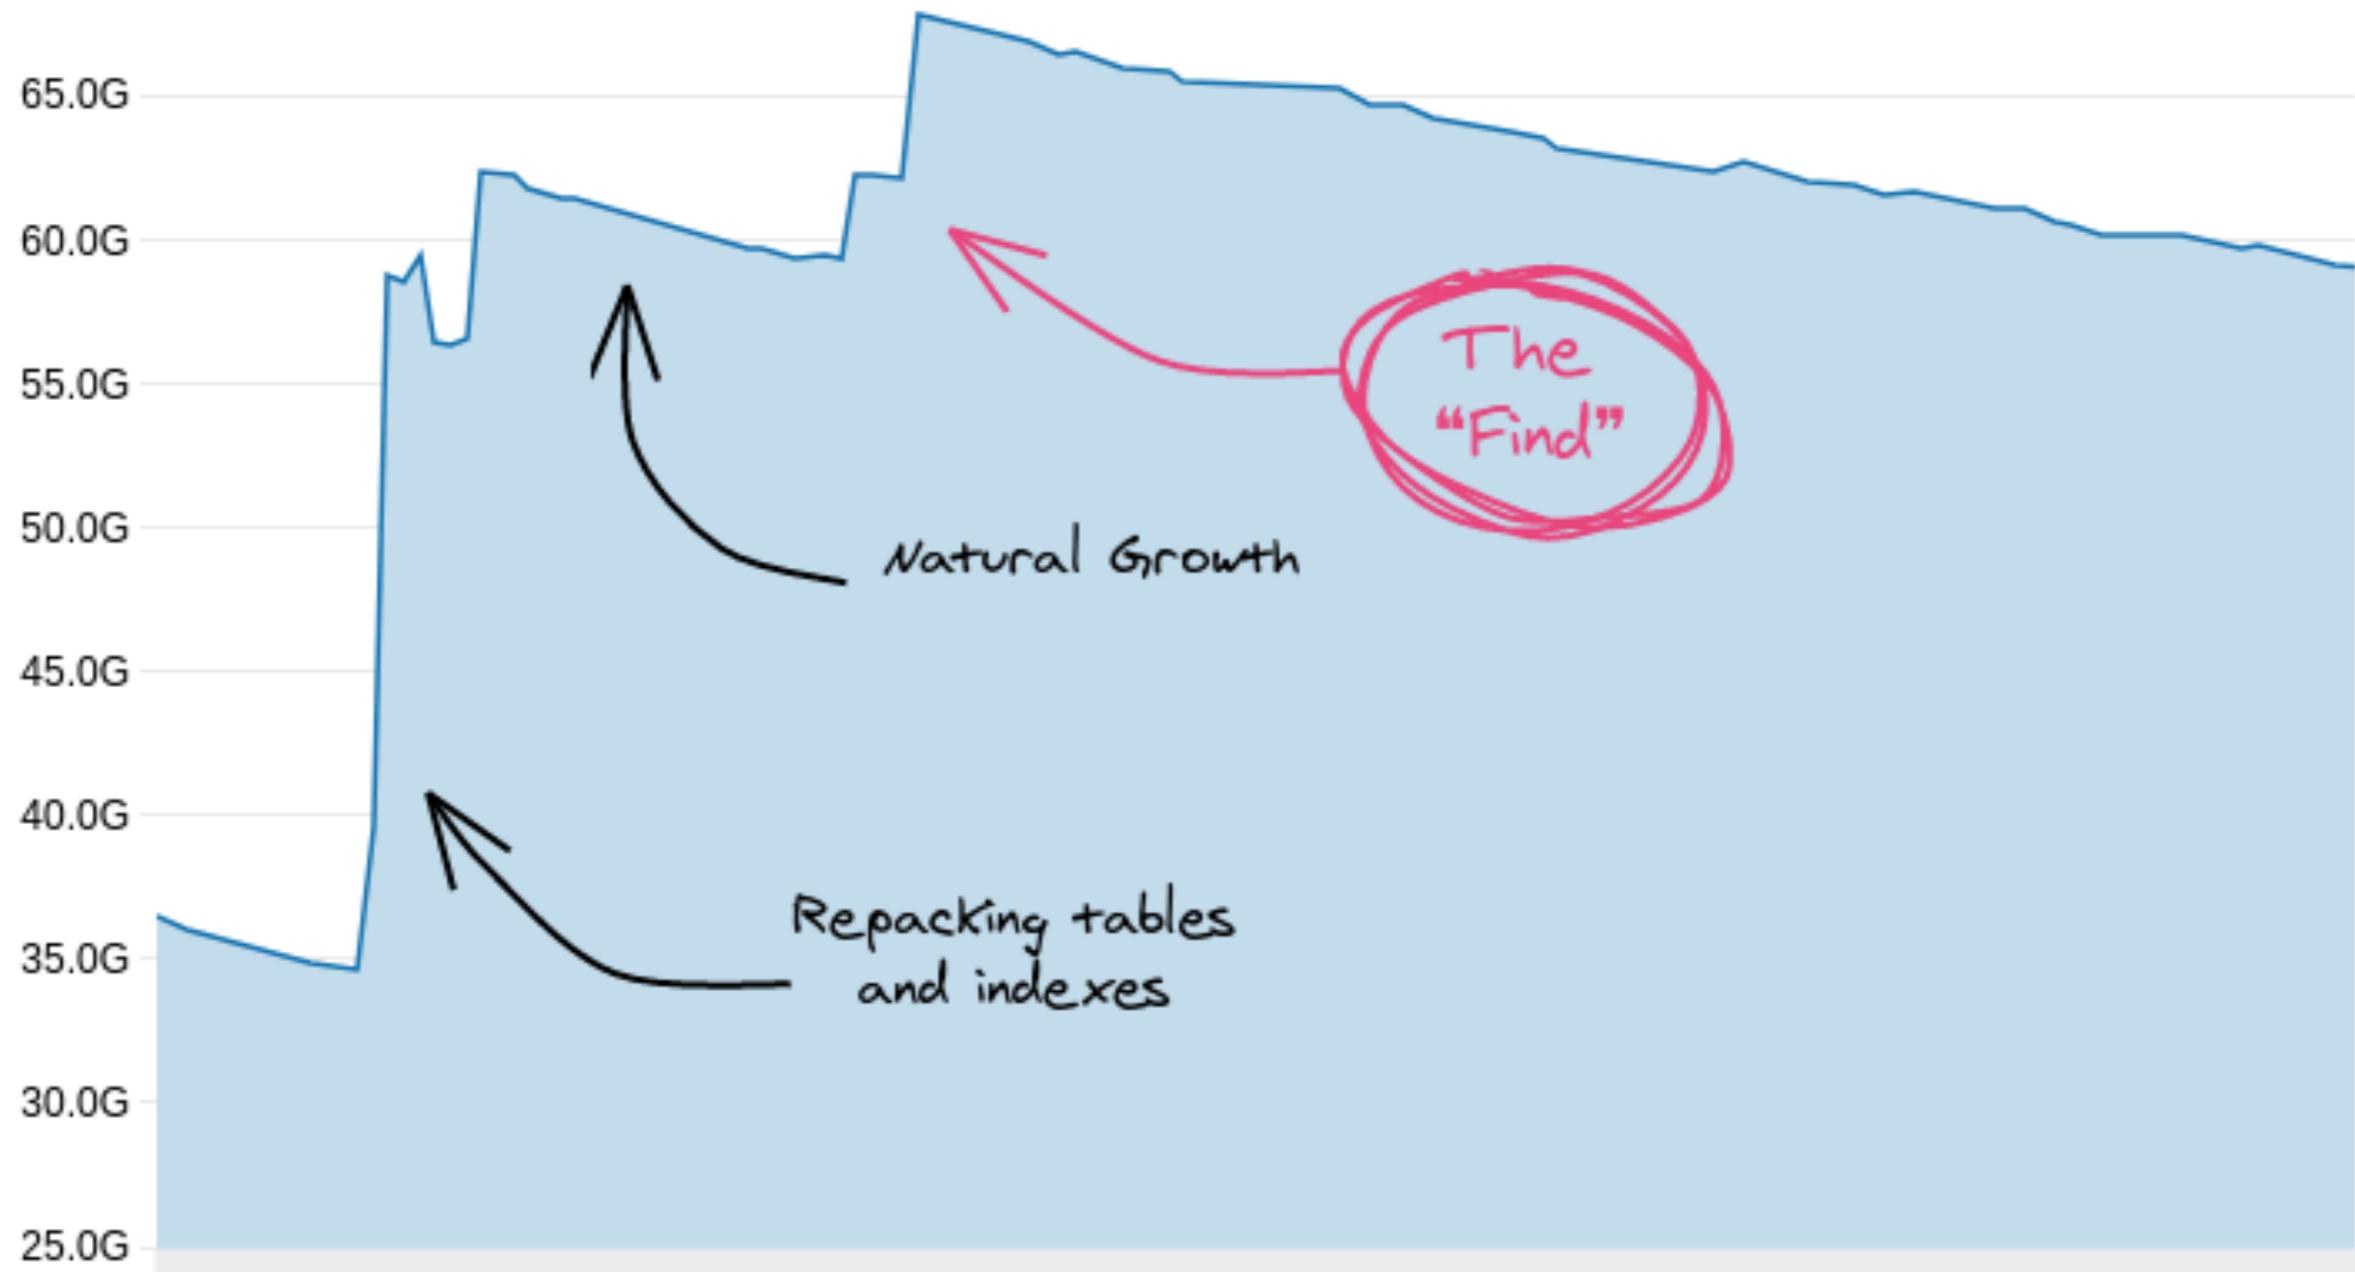
Factor to select a column to create an index

The number of **distinct values** in the indexed column divided by the number of records in the table is called a selectivity of an index

***Prefer indexing columns with selectivity greater than > 0.85 (use Index scan)***



# Unused Index ?



<https://hakibenita.com/postgresql-unused-index-size>



# Types Index in PostgreSQL

B-Tree

Hash

Block range indexes (BRIN)

Generalized inverted index (GIN)

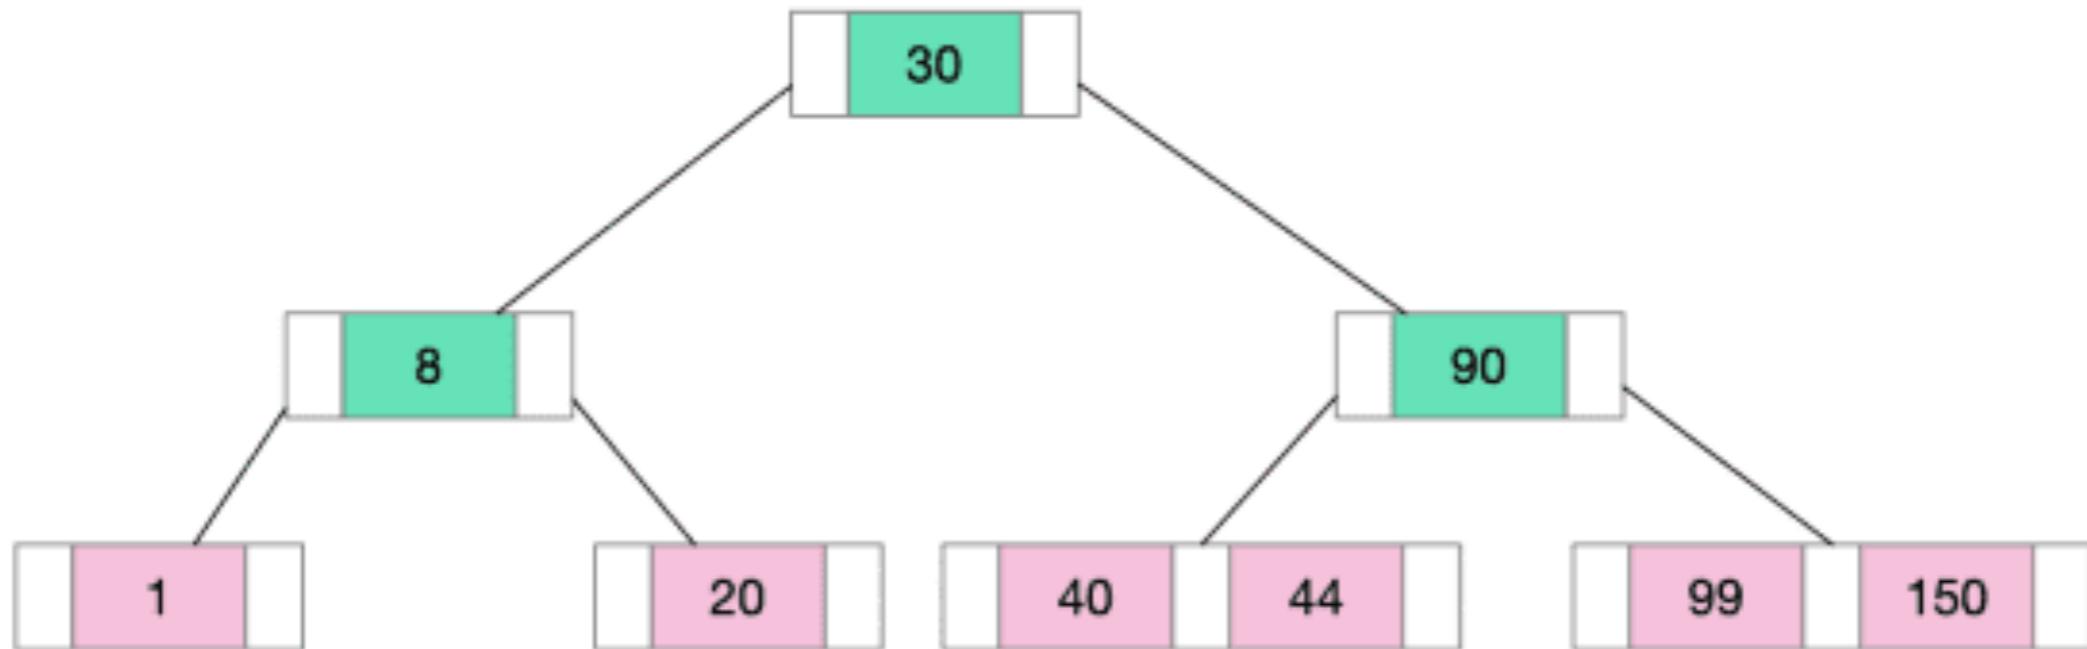
Generalized inverted search tree index (GiST)

Space partitioned GiST (SP-GiST)

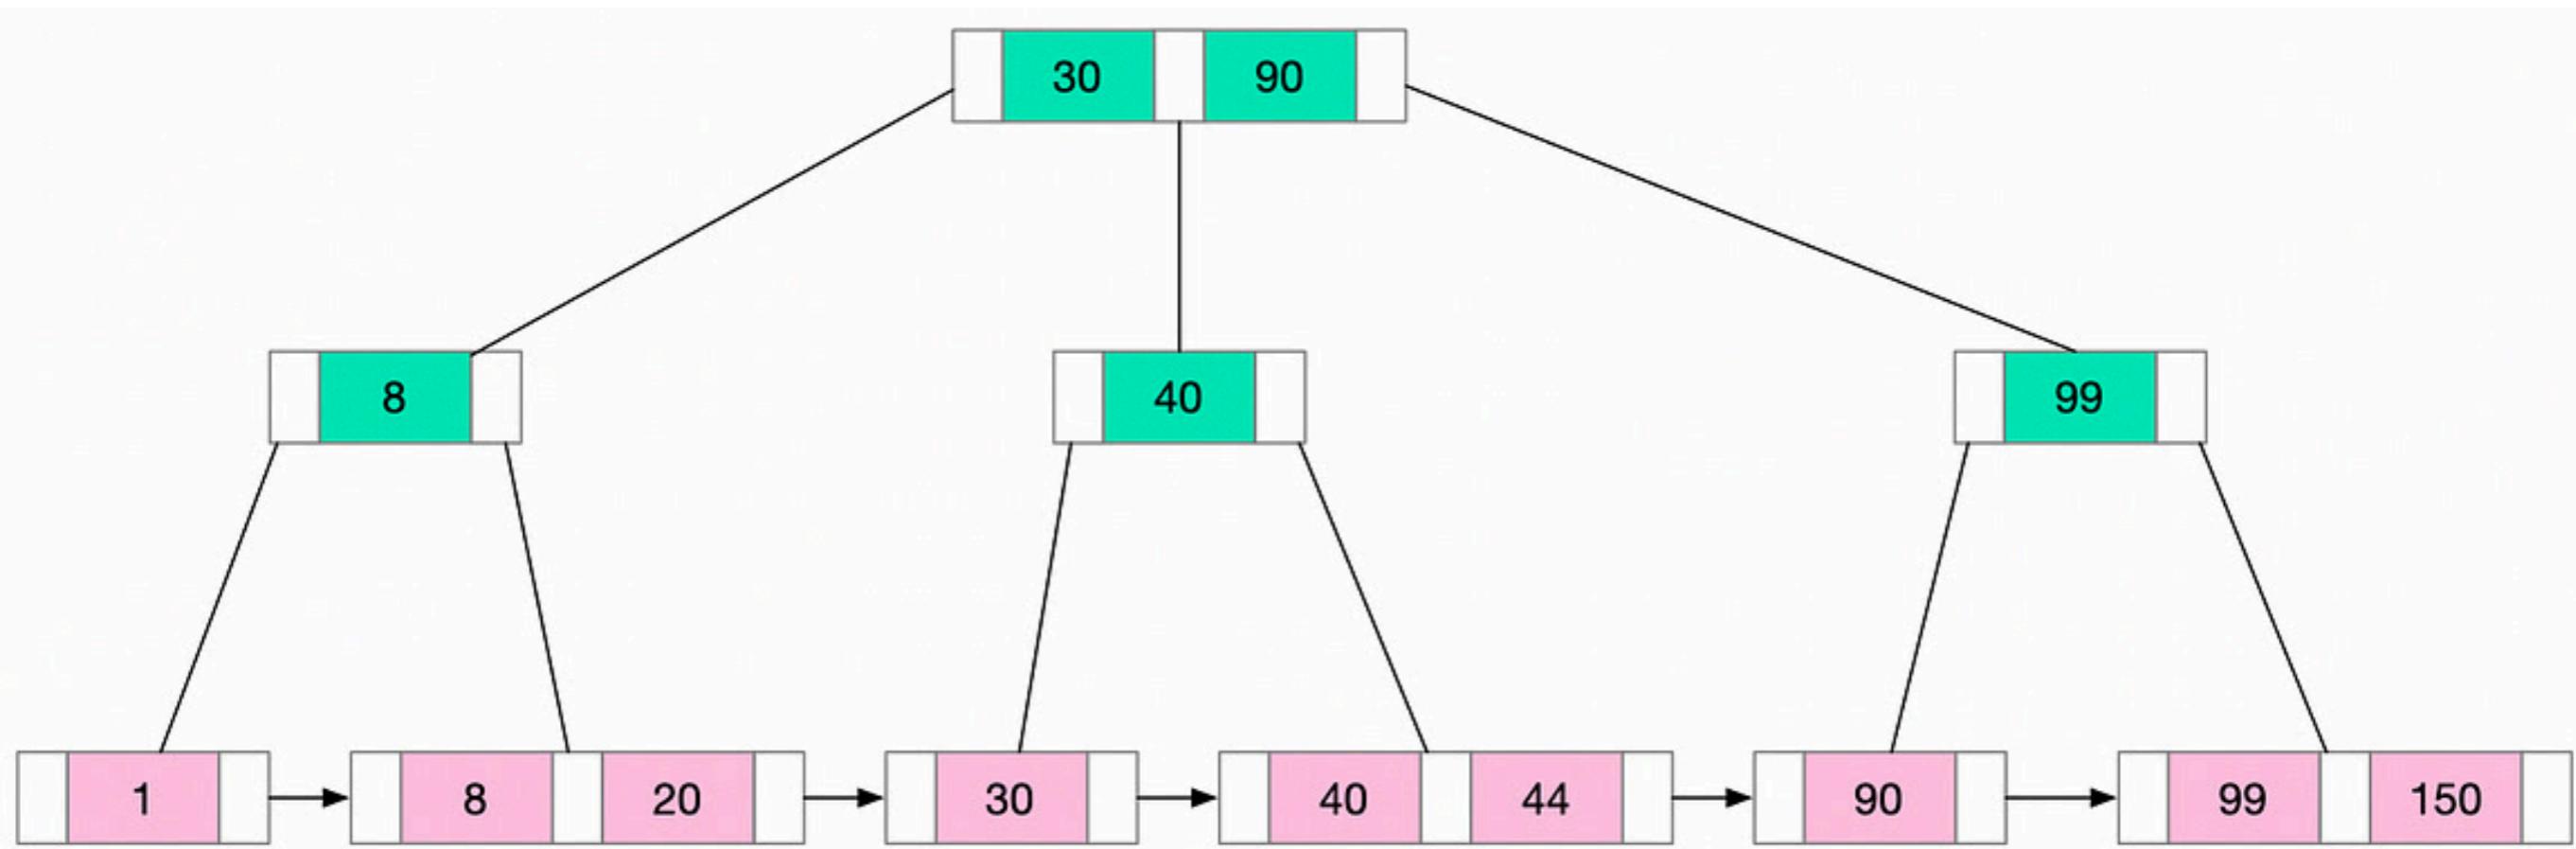
<https://www.postgresql.org/docs/current/indexes-types.html>



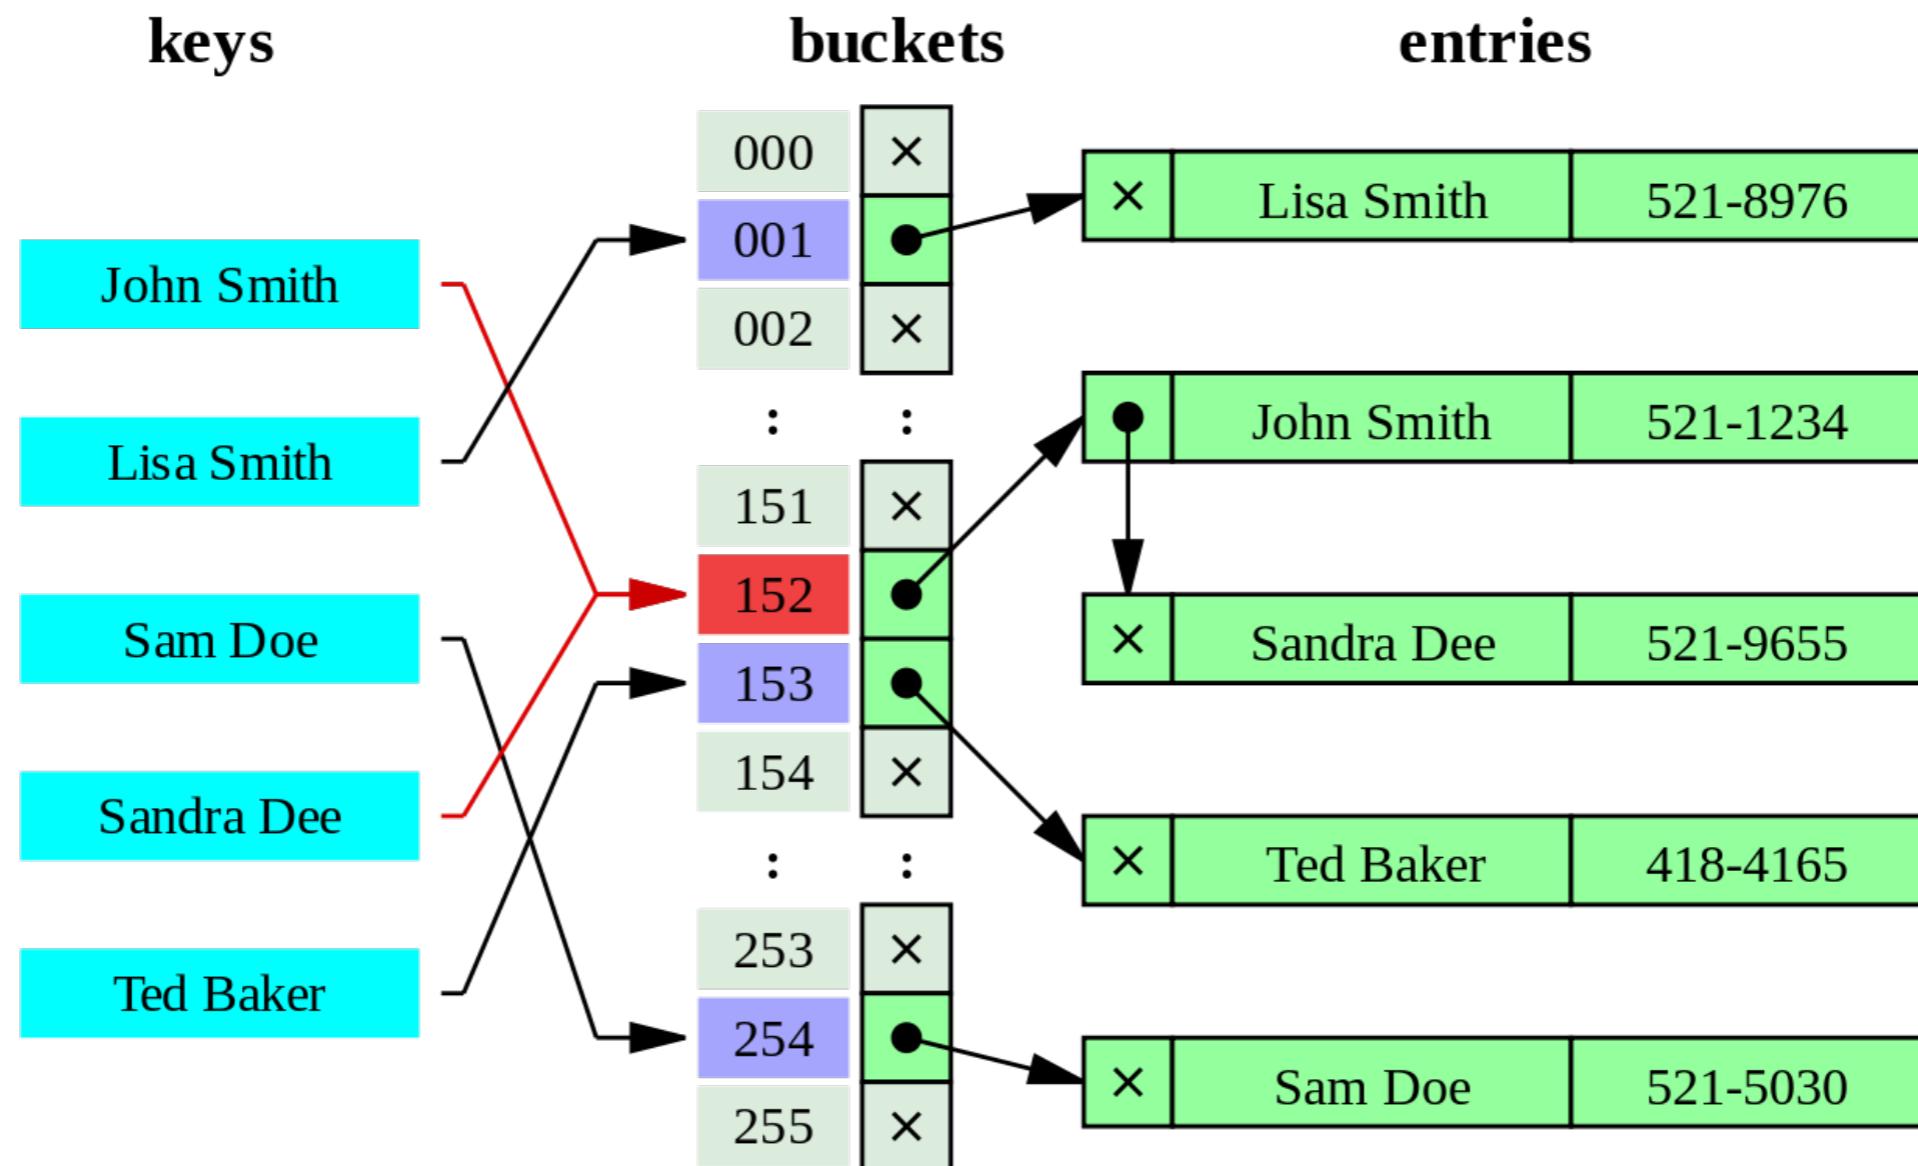
# B-Tree



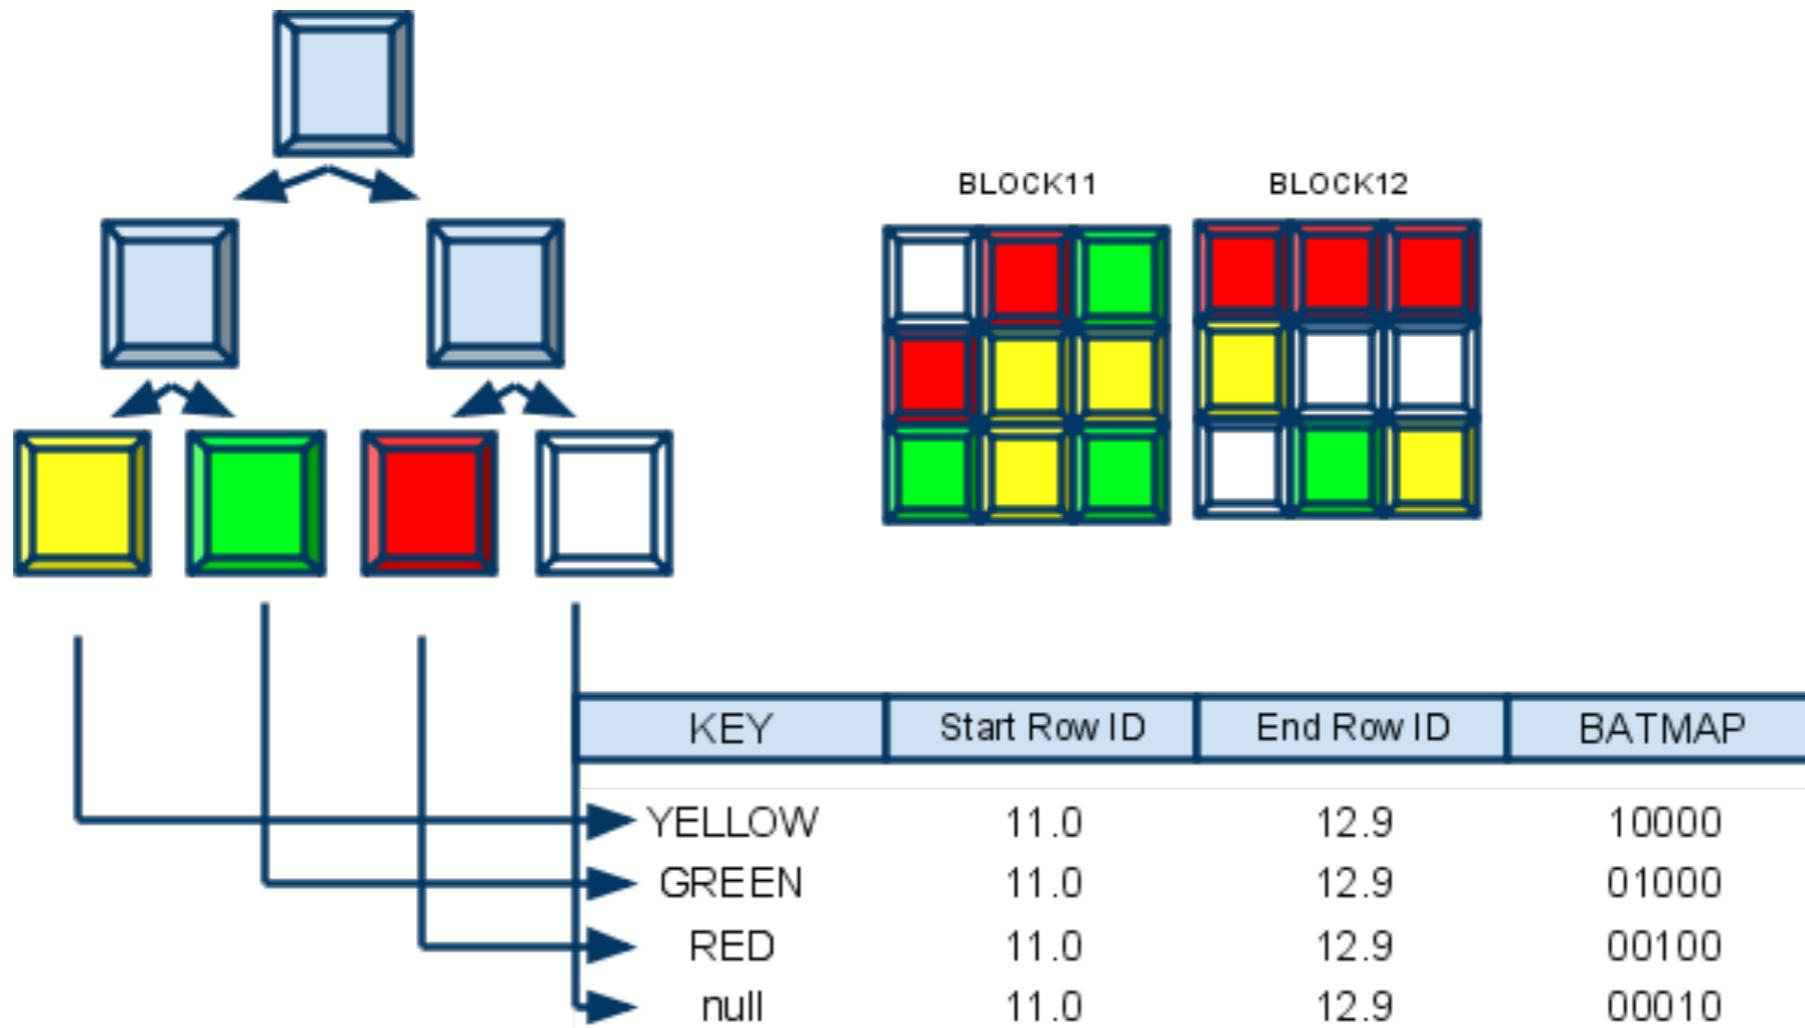
# B-Tree +



# Hash



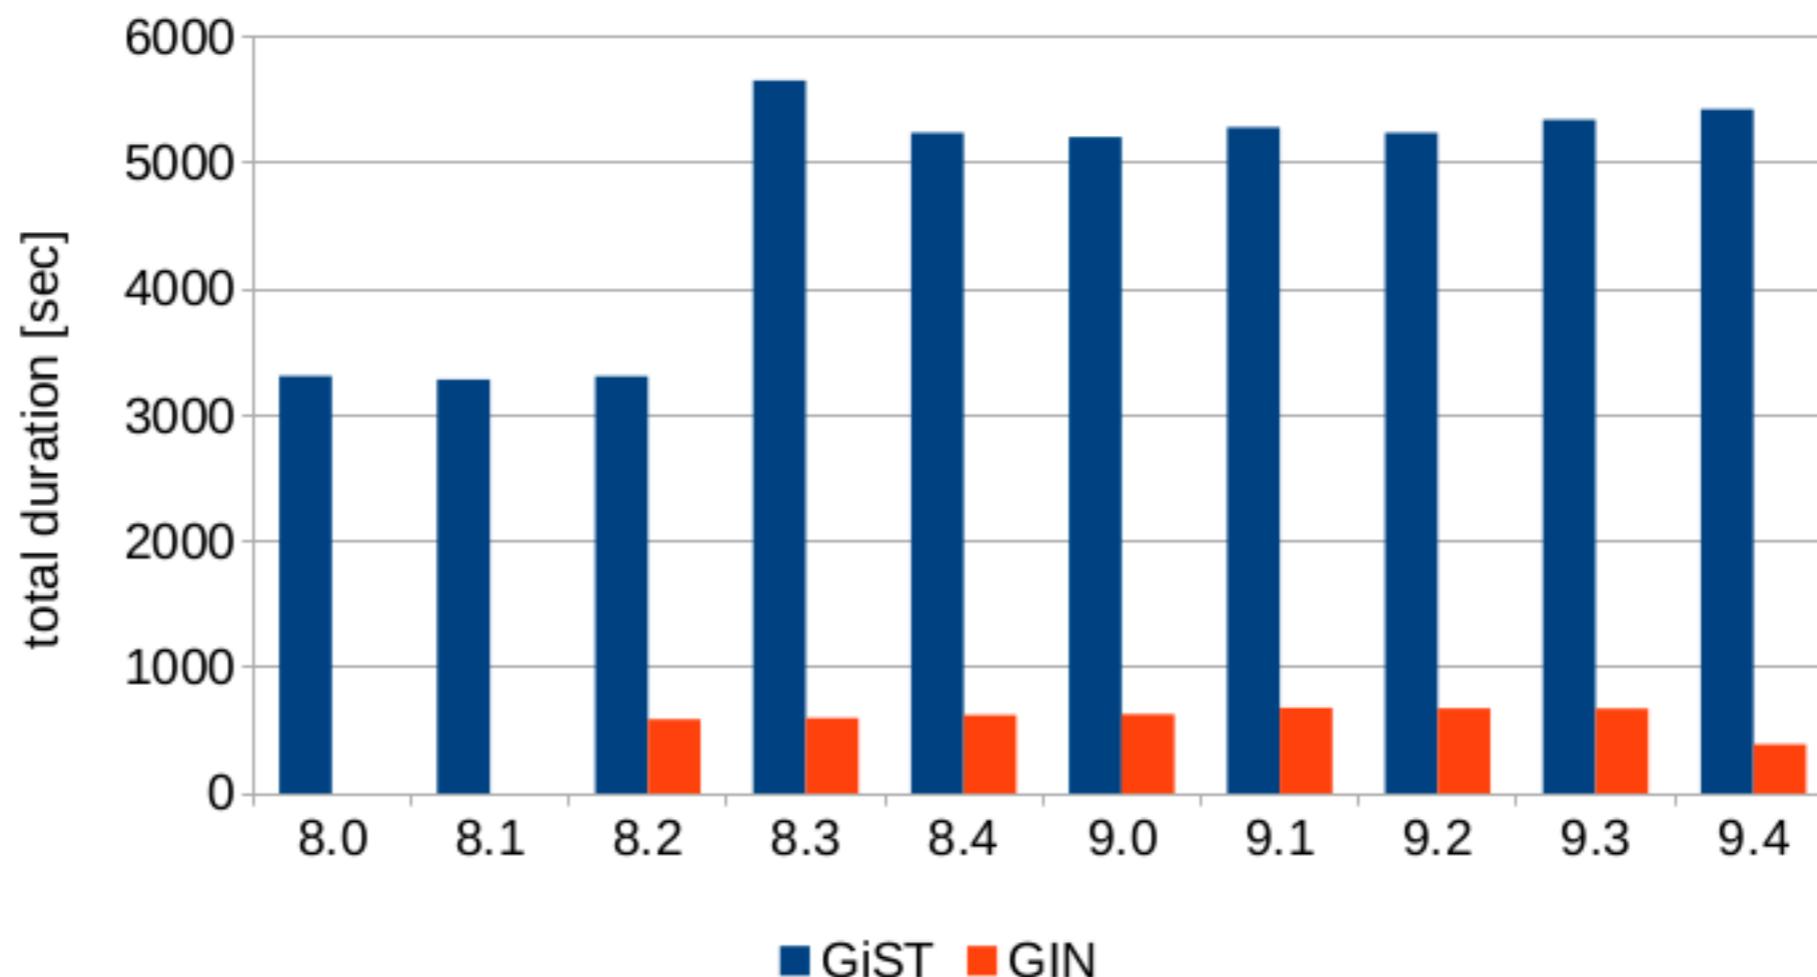
# Bitmap



# GIN vs GiST

## Fulltext benchmark - GiST vs. GIN

33k queries from postgresql.org [TOP 100]



# **Don't over-index**

# **Unused index**



# Tuning Performance



# Tuning performance

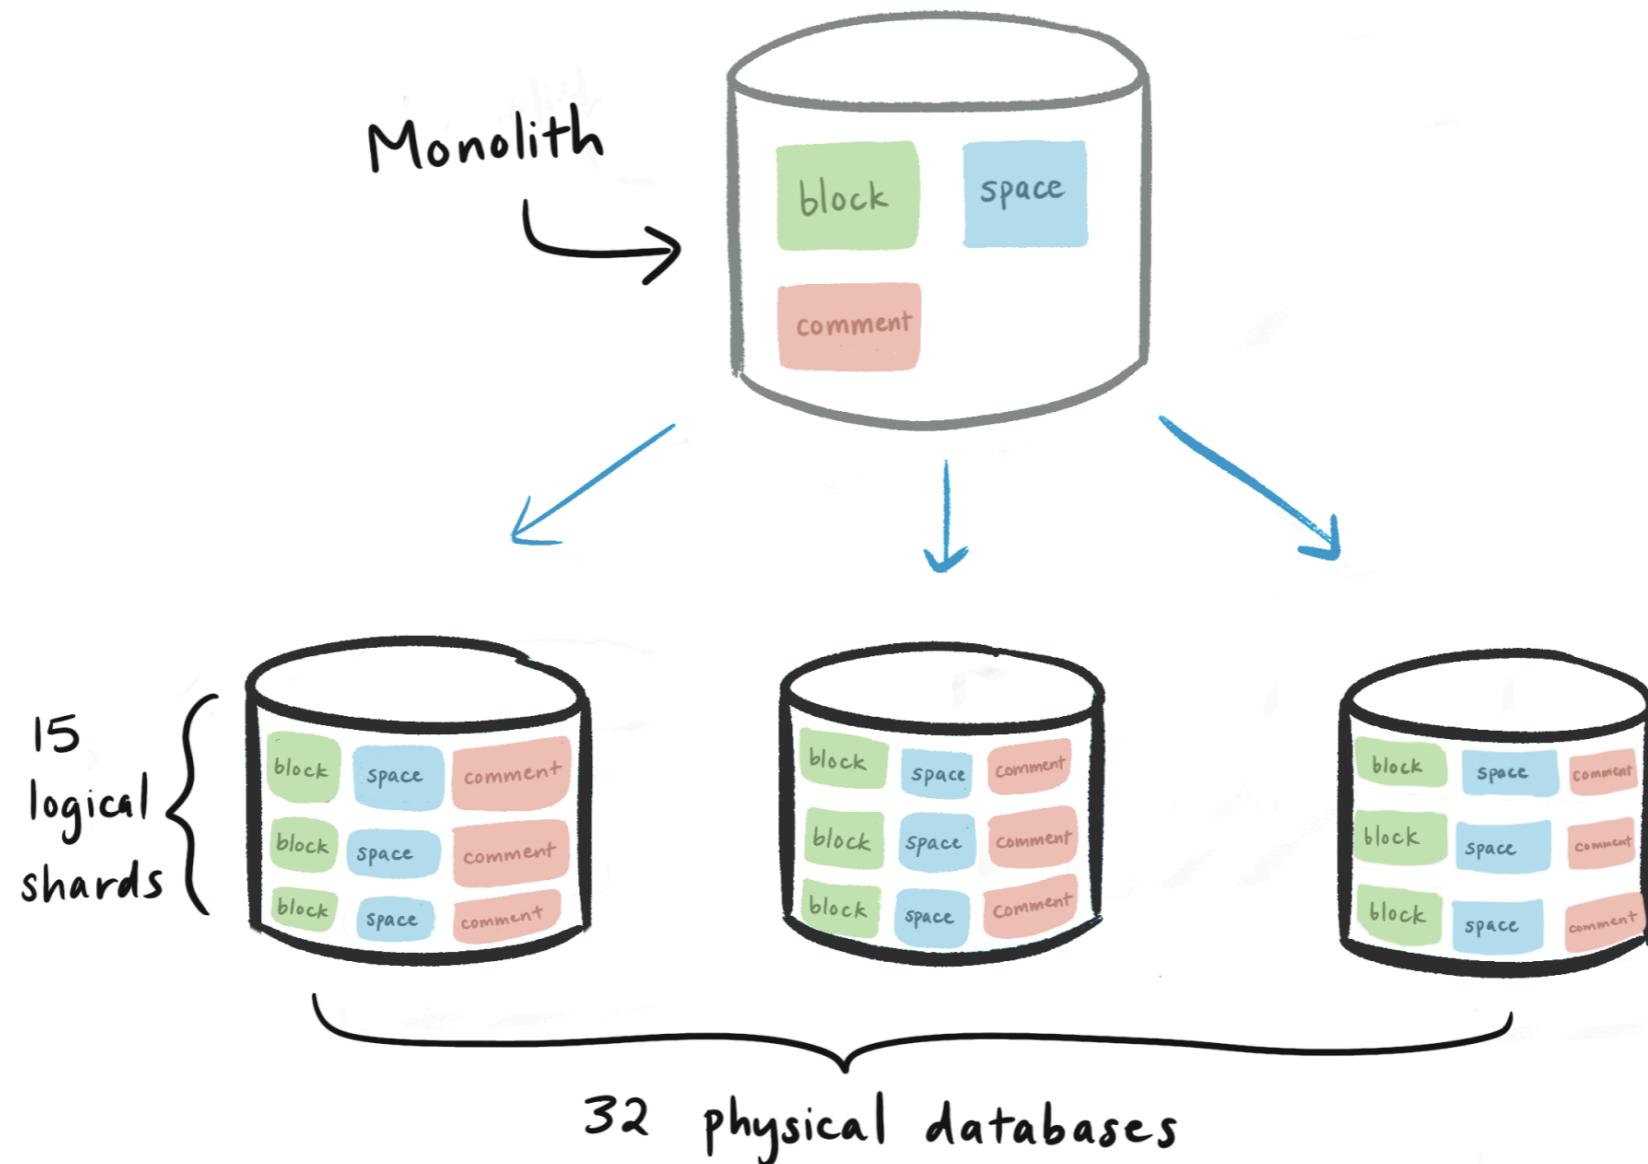
Sharding  
Partition  
Denormalization



# Sharding



# Sharding



<https://www.notion.so/blog/sharding-postgres-at-notion>



# Sharding techniques

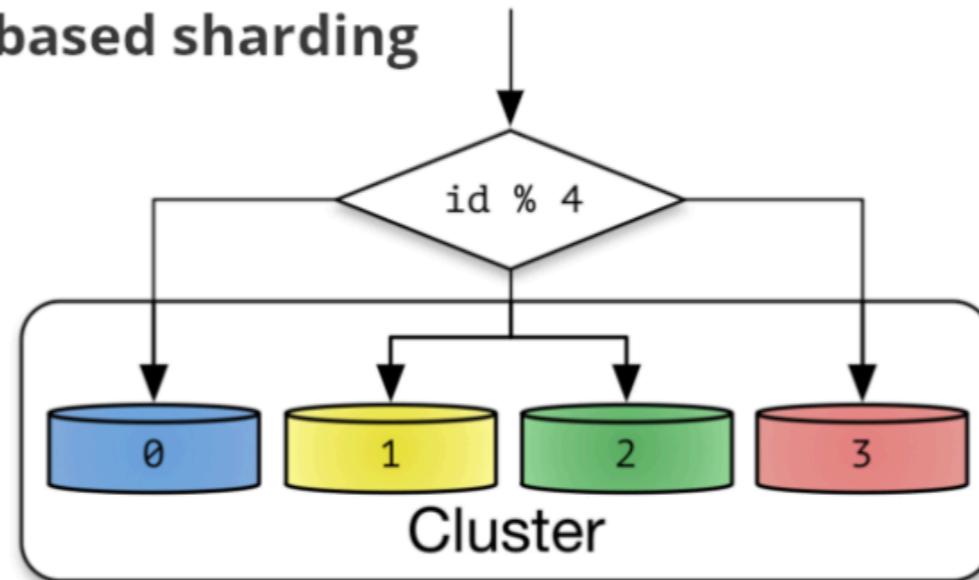
Hash  
Range  
Location  
Key/Directory

<https://www.somkiat.cc/introduction-database-sharding/>

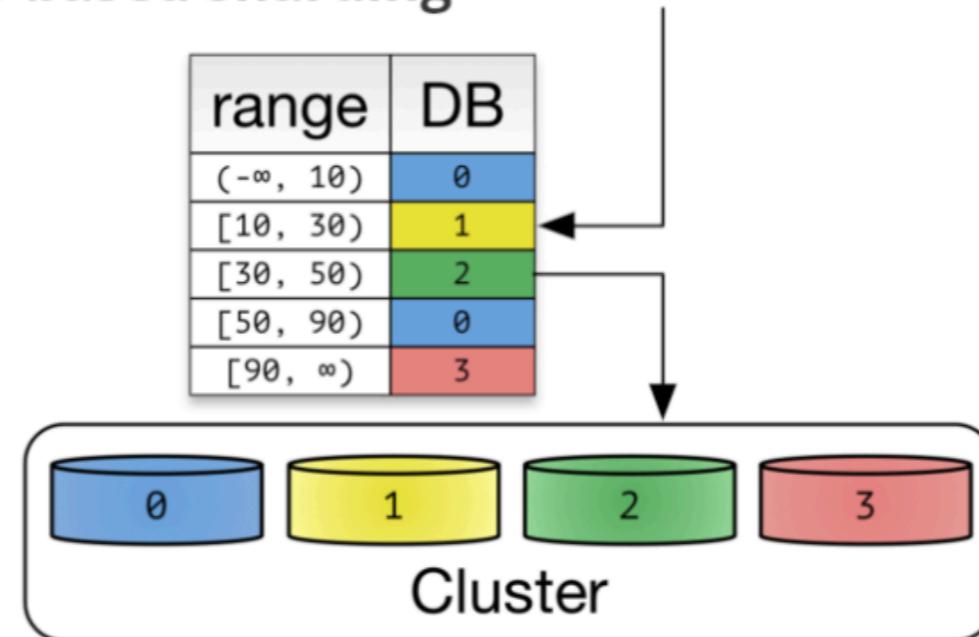


# Sharding

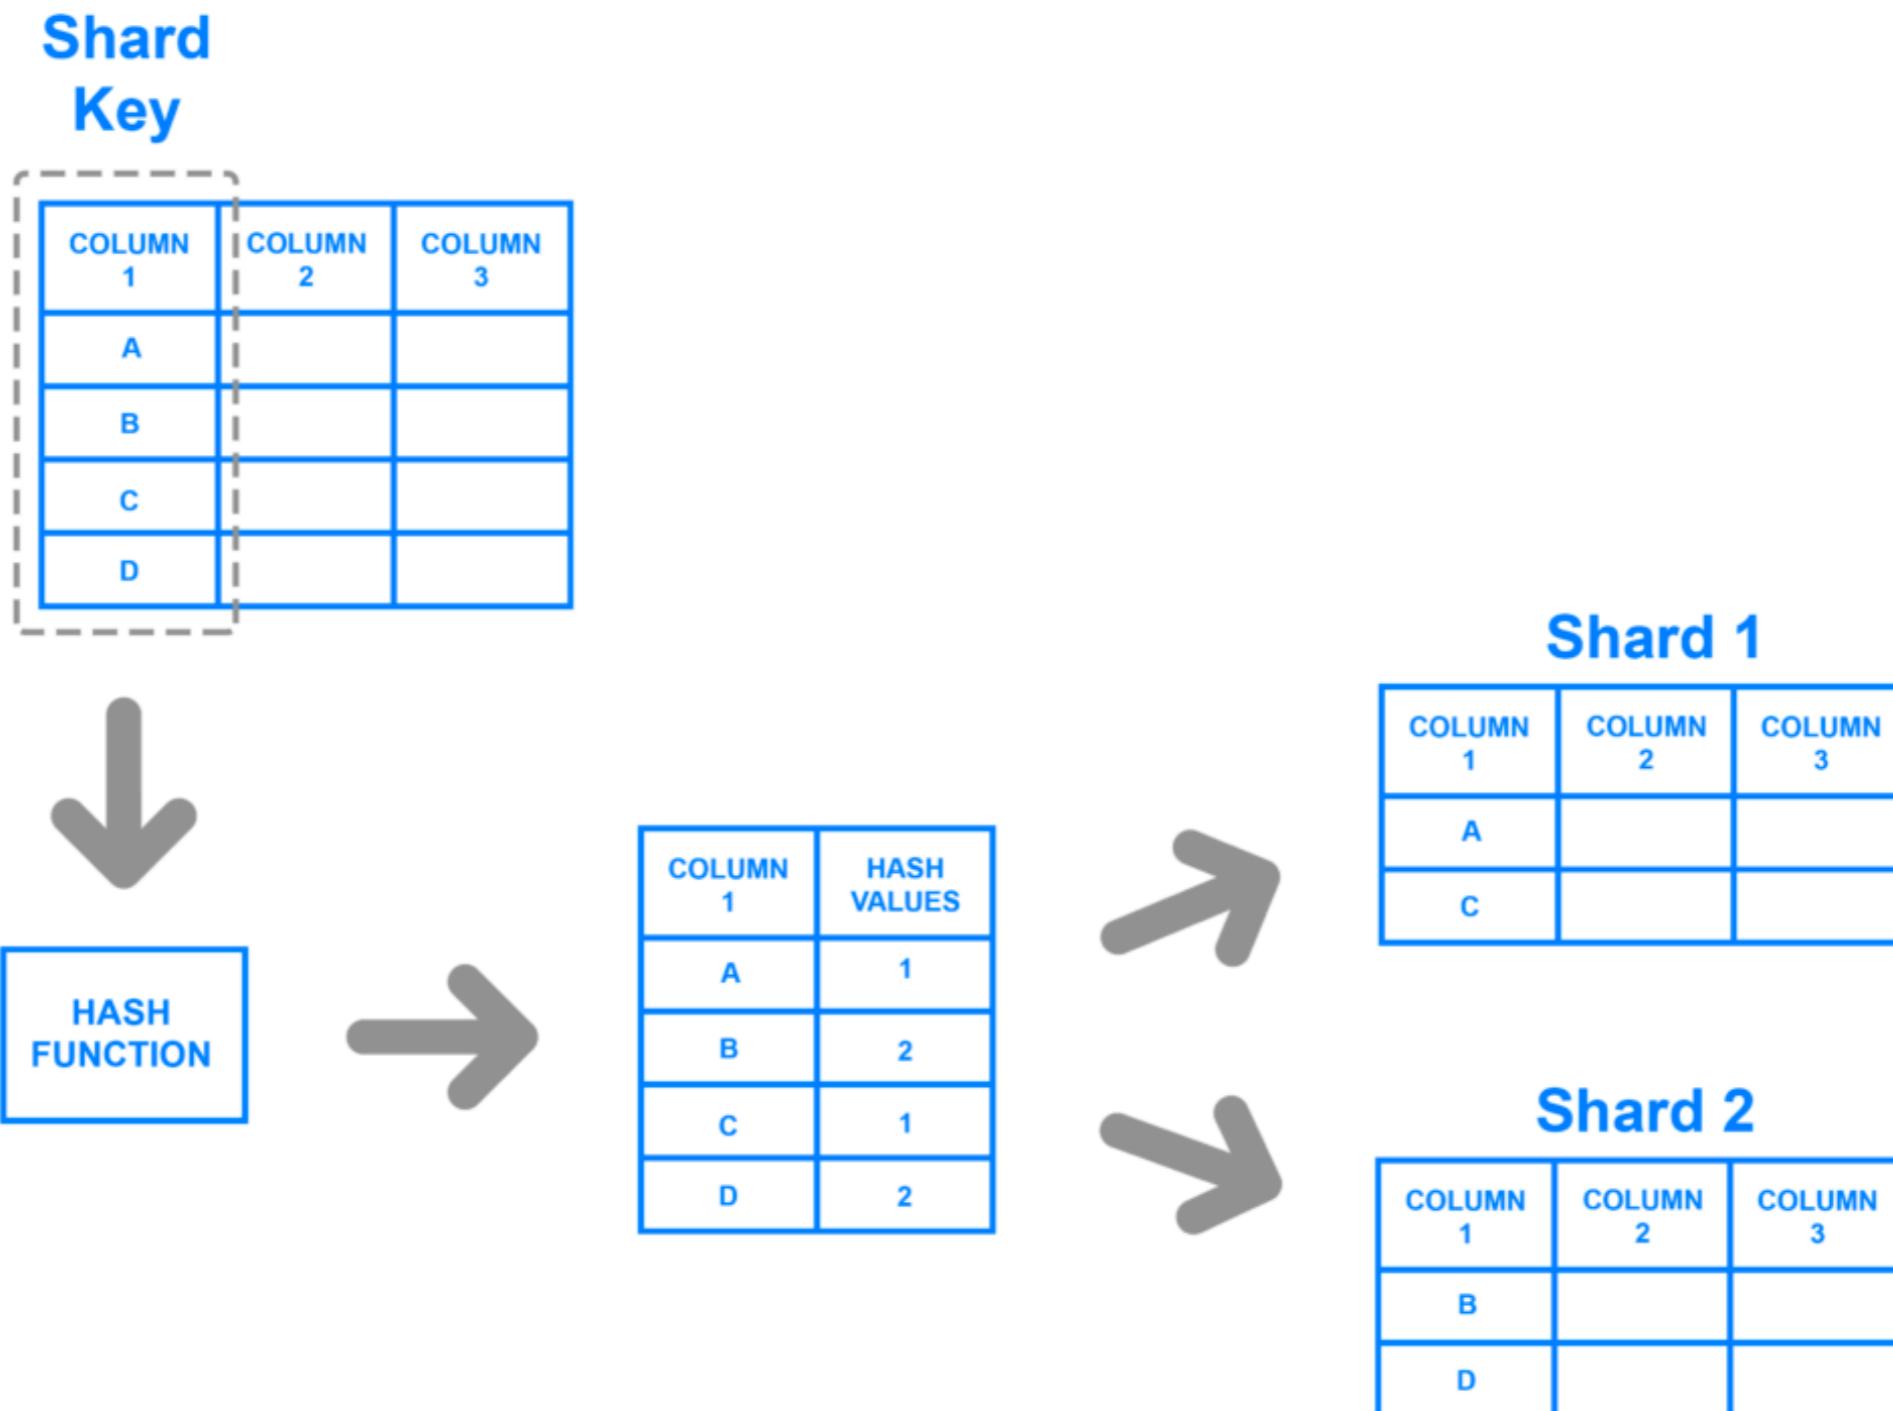
(a) Hash-based sharding



(b) Range-based sharding



# Shard Key



# Range

| PRODUCT     | PRICE |
|-------------|-------|
| WIDGET      | \$118 |
| GIZMO       | \$88  |
| TRINKET     | \$37  |
| THINGAMAJIG | \$18  |
| DOODAD      | \$60  |
| TCHOTCHKE   | \$999 |



(\$0-\$49.99)

| PRODUCT     | PRICE |
|-------------|-------|
| TRINKET     | \$37  |
| THINGAMAJIG | \$18  |



(\$50-\$99.99)

| PRODUCT | PRICE |
|---------|-------|
| GIZMO   | \$88  |
| DOODAD  | \$60  |



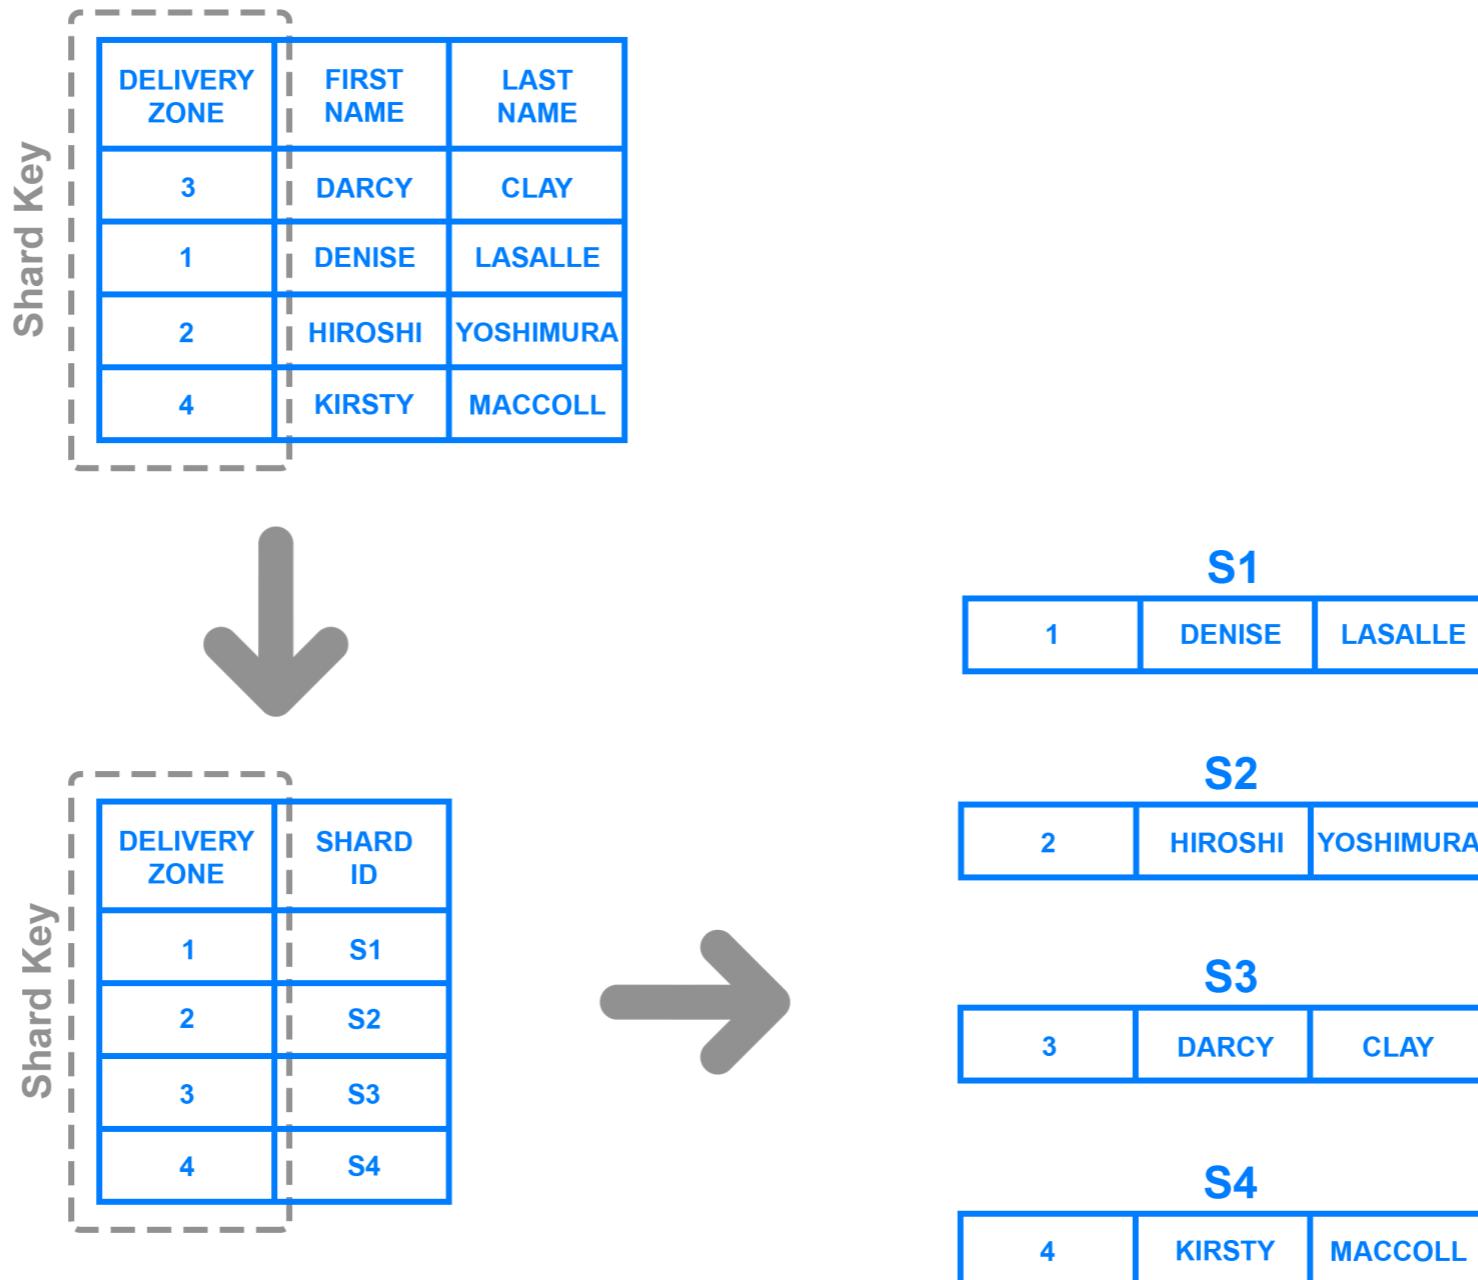
(\$100+)

| PRODUCT   | PRICE |
|-----------|-------|
| WIDGET    | \$118 |
| TCHOTCHKE | \$999 |



# Directory

## Pre-Sharded Table



# Sharding

Shard 1

| Stores |                |
|--------|----------------|
| id     | name           |
| 1      | my book store  |
| 5      | my other store |

| Products |      |          |
|----------|------|----------|
| id       | name | store_id |
| 1        | foo  | 1        |
| 2        | bar  | 1        |
| 3        | baz  | 1        |

| Purchases |            |          |       |
|-----------|------------|----------|-------|
| id        | product_id | store_id | price |
| 1         | 2          | 1        | 1000  |
| 2         | 1          | 1        | 1200  |
| 3         | 3          | 1        | 1199  |

Shard 2

| Stores |               |
|--------|---------------|
| id     | name          |
| 2      | my sock store |
| 6      | old things    |

| Products |           |          |
|----------|-----------|----------|
| id       | name      | store_id |
| 33       | new socks | 2        |
| 34       | old socks | 6        |
| 35       | old tie   | 6        |

| Purchases |            |          |       |
|-----------|------------|----------|-------|
| id        | product_id | store_id | price |
| 102       | 35         | 6        | 600   |
| 43        | 33         | 2        | 800   |

<https://www.citusdata.com/blog/2016/08/10/sharding-for-a-multi-tenant-app-with-postgres/>



# Partition

<https://hakibenita.com/postgresql-unused-index-size>



# Partition

Table partitioning refers to **splitting** what is logically one large table into **smaller physical pieces**

Reduce index size



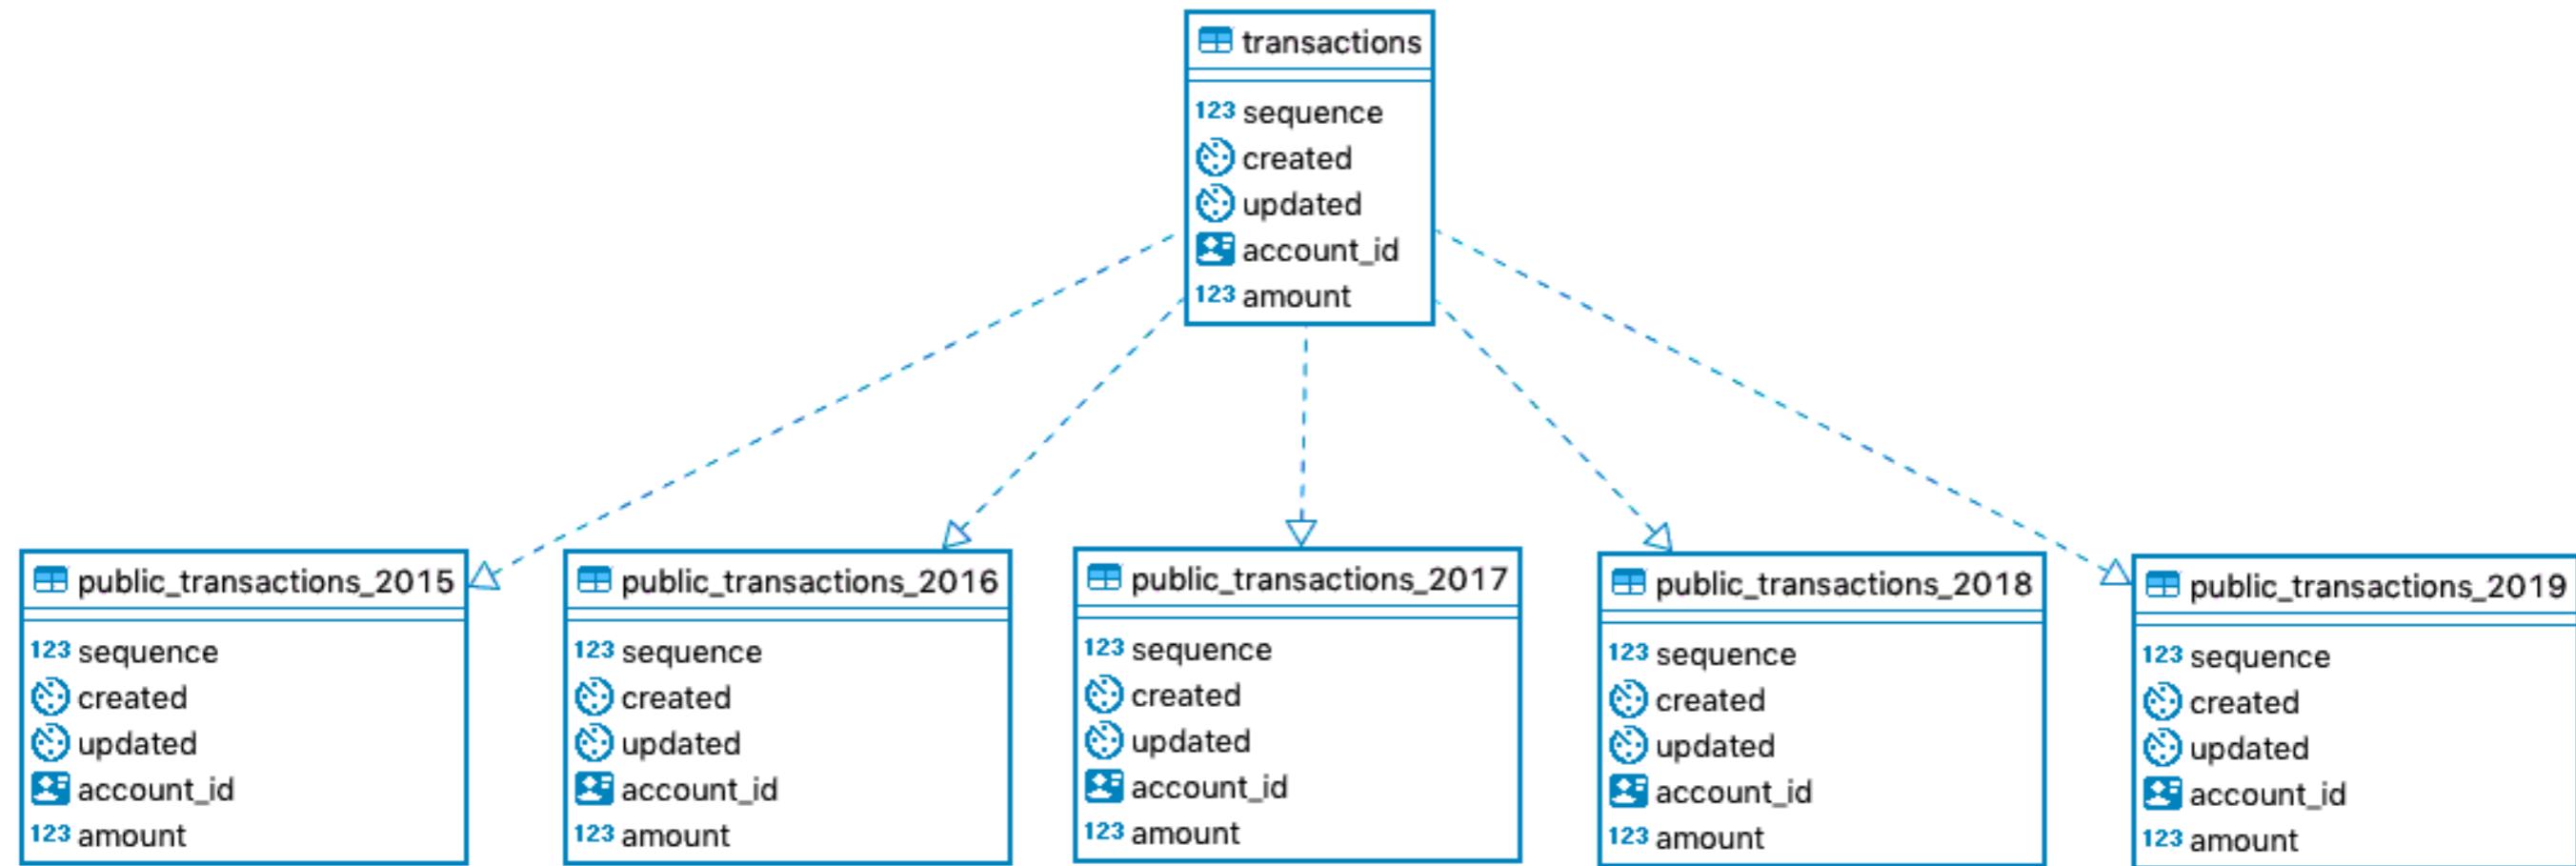
# Partition

Horizontal partitioning  
Vertical partitioning



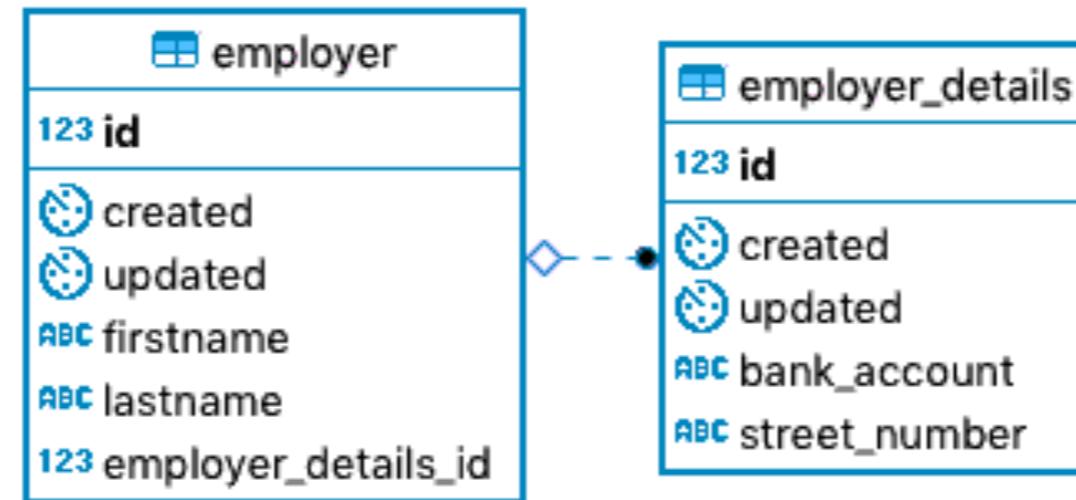
# Horizontal partitioning

Putting different rows into different table



# Vertical partitioning

## Horizontal partitioning



# Types Horizontal partitioning

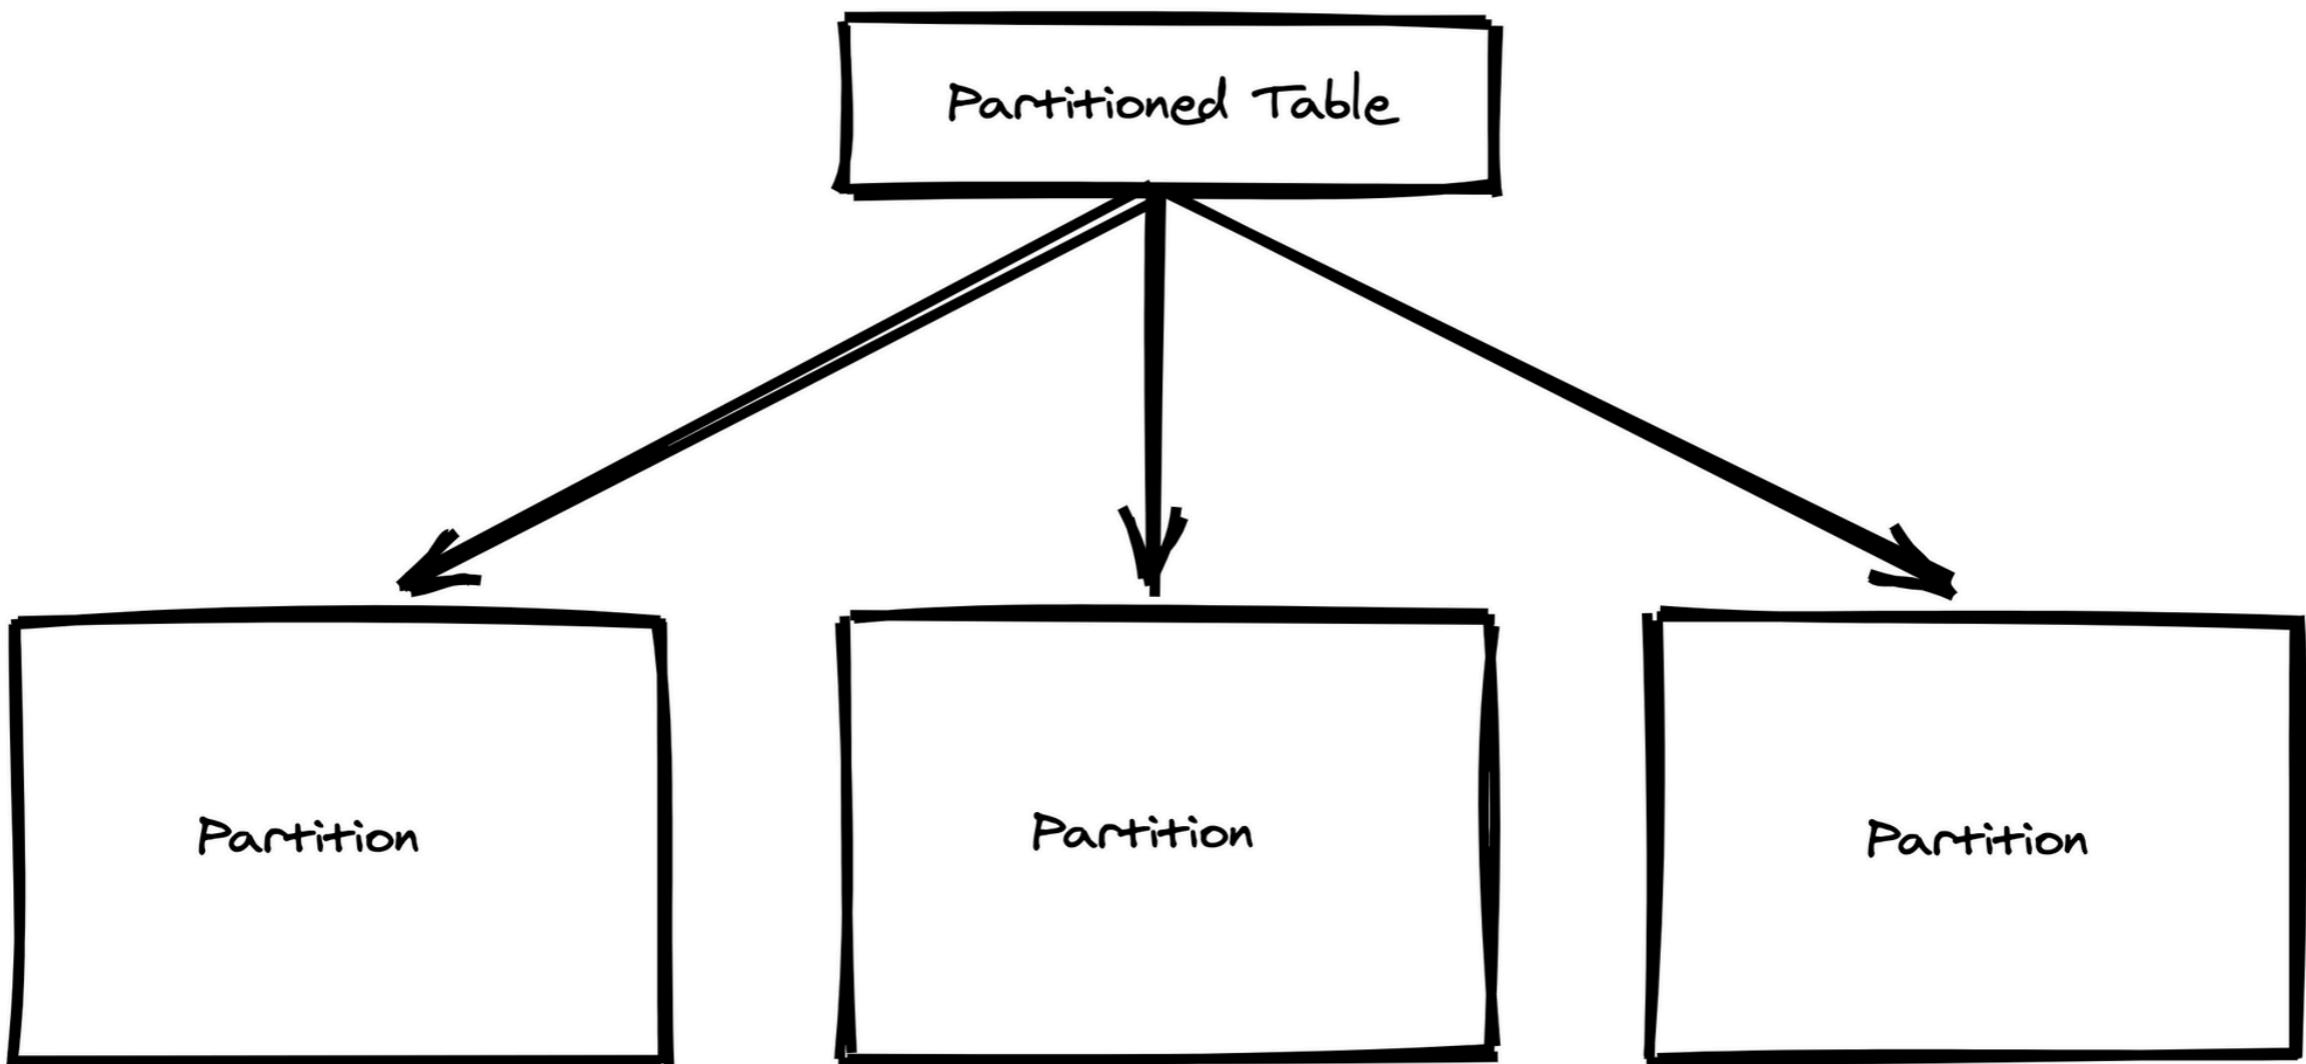
Range  
List



# Example



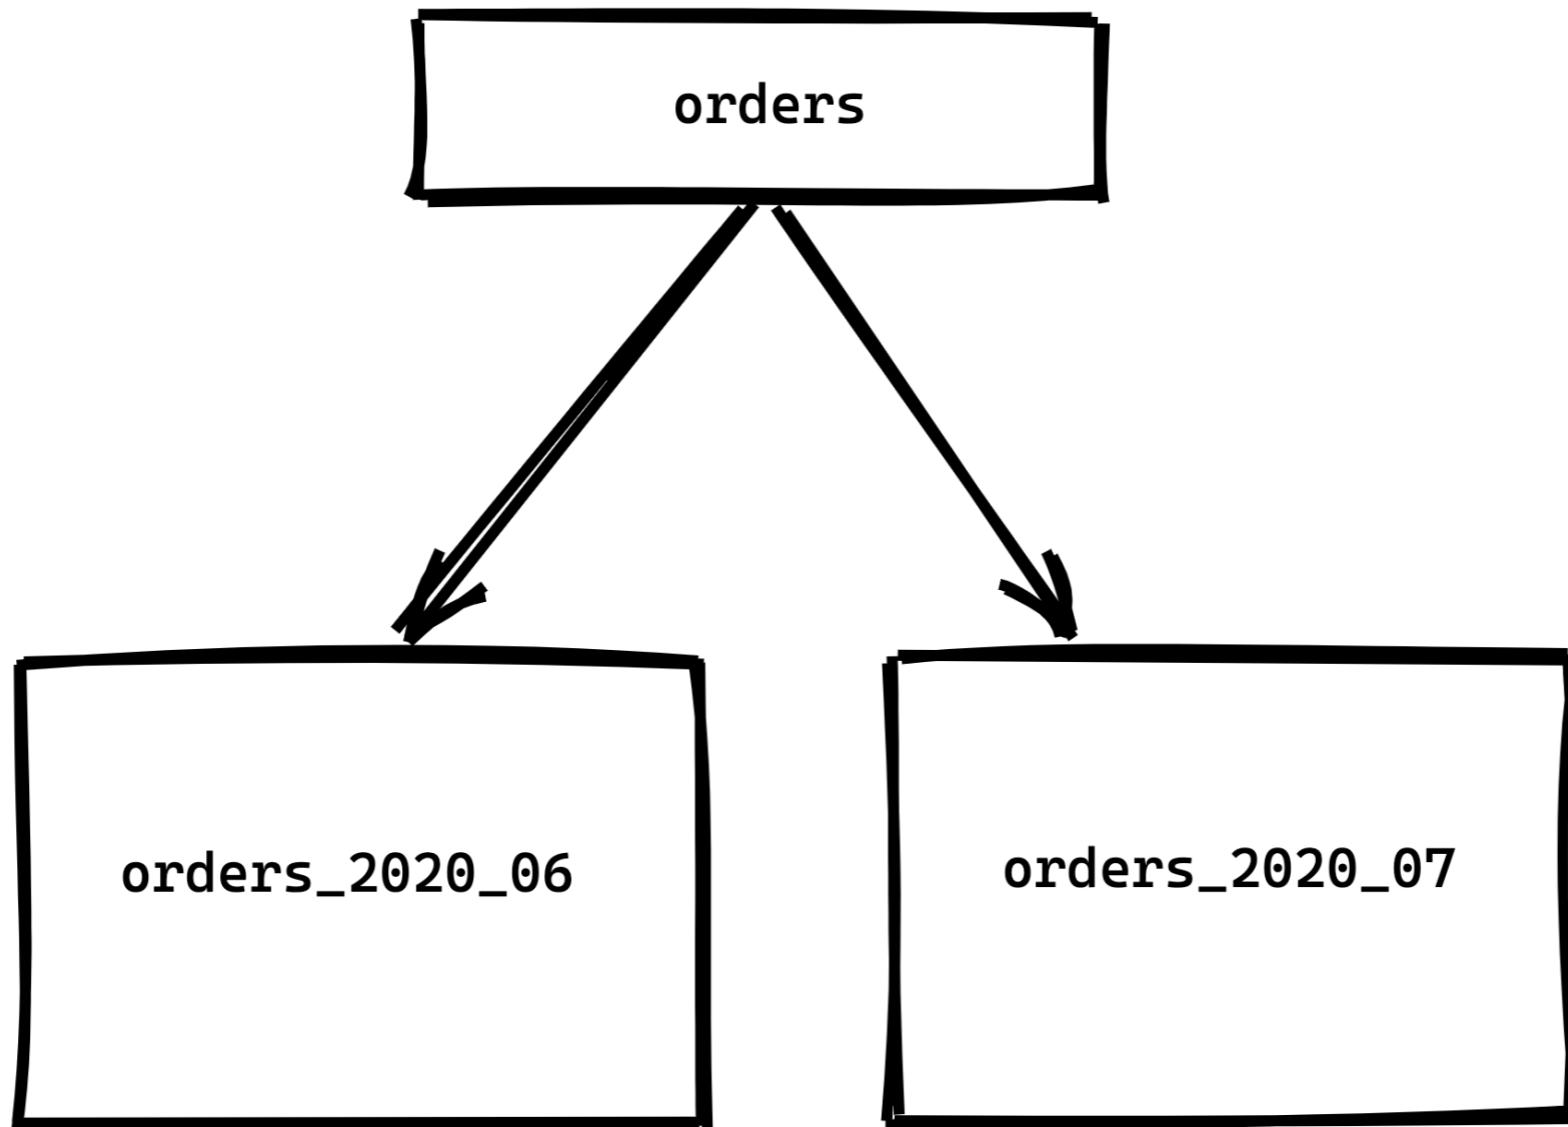
# Partition tables



<https://chriserwin.com/table-partitioning/>



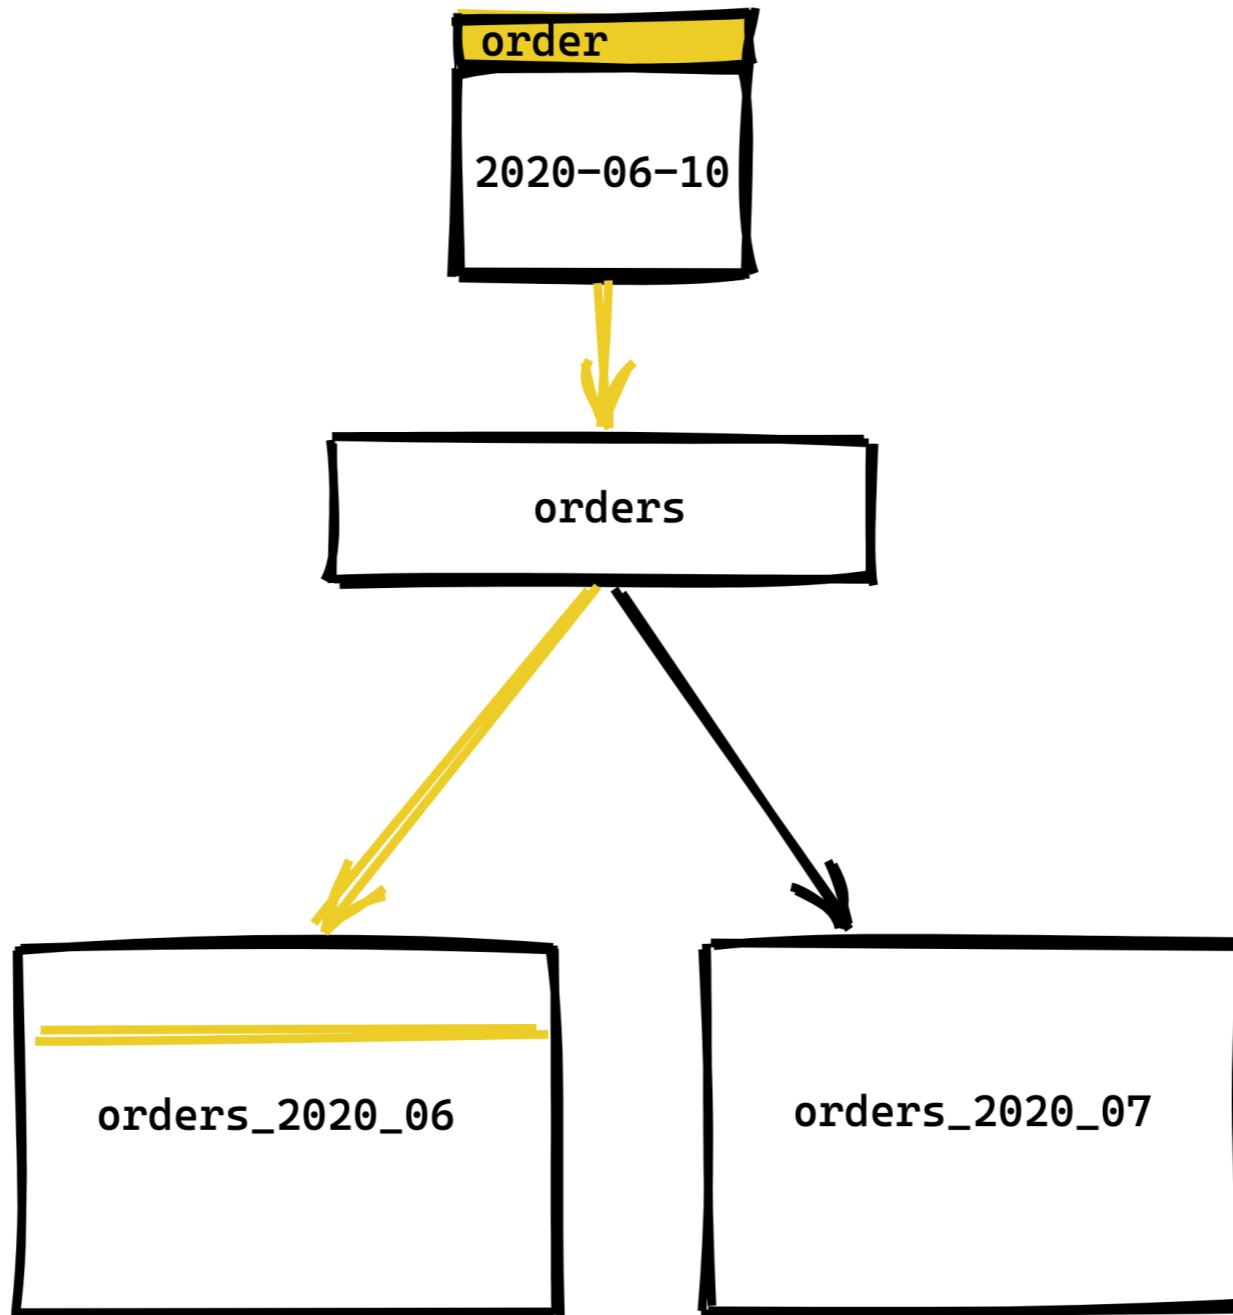
# Orders table



<https://chriserwin.com/table-partitioning/>



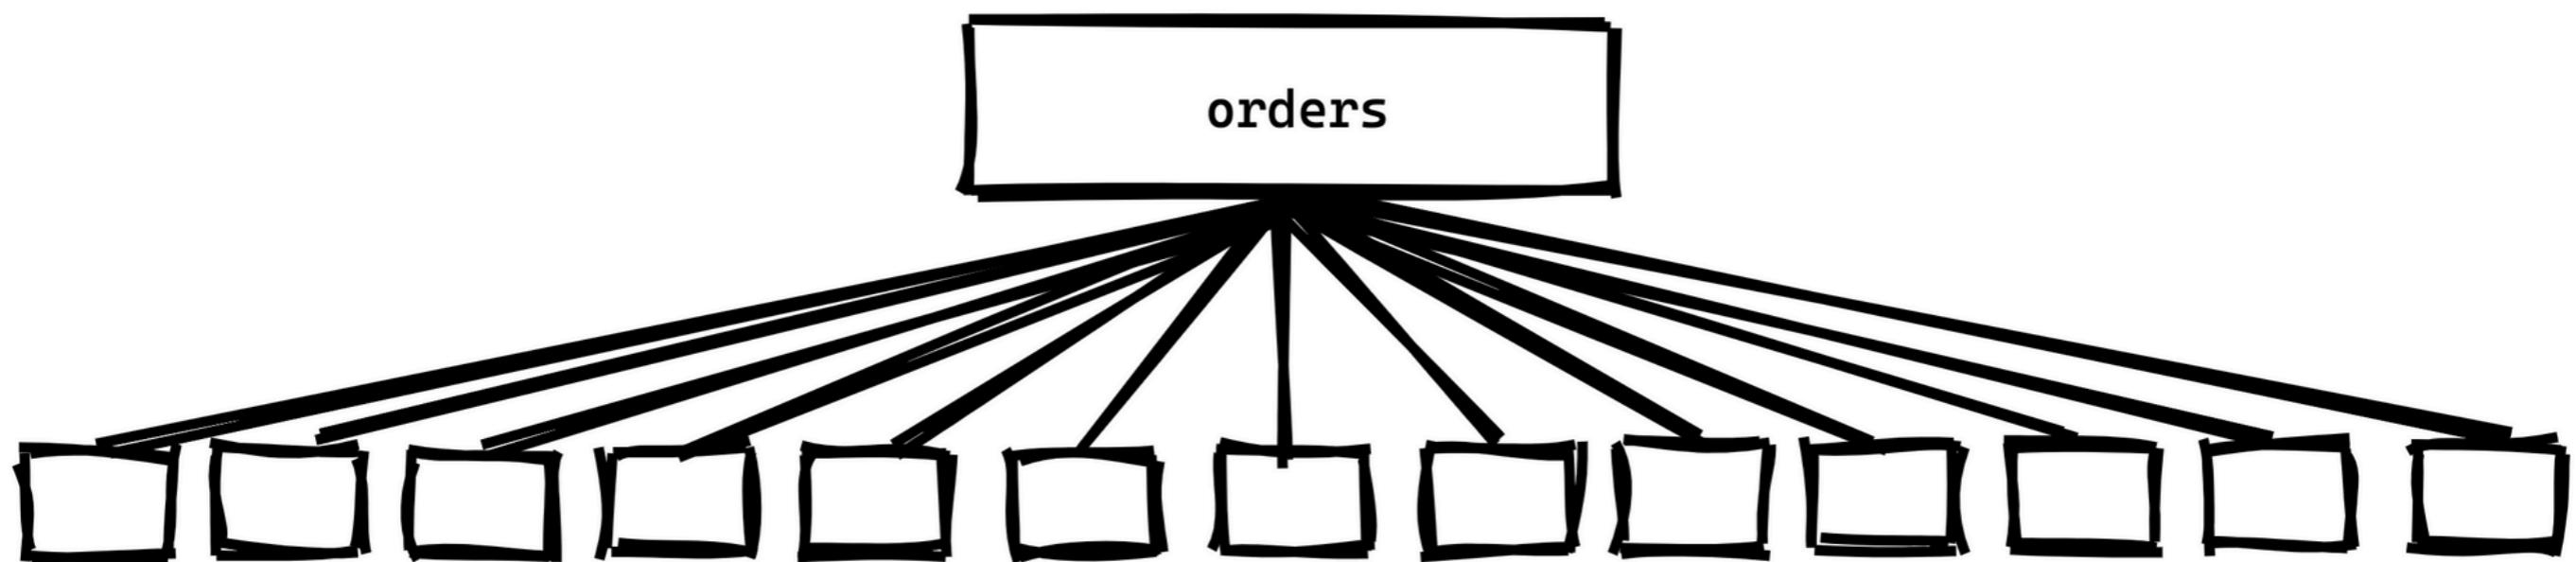
# Insert data into orders



<https://chriserwin.com/table-partitioning/>



• • •



<https://chriserwin.com/table-partitioning/>



# Workshop



# Denormalization



# Denormalization techniques

Redundant column/Pre-joining tables

Table splitting (horizontal/vertical)

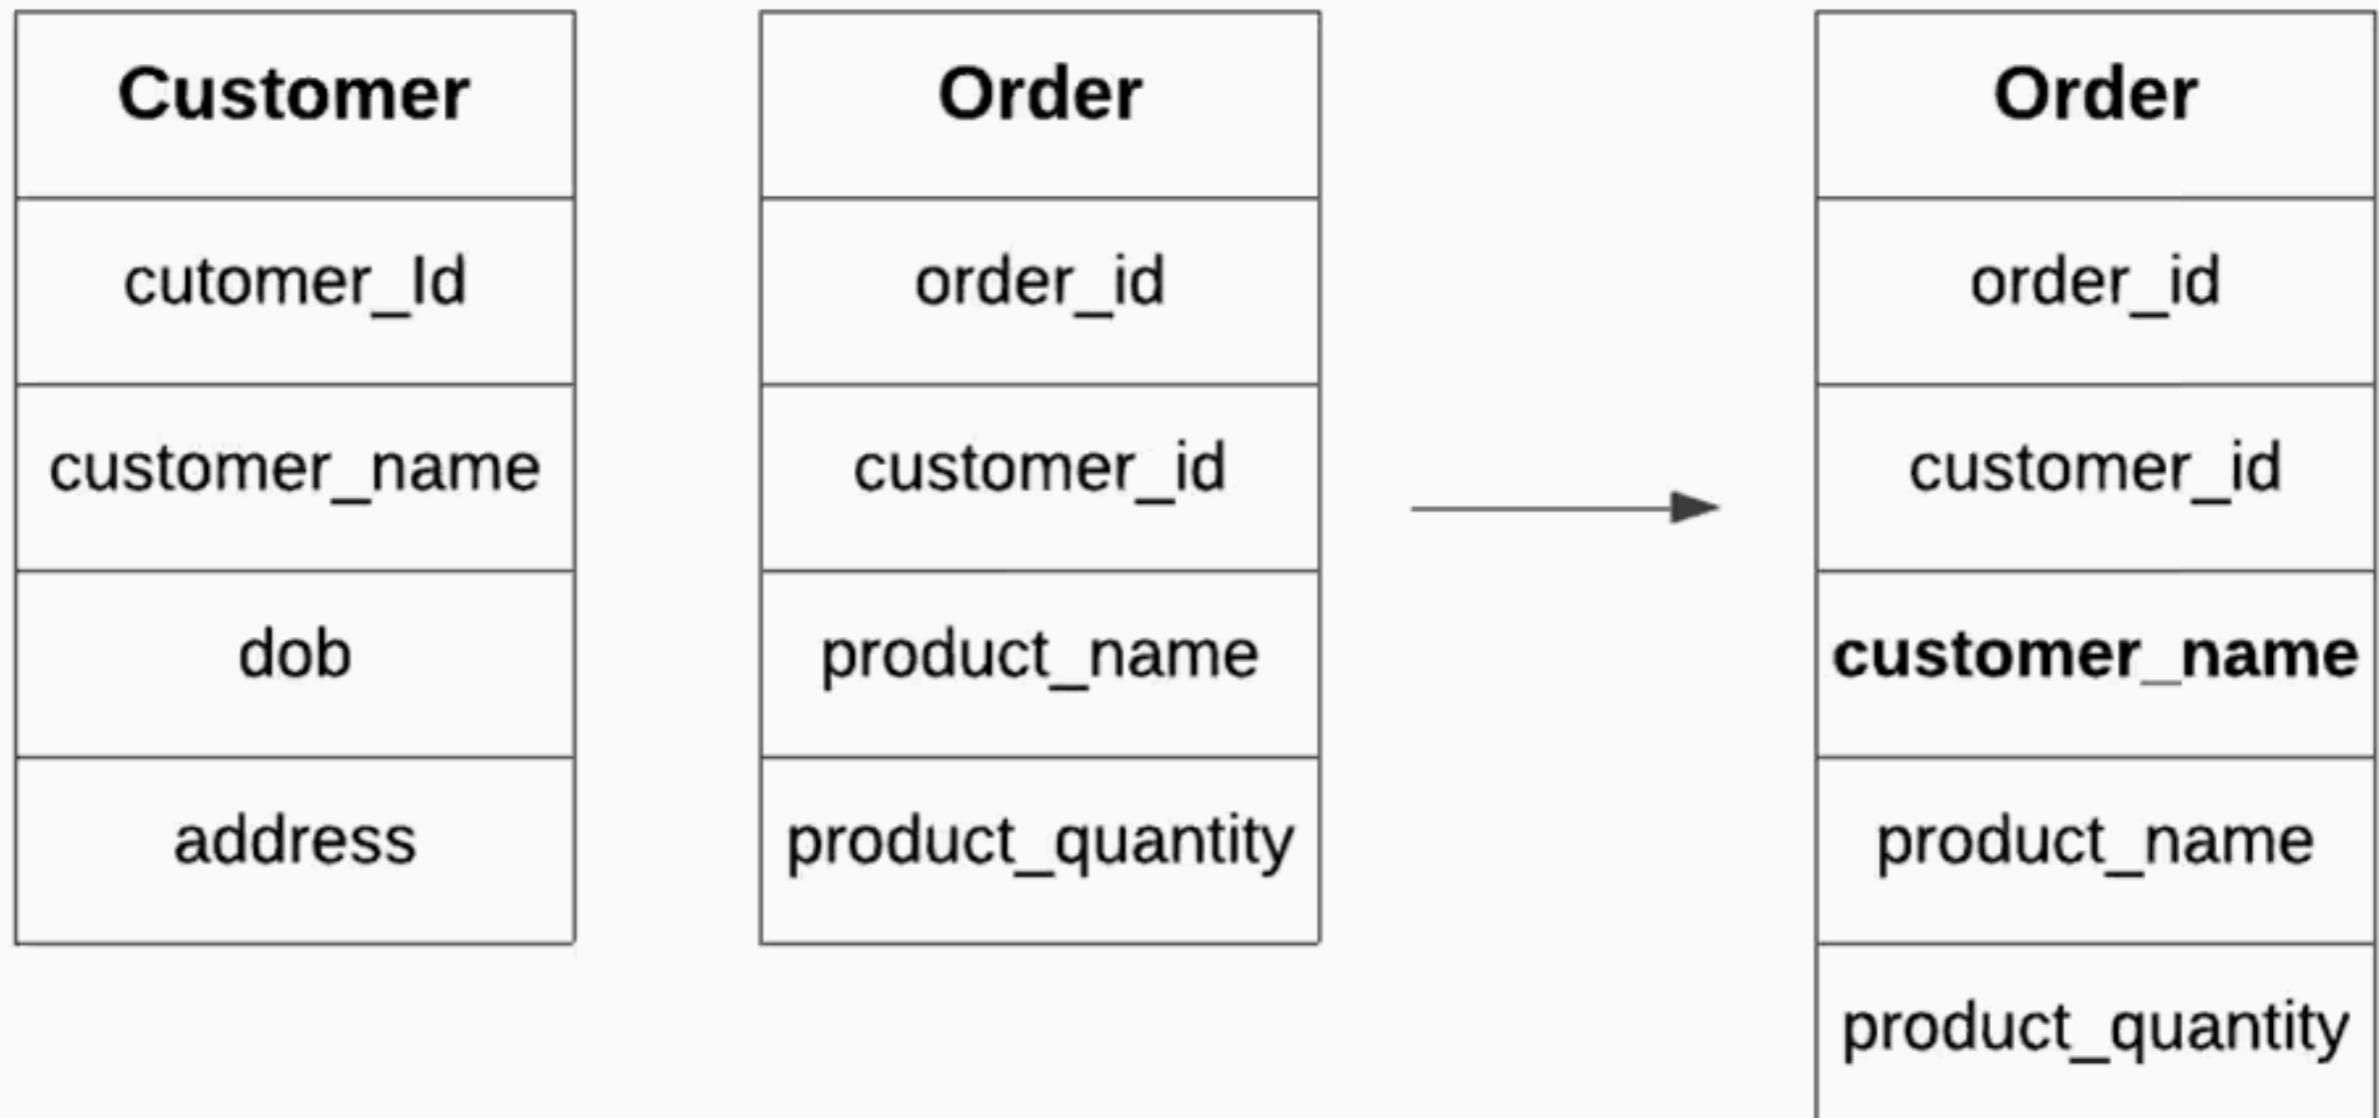
Add derived columns

Mirrored tables

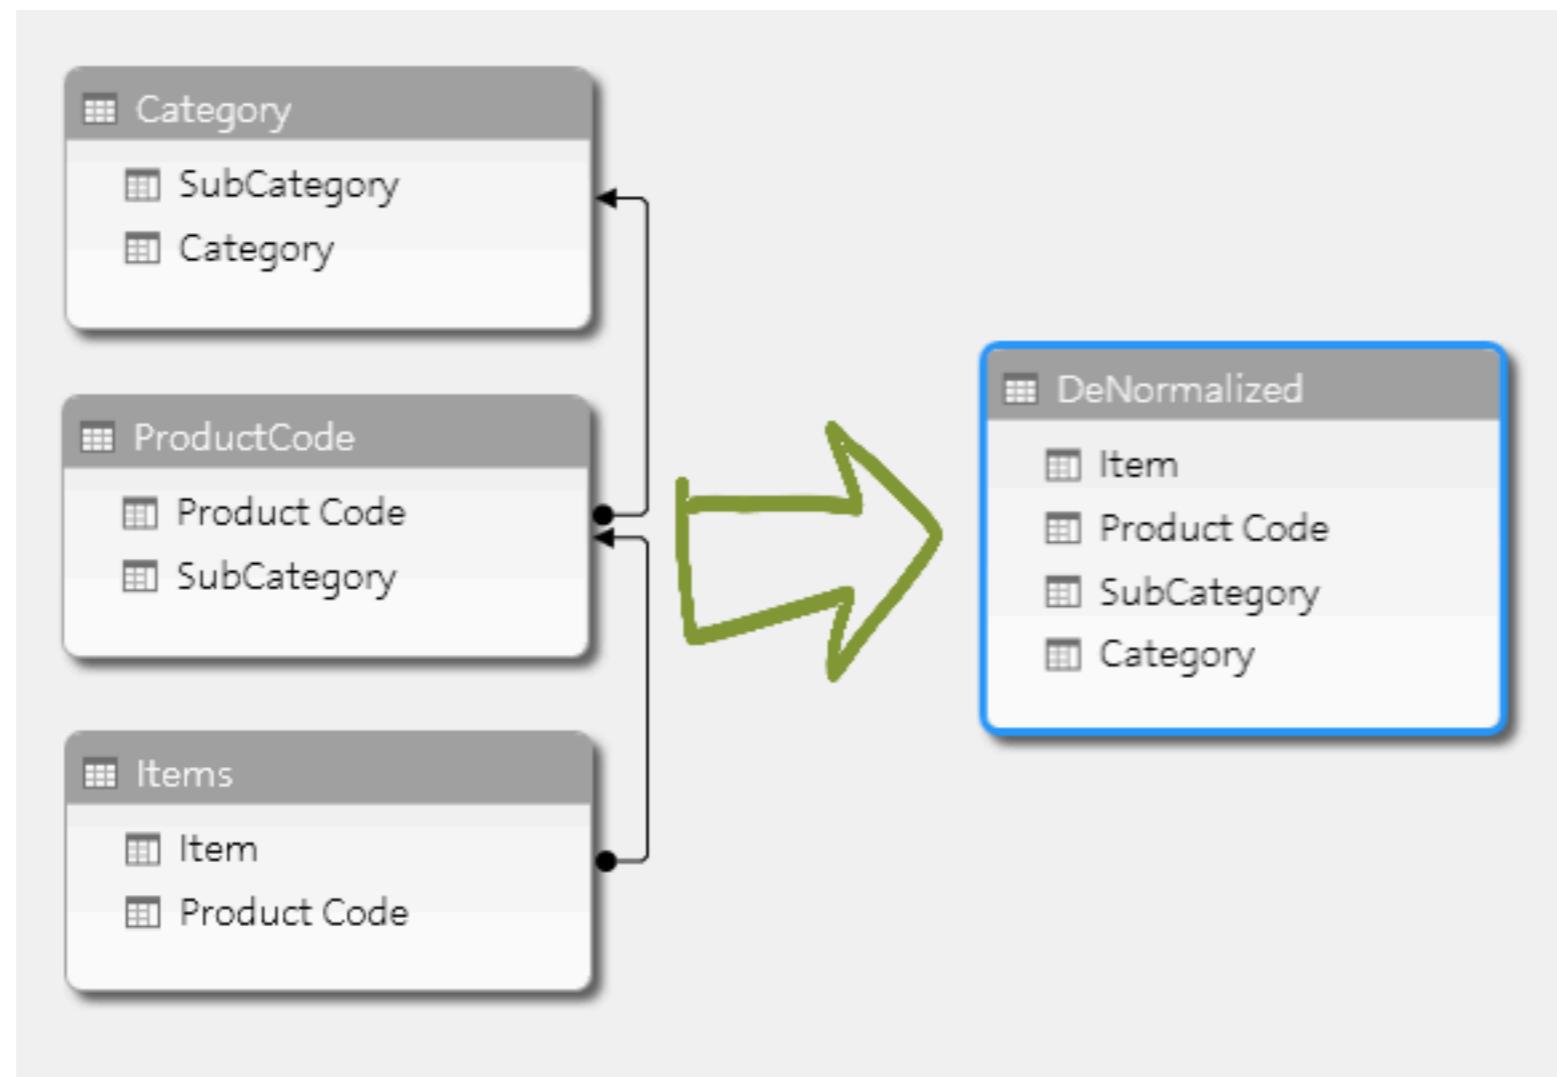
Materialized views



# Redundant column/Pre-joining tables



# Redundant column/Pre-joining tables



# Horizontal splitting

| <b>Student_ID</b> | <b>Department_ID</b> | <b>Student_Name</b> |
|-------------------|----------------------|---------------------|
| 1                 | 1                    | Alex                |
| 2                 | 2                    | Marie               |
| 3                 | 1                    | Sam                 |
| 4                 | 3                    | Sara                |

The diagram illustrates the process of horizontal splitting a single table into three smaller tables based on department categories. Three arrows point from the original table to the three resulting tables.

**Computer Science**

| <b>Student_ID</b> | <b>Department_ID</b> | <b>Student_Name</b> |
|-------------------|----------------------|---------------------|
| 1                 | 1                    | Alex                |
| 3                 | 1                    | Sam                 |

**Chemistry**

| <b>Student_ID</b> | <b>Department_ID</b> | <b>Student_Name</b> |
|-------------------|----------------------|---------------------|
| 2                 | 2                    | Marie               |

**Botany**

| <b>Student_ID</b> | <b>Department_ID</b> | <b>Student_Name</b> |
|-------------------|----------------------|---------------------|
| 4                 | 3                    | Sara                |



# Vertical splitting



# Add derived columns

| <b>Student_ID</b> | <b>Name</b> |
|-------------------|-------------|
| 1                 | Alex        |
| 2                 | Marie       |

| <b>Student_ID</b> | <b>Assignment_ID</b> | <b>Mark</b> |
|-------------------|----------------------|-------------|
| 1                 | 1                    | 20          |
| 1                 | 2                    | 35.50       |
| 2                 | 1                    | 45          |
| 2                 | 2                    | 45          |



| <b>Student_ID</b> | <b>Name</b> | <b>Total_Marks</b> |
|-------------------|-------------|--------------------|
| 1                 | Alex        | 55.5               |
| 2                 | Marie       | 90                 |



# Pros of data denormalization

Improve query performance (UX)

Reduce complexity of data model

Improve application scalability

Generate data report faster



# Cons of data denormalization

Increase data redundancy

Inconsistency between data sets

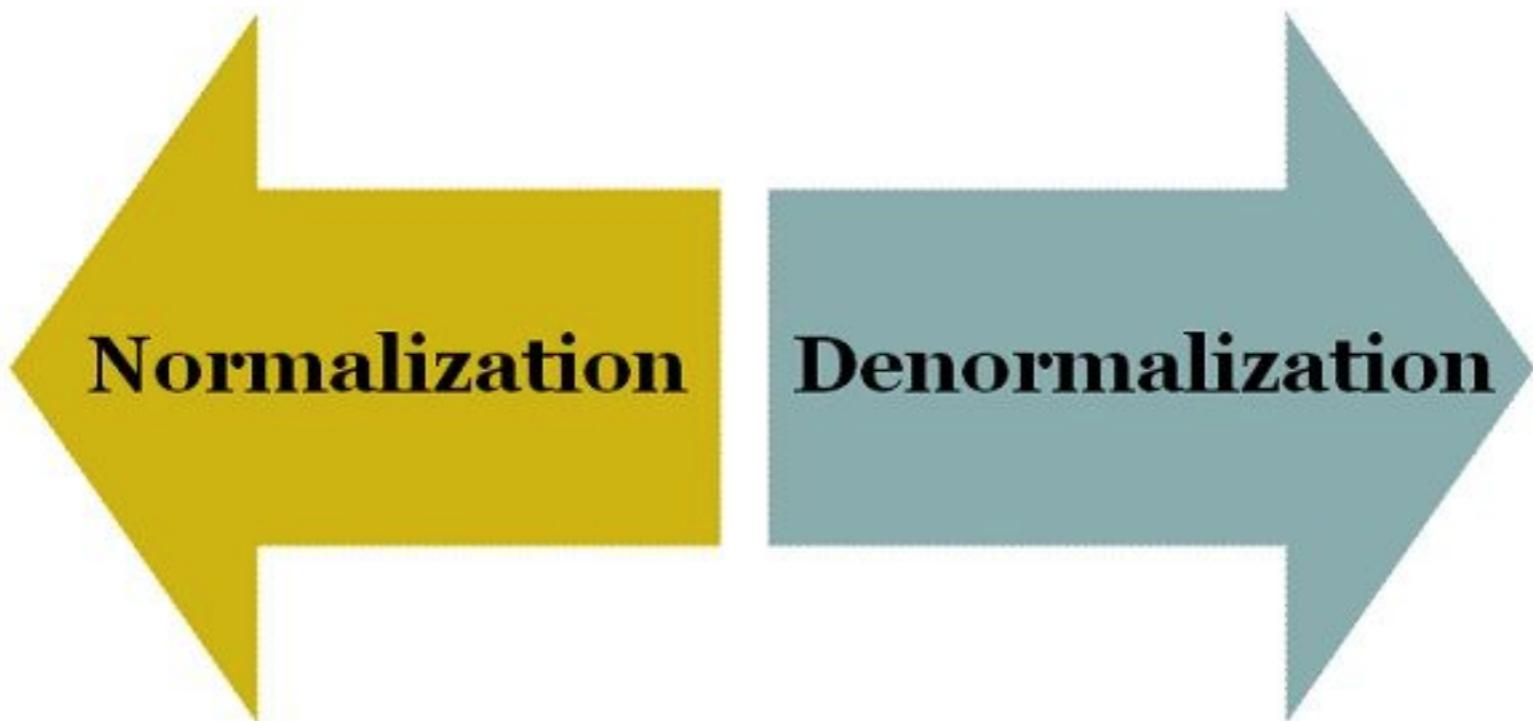
Require more disk space (split tables)

More data model (hard to maintain)

Difficult to insert and update data

Maintenance costs





# Workshop

